

Foundations of High Performance Computing

Luca Tornatore

Code optimization

Part III – Coding

Outline

General recap and some more things..

- Memory optimizations
- Loops optimizations
- Other optimizations

First things first

First things first

“Optimization” may be a tag for several different concepts, as we have seen at the beginning of this course.

Many concurrent facts and factors must be kept together, and it is quite difficult to give general statements about which ones are more fundamental.

Top-level design plays a key role, as well as algorithm choice and implementation details.

Sometime even a single line of code – for instance a prefetching – may have brilliant consequences

...

First things first

However, unless you are seeking some extreme performance, as a general guideline just keep in mind that “optimization” reads

“let the compiler squeeze the maximum from your code”

Compilers are quite good indeed, and know the hardware they are running on better than you (99% of times).

So, as first, just learn how to :

- write non-obfuscated code
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures

First things first

- write non-obfuscated code
 - avoid memory aliasing
 - make it clear what a variable is used for and when
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures

First things first

- write non-obfuscated code
- design a good data structure layout
 - be cache-conscious
 - be NUMA-conscious
 - avoid race conditions in multi-threaded codes
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures

First things first

- write non-obfuscated code
- design a good data structure layout
- design a “good” workflow
 - compiler will be able to optimize branches and memory access patterns
 - prefetching will work better
 - make it easier to use multi-threading
- take advantage of the modern out-of-order, super-scalar, multi-core architectures

First things first

- write non-obfuscated code
- design a good data structure layout
- design a “good” workflow
- take advantage of the modern out-of-order, super-scalar, multi-core architectures
 - let the compiler exploit pipelining through operation ordering and unloop
 - let the compiler exploit the vectorization capabilities of CPUs
 - think task-based, data-driven

Some C-specific hints

Storage class

May load/store latency of heavily used variables

- **extern**
Global variables, they exist forever
- **auto**
Local variables, allocated on the stack for a limited scope, and then destroyed. They must be initialized
- **register**
Suggests that the compiler puts this variable directly in a CPU register

Variable qualifiers

- **const**
Global variables, they exist forever
- **volatile**
Indicates that this variable can be accessed from outside the program.
- **restrict**
A memory address is accessed only via the specified pointer

Some C-specific hints

:: note about **restrict** qualifier ::

```
void my_function( double *a, double *b, int n)
{
    for( int i = ; i < n; i++ )
        a[ i ] = s * b[ i - 1 ];
}
```

The compiler can not optimize the access to **a** and **b** because it can not assume that **a** and **b** are pointing to the same memory locations.
That is called *aliasing*, formally forbidden in fortran: which is the reason why in some cases fortran may compile in faster executables.

Help your C compiler in doing the best effort, either writing a clean code or using **restrict** or using **-fstrict-aliasing** **-Wstrict-aliasing** options.

→ code snippet :: **pointers_aliasing.c**

Some C-specific hints

Memory allocation

We have already discussed that... try to **allocate contiguous memory** in large bunch, and to **reuse it** efficiently.

C++

That is a big topic, there's a lot of literature on it.

You're learning it in a dedicated course, so you know about that.

Let me just stress that a common illusion among non-professionals about C++ is that the compiler can clearly seen through all the abstraction that an advanced C++ programming style contains.

However, C++ should be preferably seen as a language that enables complexity management: most of its sophisticated features are not well suited for a extremely efficient low-level code.

Use of the compiler

Language standard

Some constructs can be better implemented if you specify a specific standard to the compiler.

Architecture specification

The compiler is able to detect the architecture it runs on, but – to maximize executable's compatibility – it will not turn on very specific optimization unless you tell him so.

Read *carefully* the manual of your compiler.

In `gcc`, for instance, `-march=native` turns on a lot of sophisticated optimizations

Use of the compiler

Optimization level: -On

It is not granted that **-O3**, although often generating a faster code, is what you really need.

For instance, sometimes expensive optimizations may generate more code that on some architecture (e.g. with smaller l caches) run slower, and using **-Os** may bring surprising results.

Take into accounts that modern compilers allows for local specific optimizations or compilation flags. In gcc for instance:

```
__attribute__ ((__option__ ("...")))  
__attribute__ ((optimize(n)))
```

Special options

Not all the optimizations are enabled via **-On**, even for the highest *n*. Check your manual.

Use of the compiler

Profile-guided optimization

Compilers (**gcc**, **icc** and **clang**) are able to instrument the code so to generate run-time information to be used in a subsequent compilations. Knowing the typical execution patterns enables the compiler to perform more focused optimizations, especially if several branches are present.

For **gcc**:

```
gcc -fprofile-arcs  
< ... run ... >  
gcc -fbranch-probabilities
```



Specific for branch prediction

```
gcc -fprofile-generate  
< ... run ... >  
gcc -fprofile-use
```



More general; enables also
-fprofile-values **-freorder-functions**

Outline

■ Memory optimizations

- > the importance of memory access
- > the importance of the cache
- > cache access optimization in loops
- > cache-oblivious algorithms
- > cache-oblivious data structures

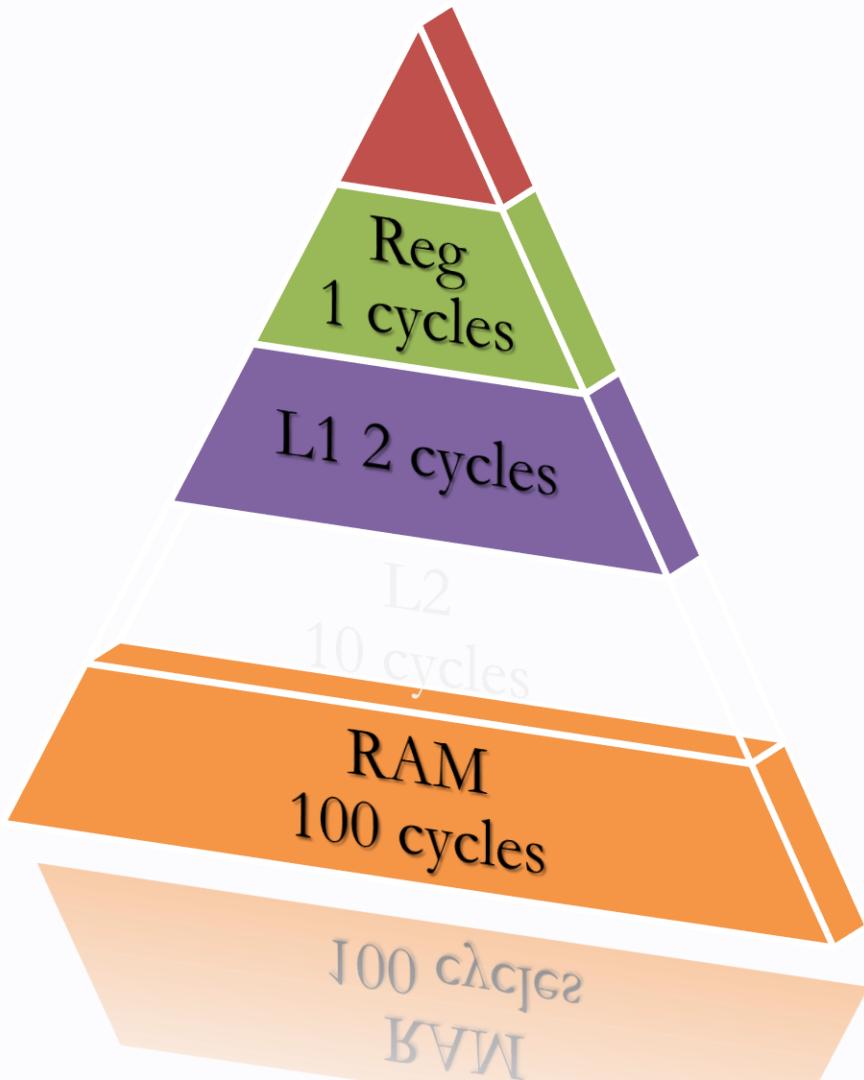
■ Loops optimizations

■ Other optimizations

The memory access
Again on cache

The importance of memory access

Impact of cache misses on average memory access time



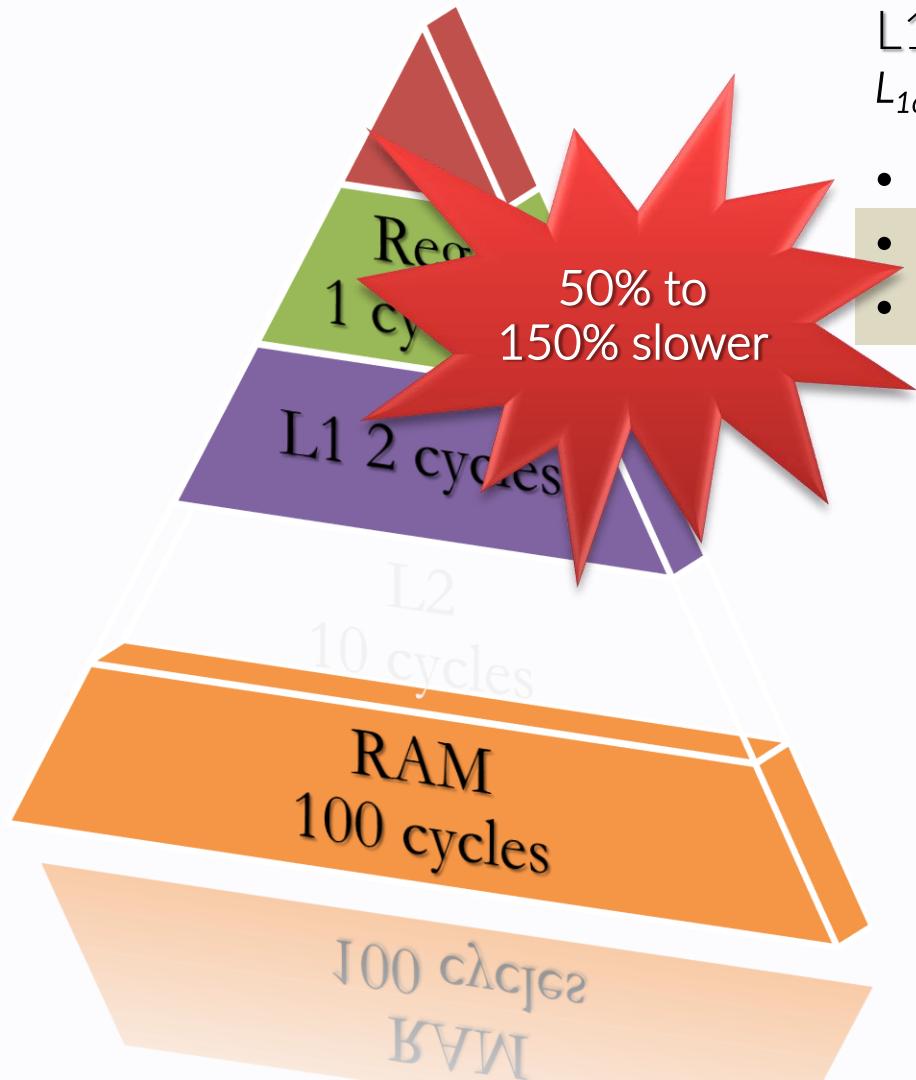
L1 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times RAM_{\text{cost}}$$

- 100% L1 hit → 2 cycles
- 99% L1 hit → 3 cycles
- 97% L1 hit → 5 cycles

The importance of memory access

Impact of cache misses on average memory access time



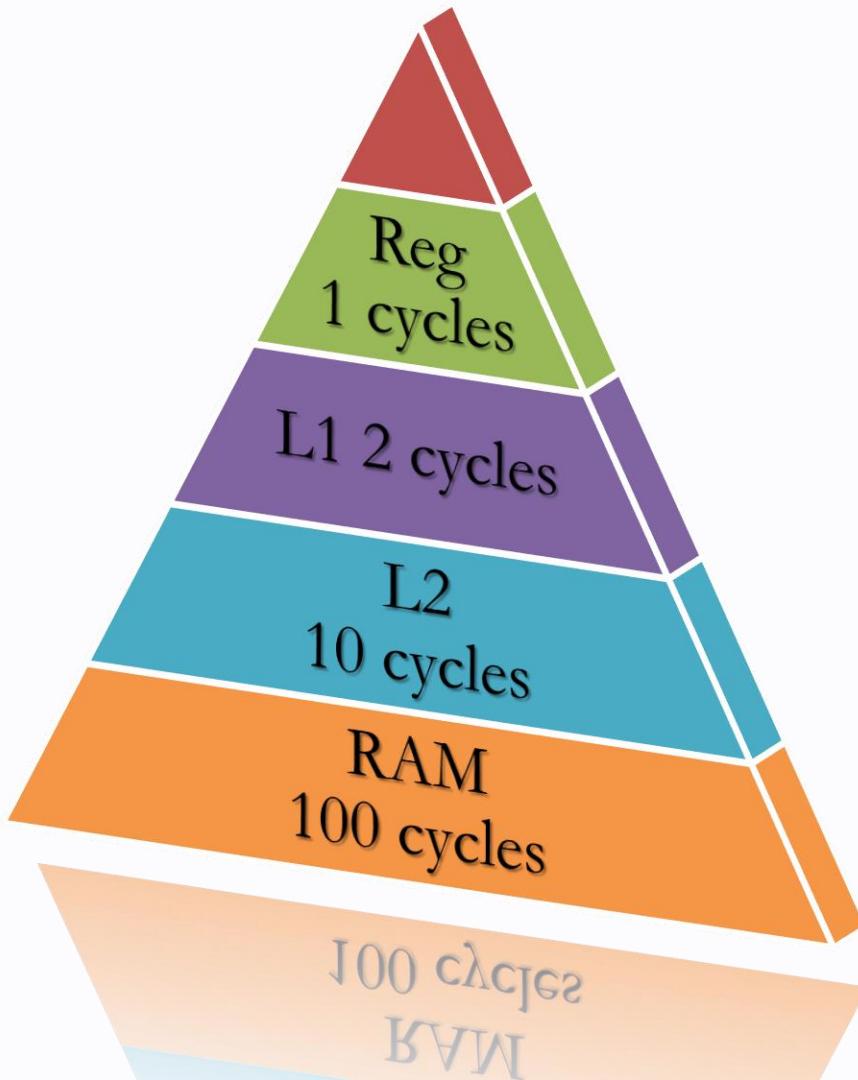
L1 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times RAM_{\text{cost}}$$

- 100% L1 hit → 2 cycles
- 99% L1 hit → 3 cycles
- 97% L1 hit → 5 cycles

The importance of memory access

Impact of cache misses on average memory access time



L1 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times RAM_{\text{cost}}$$

- 100% L1 hit → 2 cycles
- 99% L1 hit → 3 cycles
- 97% L1 hit → 5 cycles

L1 cache + L2 cache + RAM

$$L_{1\text{cost}} + \text{Miss}_1 \times (L2_{\text{cost}} + \text{Miss}_2 \times RAM_{\text{cost}})$$

- 100% L1 hit → 2 cycles
- 99% L1 hit, 100% L2 hit → 2.1cycles
- 97% L1 hit, 100% L2 hit → 2.3 cycles
- 90% L1 hit, 97% L2 hit → 3.3 cycles

What would happen if you had not separate L1 caches for data and instructions?

The importance of cache

When the cache is hit and when it is not: a simple model

Consider a simple direct mapped **16 byte** data cache with **two cache lines**, each of size 8 bytes (two floats per line)

Consider the following code sequence, in which the array **X** is cache-aligned (that is, **X[0]** is always loaded into the beginning of the first cache line) and accessed twice in consecutive order:

```
float X[8];  
for(int j=0; j<2; j++)  
    for(int i=0; i<8; i++)  
        access(X[i]);
```



The hit-miss pattern is : MH MH MH MH MH MH MH MH
the miss-rate is 50% (the first miss is compulsory miss).

The importance of cache

Let's consider another code sequence that access the array twice as before, but with a strided access

```
float X[8];
for(int j=0; j<2; j++)
{
    for(int i=0; i<7; i+=2)
        access(X[i]);
    for(int i=1; i<8; i+=2)
        access(X[i]);
}
```

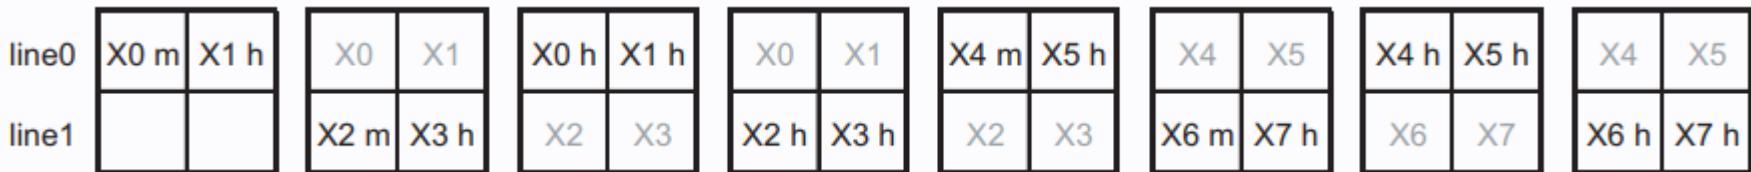


The hit-miss pattern now is : MM MM MM MM MM MM MM MM, the miss-rate is 100%

The importance of cache

Finally, consider a third code sequence that again access the array twice:

```
for(int i = 0; i < 2; i++)  
    for(int k = 0; k < 2; k++)  
        for(int j = 2*i; j < (i+1)*2; j ++)  
            access(X[ j ]);
```

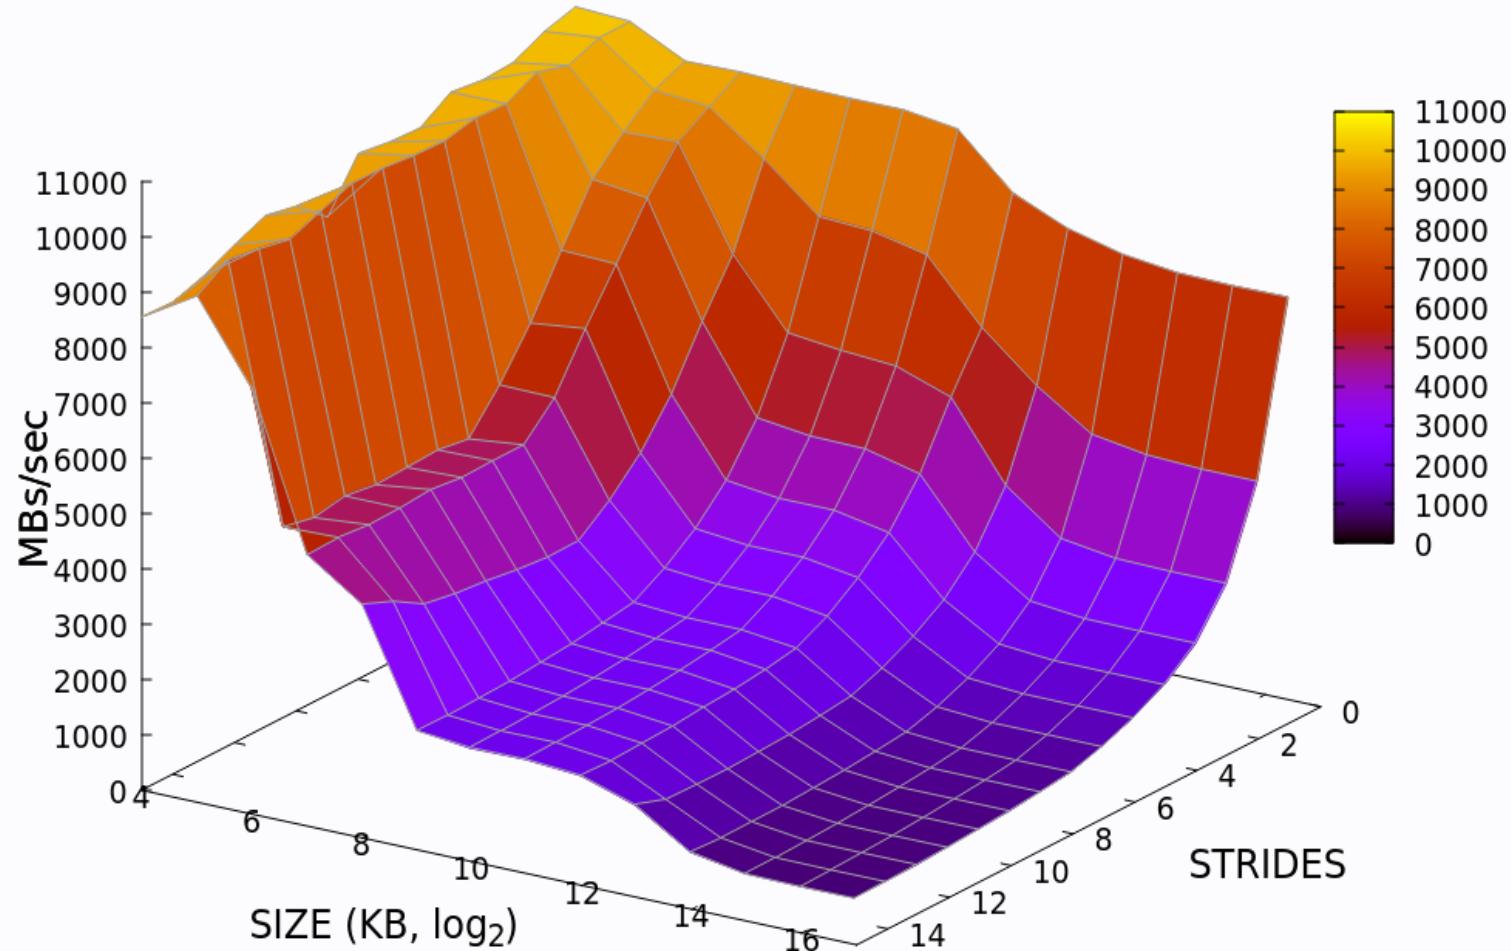


The hit-miss pattern now is : MH MH HH HH MH MH HH HH,
the miss-rate is 25%

The main message is: memory access pattern is of primary importance

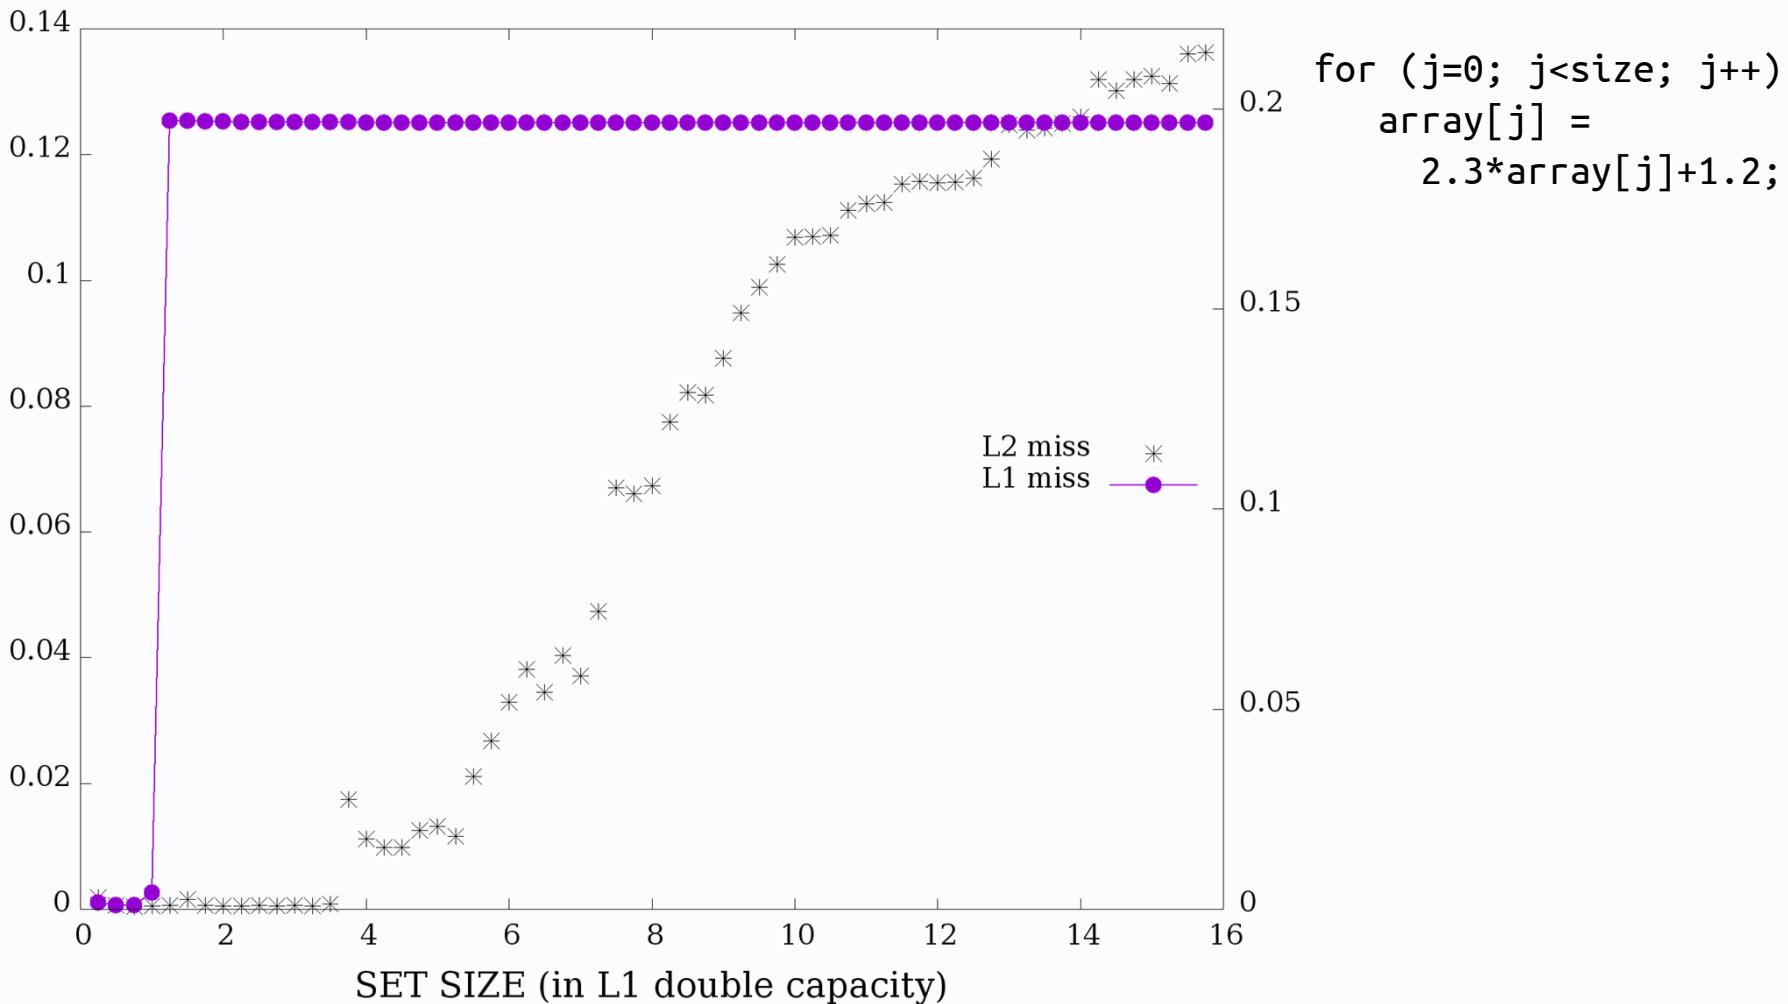
The importance of cache

The result is.. the memory mountain



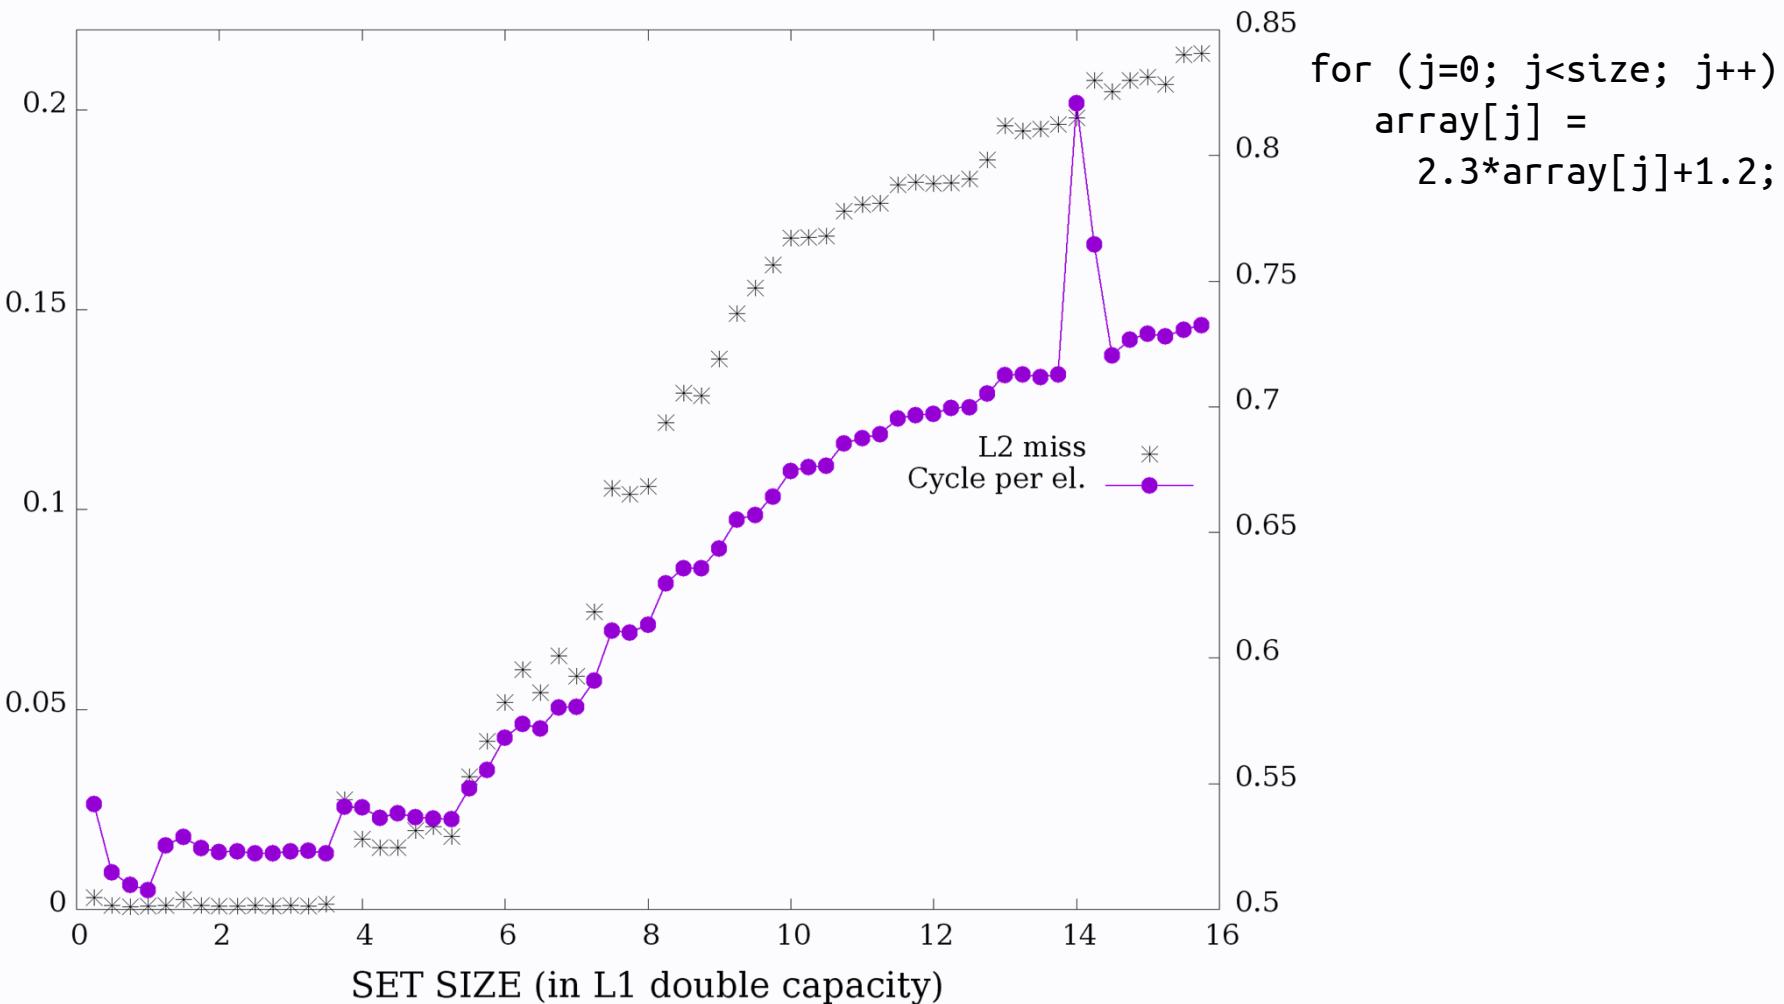
The importance of memory access pattern

Let's find our L1 - and L2- cache sizes



The importance of memory access pattern

..and the effect on cycles-per-operation



The importance of memory access pattern

A Worked example you have already seen: dense MATRIX TRANSPOSE

Naïve version:

```
for(int row = 0; row < N; row++)
    for(col = 0; col < N; col++)
        A [ col*N + row ] = B [ row*N + col ];
```

NOTE: strided access to either A or B is unavoidable.

However: is it better to have it either on *read* or on *write* ?

The importance of memory access pattern

A Worked example you have already seen: dense MATRIX TRANSPOSE

Naïve version:

```
for(int row = 0; row < N; row++)
    for(col = 0; col < N; col++)
        A [ col*N + row ] = B [ row*N + col ];
```

NOTE: strided access to either A or B is unavoidable.

However: is it better to have it either on *read* or on *write* ?

Due to write-allocate transactions in the cache, **strided writes are more expensive than strided loads.**

The importance of memory access pattern

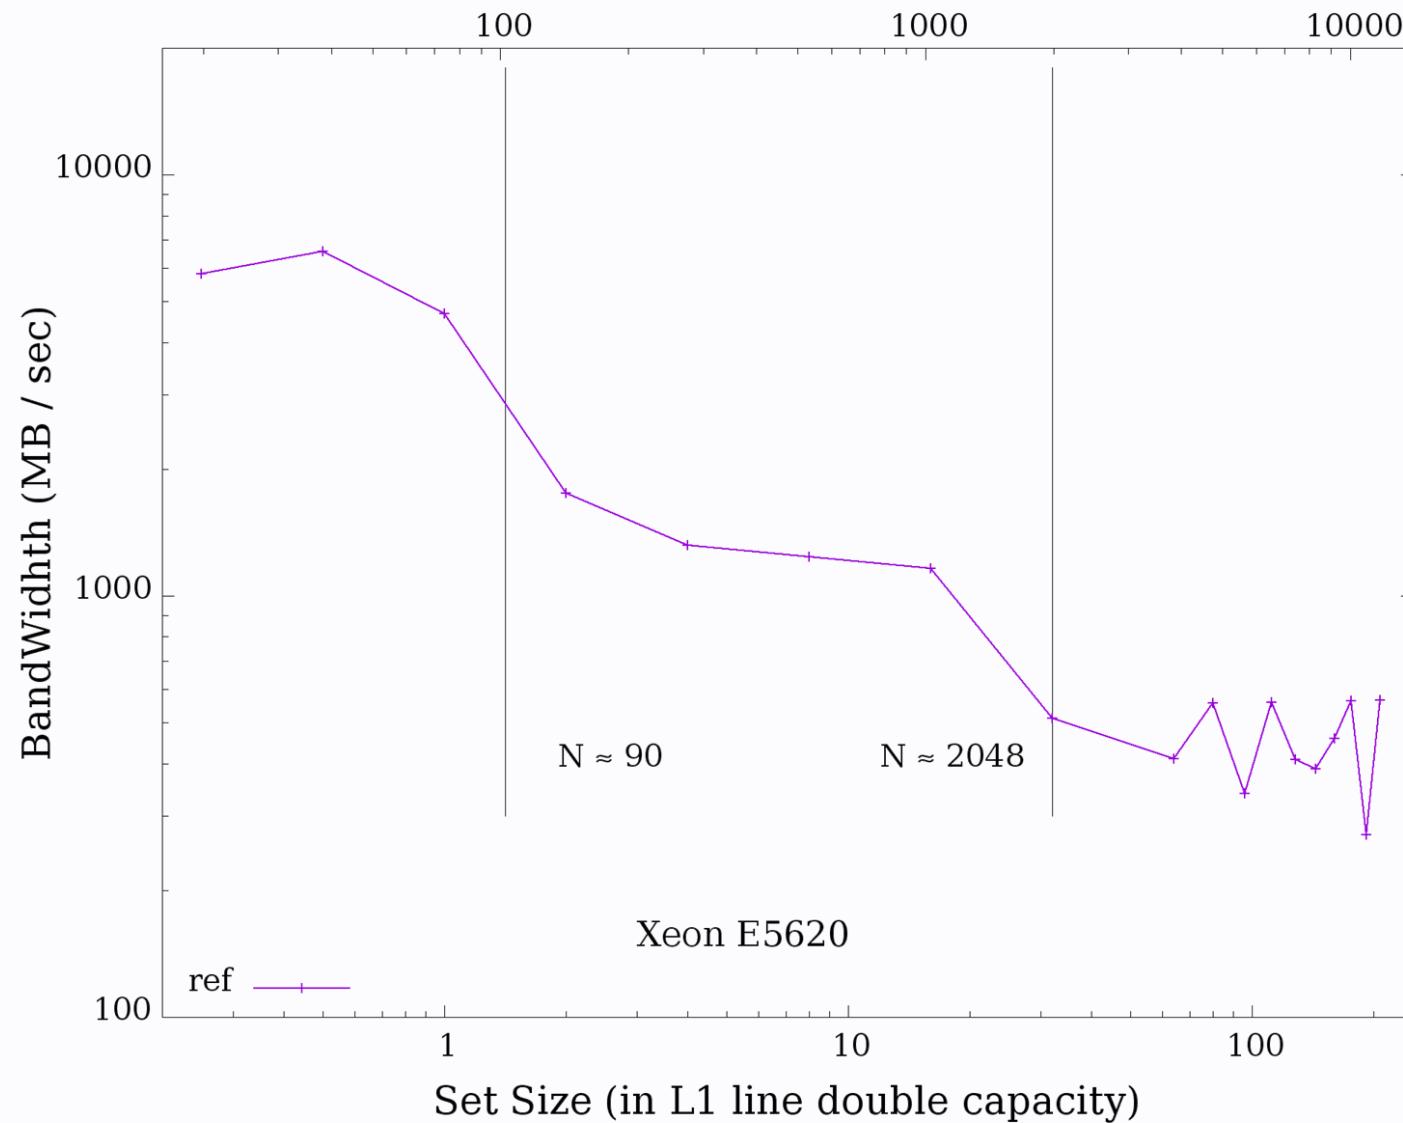
A Worked example: dense MATRIX TRANSPOSE

Let C be the cache size and L_C the cache line size, we should expect 3 different regimes:

1. $2 \times N^2 < C$
both matrices can fit in the cache, traversal order and locality does not impact on performance (bandwidth \sim maximum)
2. $N \times L_C < C$
strided write is alleviated by fraction of column fitting in the cache
3. $N > C \times L_C$
Each access to A determines a cache miss and a *write-allocate*.
A sharp drop in performance is expected since basically only one entry per line will be used.

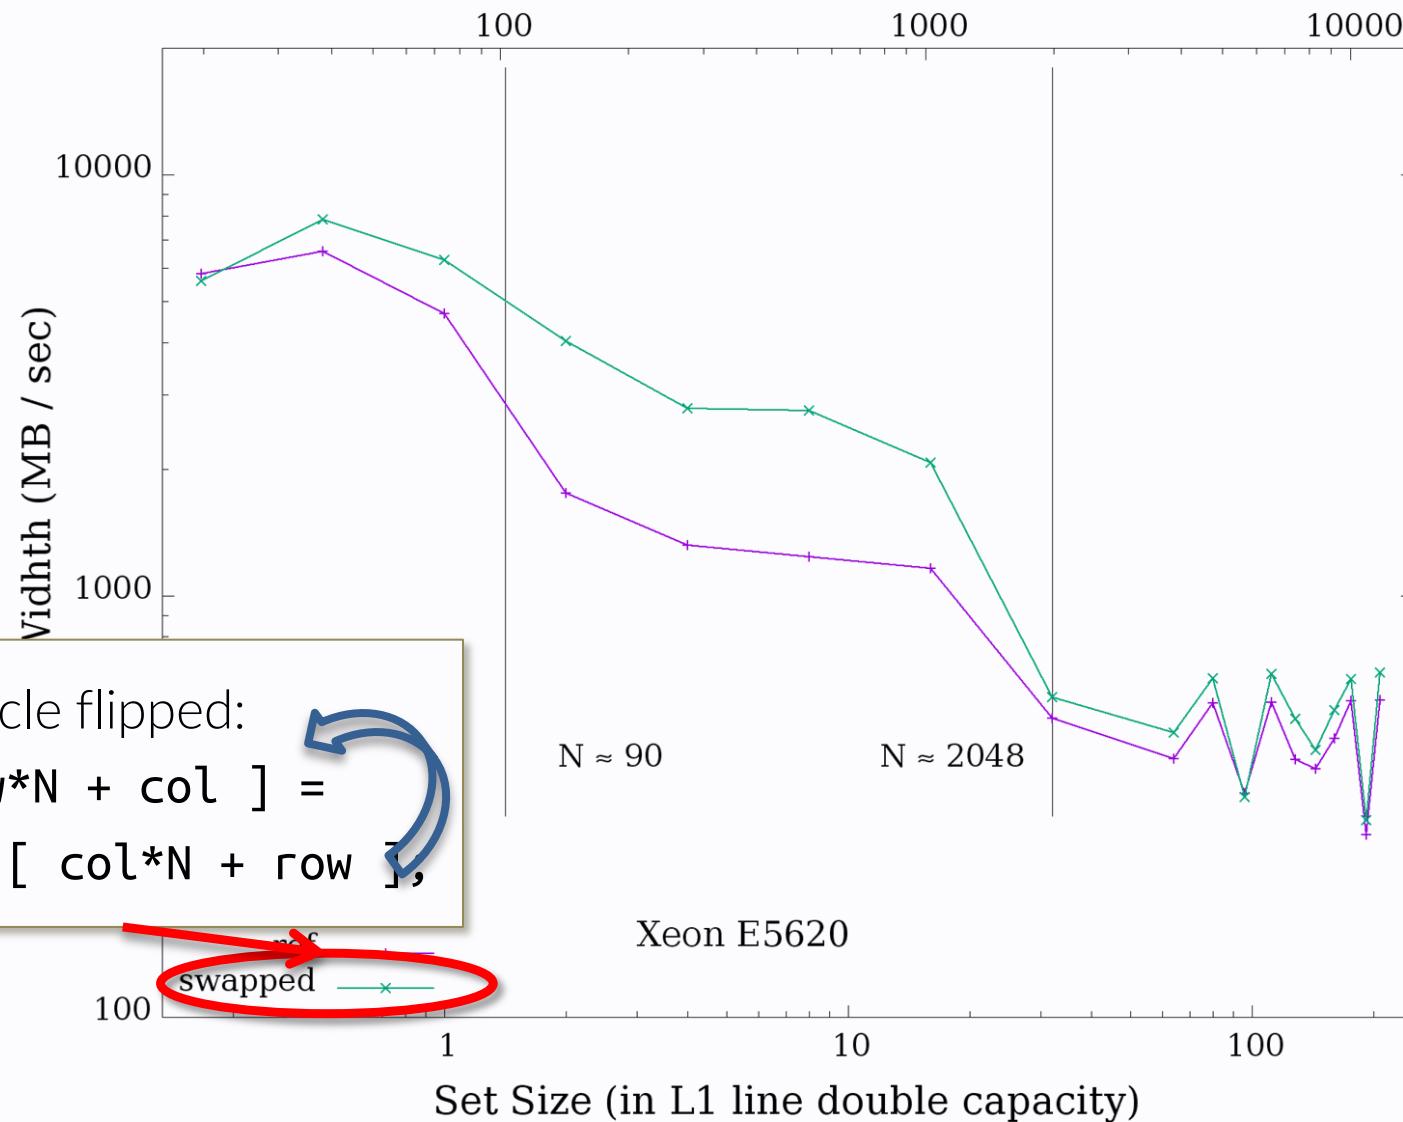
The importance of memory access pattern

..more or less, you already know from previous lectures..



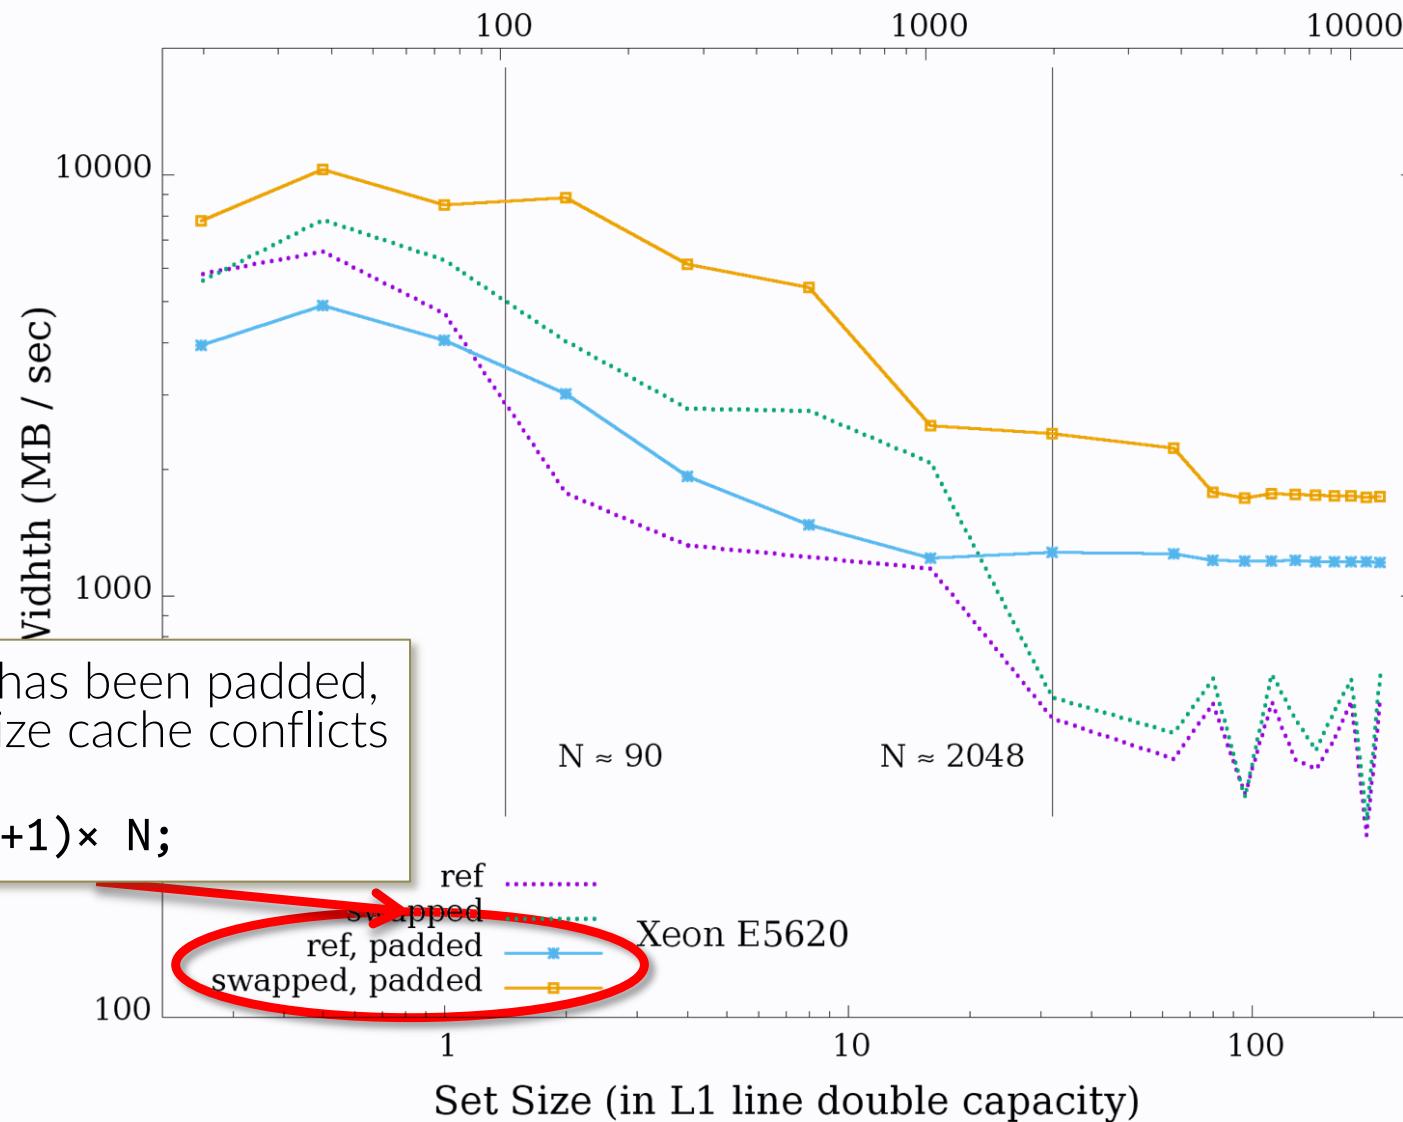
The importance of memory access pattern

Matrix transpose: avoiding strided access on write



The importance of memory access pattern

Matrix transpose: avoiding cache associativity conflicts



Cache resonance

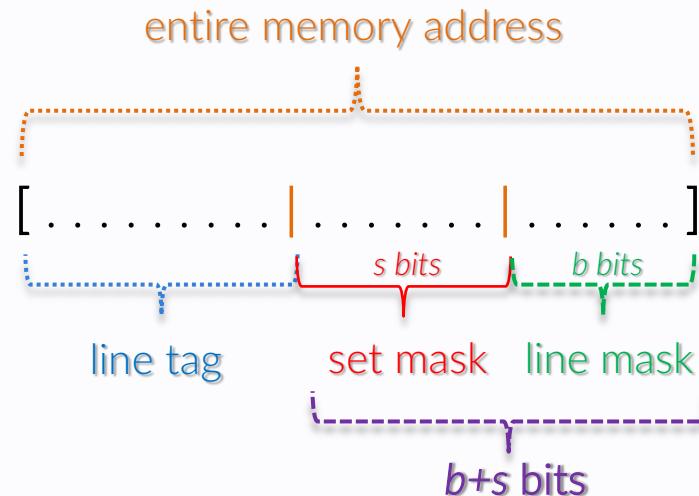
Cache associativity conflicts (or “resonances”): what are they ?

You know that your cache of size c bytes is w -way associative: it means that your cache is made up by lines of size L bytes, grouped in w -sized sets.

Tipical figures nowadays are: $L = 64B (=2^b$ bytes), $w = 4-8$, $c = 32KB$ (i.e. there are $c / (L \times w) = 64-128 (= 2^s$ sets).

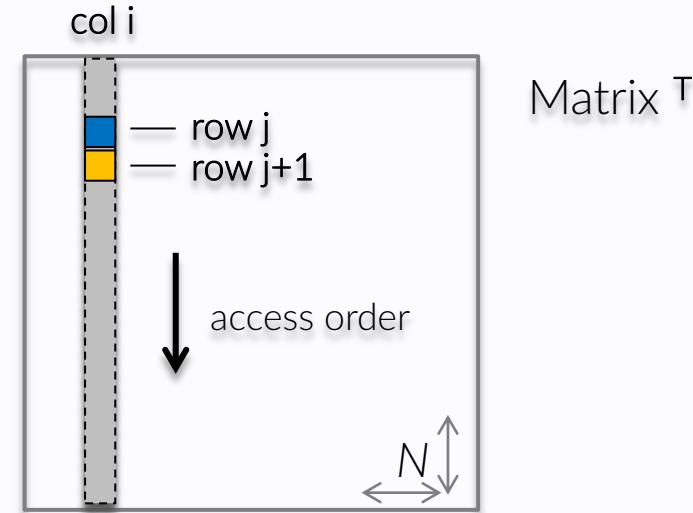
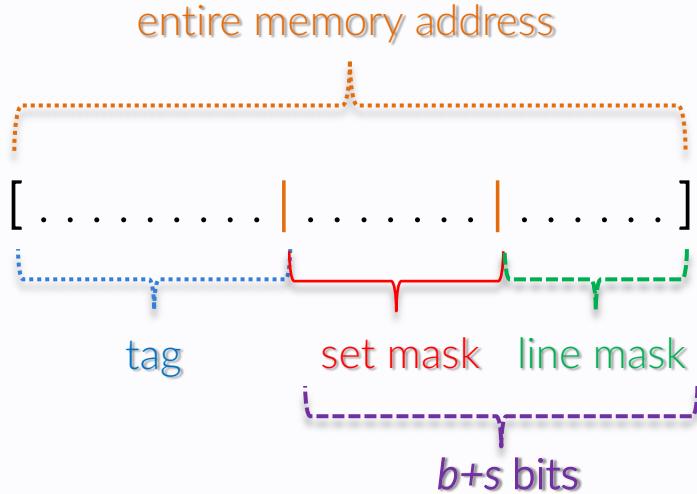
“ways” of associativity

To map the memory addresses into the cache, the first $b+s$ bits are used to determine to what byte in a free line of what cache set that address will be stored.



Cache resonance

Cache associativity conflicts (or “resonances”): what are they ?



Accessing the element $i, j+1$ of the transposed matrix, the stride with respect to previously accessed element i, j will amount to:

$$\text{offset} = N \times W \text{ bytes}$$

where N is the matrix size and W is the type size (e.g. double, 8 bytes).

If N is a power of 2, $N = 2^n$; if $W = 2^d$ bytes.

$$\text{offset} = 2^{n+d} \text{ bytes}$$

then if $n+d > b+s$, the address of $i, j+1$ is mapped in the same set than i, j and then (since we assume $N \gg w$) at least every w accesses there is a cache conflict.

Cache resonance: padding

Padding is a simple but effective cure for the issue of cache resonance.

Adding one column to the transposed matrix, which then is $(N+k) \times N$, means that when the element $i, j+1$ is accessed, the stride with respect to previously accessed element i, j amounts to:

$$\text{offset} = (N+k) \times W \text{ bytes} = 2^{n+d} + k2^d \text{ bytes}$$

If N is a power of 2, $N = 2^n$, with $W = 2^d$ bytes,

$$\text{offset} = 2^{n+d} + k2^d \text{ bytes}$$

then if

$$2^{b+s} < 2^{n+d} + k2^d$$

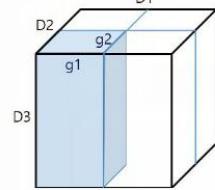
the address of $i, j+1$ is mapped on a different cache set, alleviating the cache conflicts

The importance of data layout

Cache associativity conflicts (or “resonances”): what are they ?

3.2 Padding for Set-associative Caches

To extend this result to the set-associative case for all i , $1 \leq i \leq d-1$, we introduce the characteristic number g_i of dimension i with respect to the cache size. Intuitively, as depicted (square at right) for $A = 4$, we will establish that if the enclosed tile $g_1 \times \dots \times g_{d-1} \times D_d$ is free of self-interference conflicts in a direct-mapped cache, then the A -times larger tile $D_1 D_2 \dots D_d$ is free of self-interference conflicts in an A -associative cache of the same capacity.



Theorem 2 (Associative cache). *Consider a set-associative cache of capacity $C = SAB$. For all $1 \leq i \leq d-1$, let $g_i = \gcd(S/\prod_{1 \leq k \leq i-1} g_k, N_i)$. A loop nest whose tiles have a d -dimensional array footprint can fully utilize the cache and remain free of self-interference if and only if the following conditions are met:*

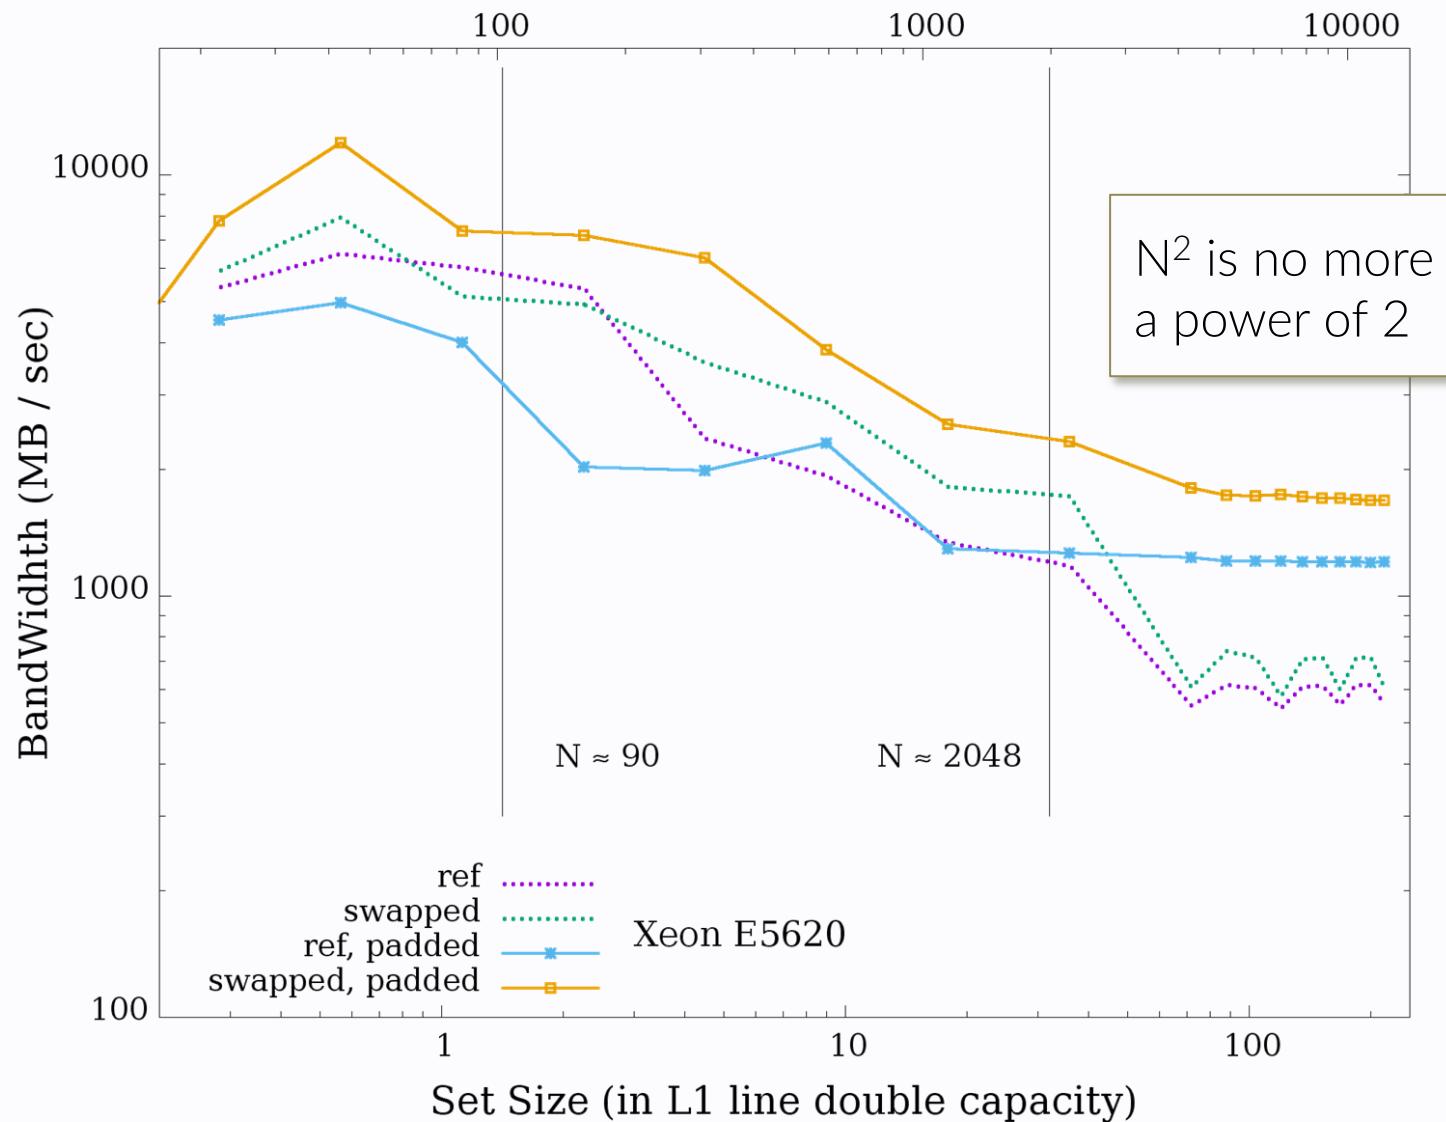
1. $\forall i, 1 \leq i \leq d-1, \exists j, 1 \leq j \leq i, \prod_{1 \leq k \leq i} g_k$ divides $D_j \prod_{1 \leq k \leq j-1} g_k$.
2. $\exists i, 1 \leq i \leq d, S$ divides $D_i \prod_{1 \leq k \leq i-1} g_k$.

People are doing a lot of impressive stuff and research about this issue.

The previous one was a simplistic explanation and a very naive solution of the problem
(with, however, a very good return-on-investment)

The importance of data layout

Matrix transpose: avoiding cache associativity conflicts



Optimization of cache access in loops

Cache access optimization in loops

Loop classification

$$A_I = \frac{f(n)}{n}$$

Arithmetic Intensity: the **ratio** between the number of performed operations and the **amount** of the requested data.

1. $O(N) / O(N)$ → optimization potential limited
2. $O(N^2) / O(N^2)$ → some more opportunities for opt.
3. $O(N^3) / O(N^2)$ → significant optimization potential

Cache optimization: $O(N)$ / $O(N)$ loops

Ex 1-level loops: Scalar products, vector additions, sparse matrix-vector multiplication

Inevitably memory-bound for very large N ; in general, improvements come from *avoiding unnecessary operations and/or repeated memory accesses, increasing data reuse*

```
for(int j=0; j<2; j++)
```

```
    A[i] = B[i] × C[i]
```

```
for(int j=0; j<2; j++)
```

```
    Q[i] = B[i] + D[i]
```



```
for(int j=0; j<2; j++)
```

```
{
```

```
    A[i] = B[i] × C[i]
```

```
    Q[i] = B[i] + D[i]
```

```
}
```

[*loop fusion*]

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

Ex 2-level loops on N: dense matrix-vector mul, matrix transpos., matrix add, ...

Improvements comes again from increasing *data reuse*,
exploiting *locality* and avoiding unnecessary operations and
memory accesses.

Example of unrolling + fusion, starting from the following code:

```
for(int i=0; i < N; i++)
    for(int j=0; j<N; j++)
        C[i] += A[i][j] * B[j];
```

$N \times N$

N^2

$N \times N$

$\rightarrow 3 \times N^2$

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

STEP 1:

Avoid unnecessary load/store operations

```
for(int i=0; i < N; i++)
{
    c_temp = C[i];
    for(int j=0; j < N; j++)
        c_temp += A[i][j] * B[j];
    C[i] = c_temp;
}
```

Now it is more clear for the compiler that **C[i]** need to be loaded and stored only 1 time

→ $2 \times N^2 + N$

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

STEP 2:

Unroll outer loop and fuse in the inner loop; there is potential for vectorization.

```
for(int i=0; i < N; i += m)
    for(j = 0; j < N; j++)
    {
        b_temp = B[j];
        C[i] += A[i][j] * b_temp;
        C[i+1] += A[i+1][j] * b_temp;
        ...
        C[i+m] += A[i+m][j] * b_temp;
    }
/* NOTE: remainder loop ignored here */
```

$\rightarrow N^2 \times (1 + 1/m) + N$

N N² N × N / m

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

NOTE :: UNROLLING AND REGISTER SPILL

If m in the previous example is too large it results in a code bloating if the target CPU does not have enough registers to keep all the needed operands.

In this case, the CPU has to spill registers' content to cache and viceversa, slowing down the computation.

→ learn to inspect the compiler's log

A too much involve and obscure loop body may hamper the compiler to effectively perform *unroll & jam* optimizations targeted to the CPU it runs on.

→ hand code effort to clarify the code

→ hints / directives to the compiler

(directives are generally not portable across different compilers)

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

Sometimes no magic wand can cure the fact that you have to access N^2 memory locations.

For instance: in matrix transpose you have to access all the source matrix and all the destination matrix once.

Unroll & Jam strategy can bring benefits as long as the cache can hold N lines.

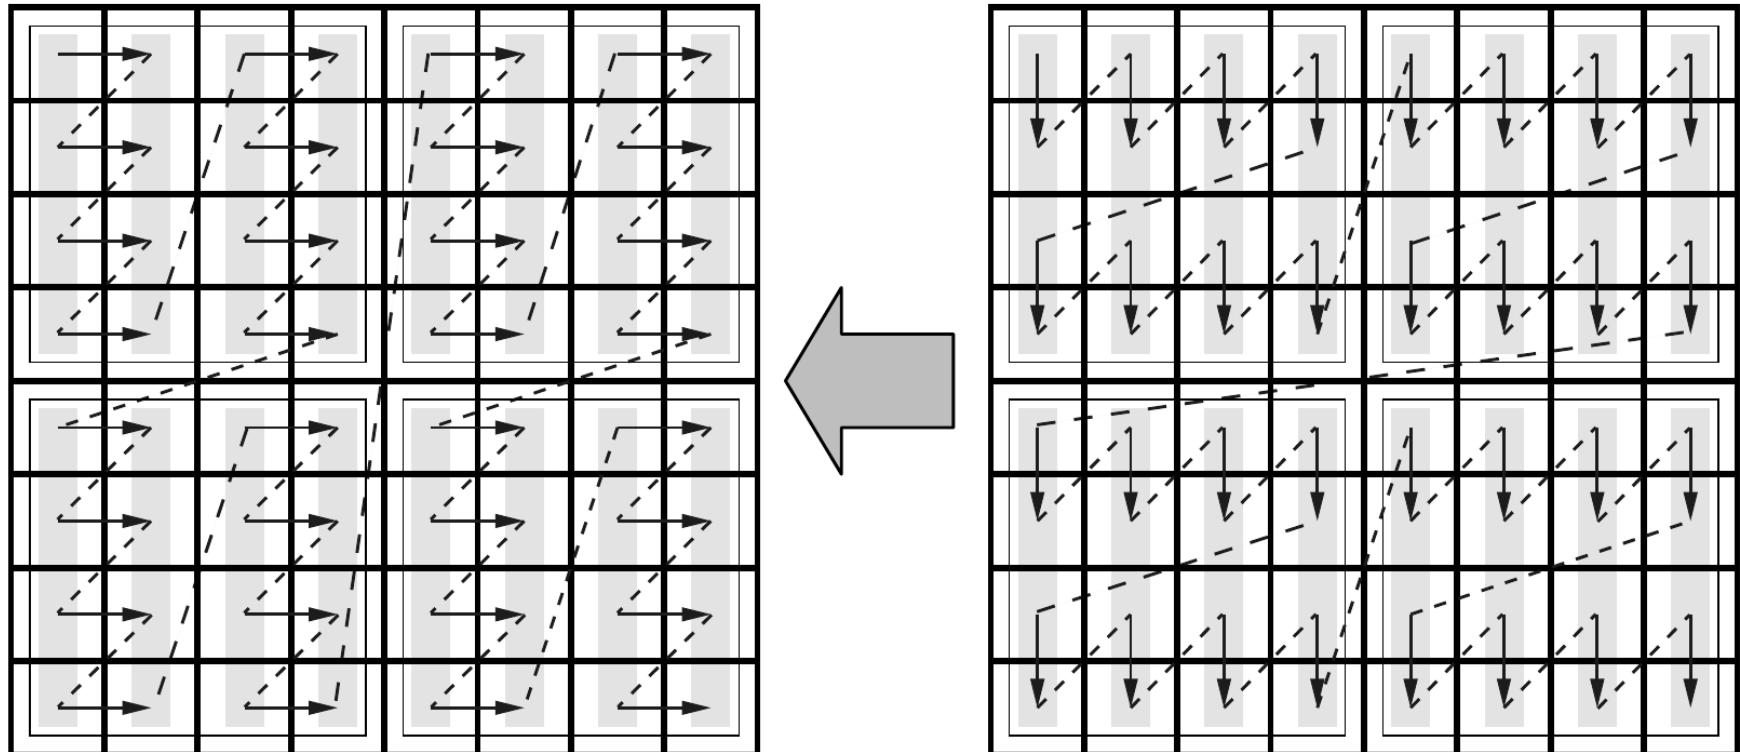
An L_C -way unrolling is too much aggressive and may easily result in register pressure.

Loop blocking is a good strategy that does not save memory loads but increase dramatically the cache hit ratio

Cache optimization: $O(N^2)$ / $O(N^2)$ loops

STEP 3:

Exploit full locality of referenced data, cut TLB misses by
accessing 2D arrays by blocks



Taken from: introduction to HPC for scientists and engineers

Cache optimization: $O(N^3)$ / $O(N^2)$ loops

These algorithms (ex: matrix-matrix multiplication or dense matrix diagonalization) are very good candidates for optimizations that lead performance very close to theoretical performance peak (MMM at the core of **linpack**).

Blocking, unroll&jam + vectorization of operations, reorganization of ops to exploit CPU's pipelines and out-of-order capability, are all used by extremely **specialized libraries**.

→ It is a brilliant idea to link those library instead of developing your own algorithm, unless some very special needs must be met.

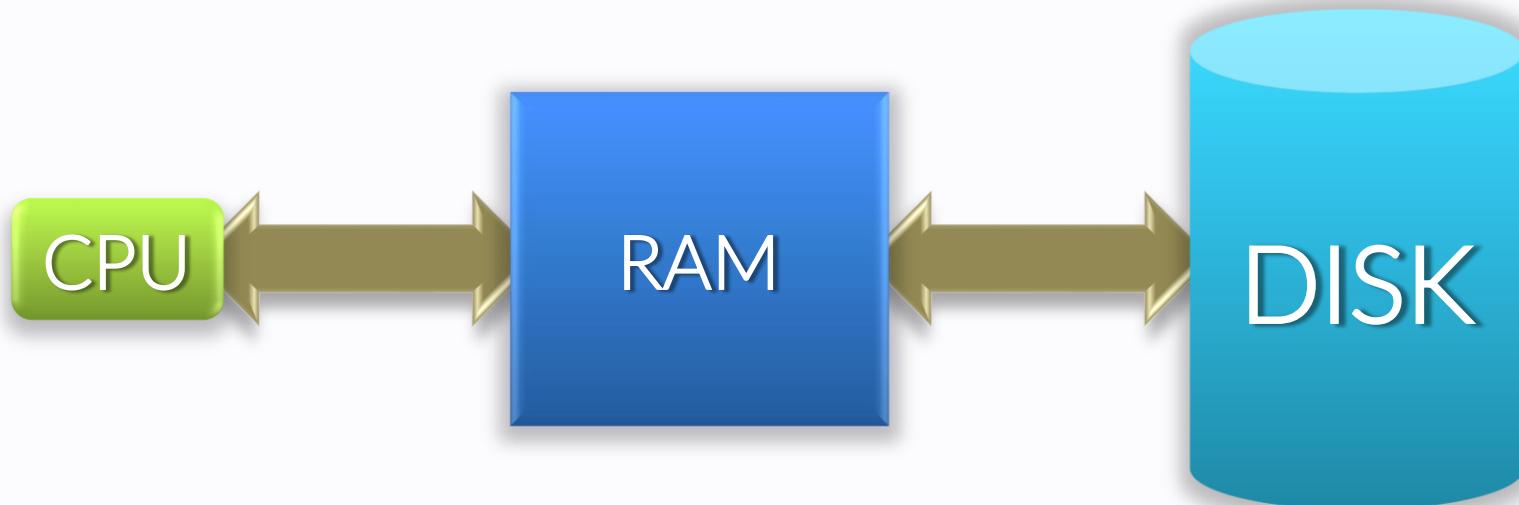
Cache-consciousness

cache-aware vs cache-oblivious

Cache – consciousness

In the good old days, computers were much simpler and much closer to a lovable Von Neumann architecture.

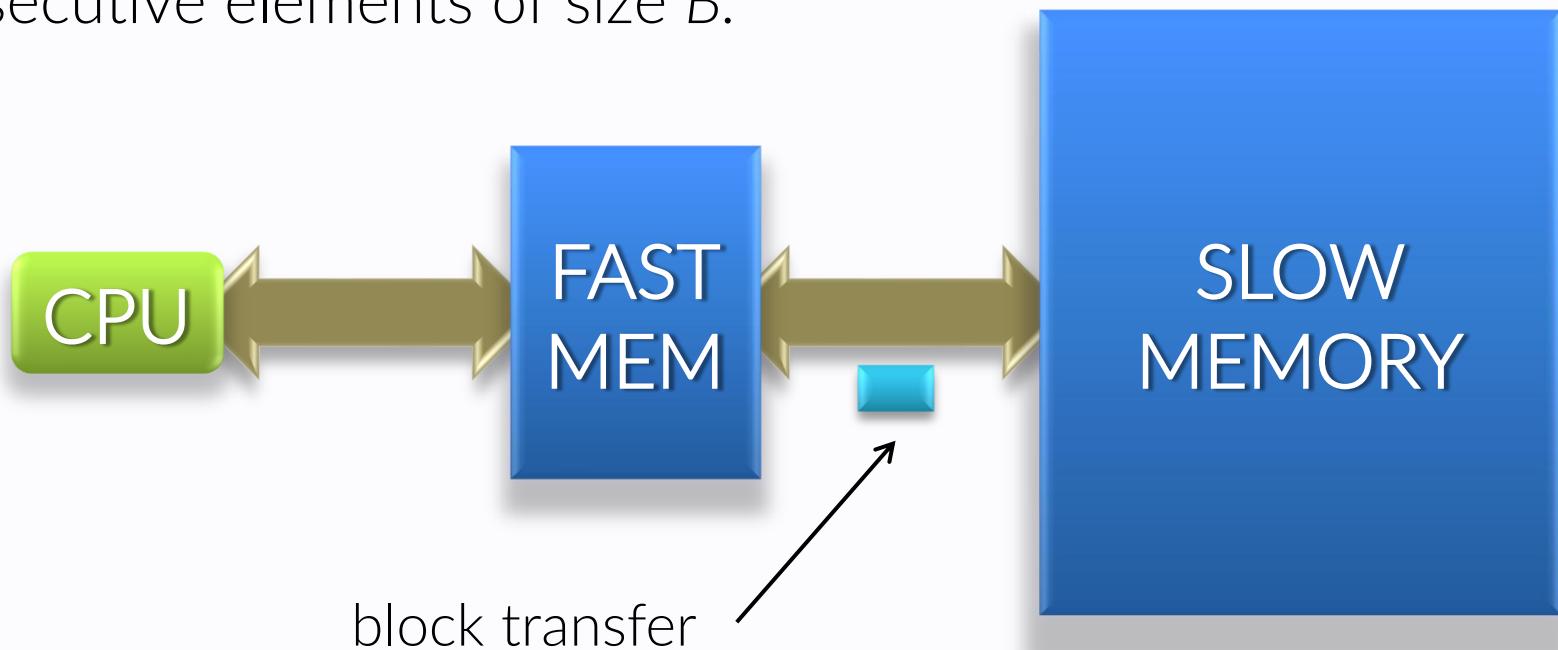
The **RAM MODEL** was more than enough to design and analyze algorithms, based on complexity measured as the *instructions* – the work $W(n)$ – needed for an input problem of a given size n .



Cache – consciousness

More realistic models became soon mandatory, and the most successful one was the simple two-level **I/O MODEL**, which was firstly shaped with memory-disk operations in mind.

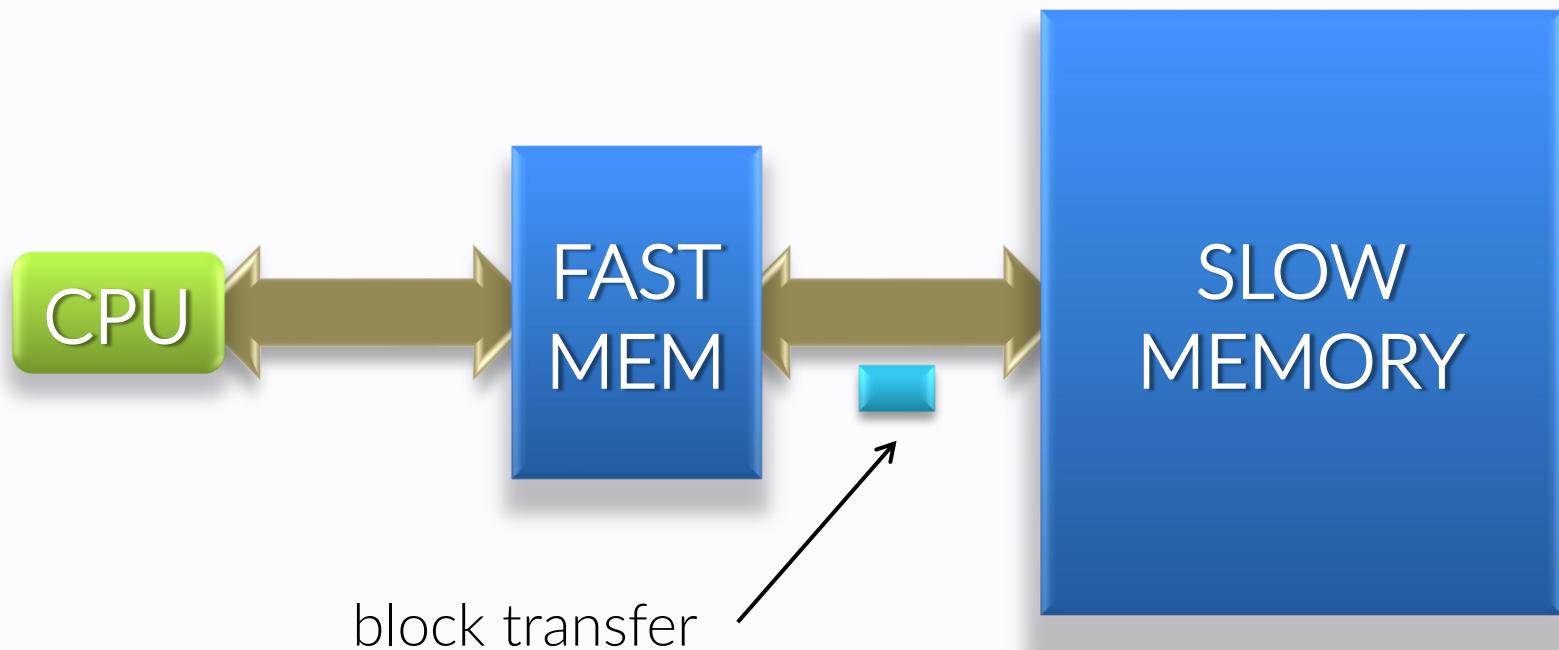
The hierarchy is assumed to consist in a fast memory of size M that exchange data with a slower infinite memory in bunches of consecutive elements of size B .



Cache – consciousness

The basic idea is to evaluate the code efficiency by the memory transfers it need to operate. It captures particularly well the case of problems whose size n largely exceeds the RAM size. When tuned to specific (M, B) , it is called **cache-aware** model.

Ex: comparison-based Sort _{M, B} (N) = $\Theta\left(\frac{N}{B} \log_{M/B} \frac{N}{B}\right)$ memory transfers



Cache – oblivious model

In 1999, Frigo et al. proposed the *cache-oblivious model* which did not required a precise knowledge of the memory hierarchy.

Basically, a cache-oblivious algorithm is an algorithm formulated for the RAM model but analyzed with the I/O model, with the requirement that the analysis holds for any (M, B) pair under the only assumption of the *tall memory*: $M = \Omega(B^2)$

The I/O-model analysis is valid for any block B and memory M , and then is valid for *all levels* of the cache hierarchy. Hence, optimizing an algorithm, or a data structure for an unknown level of the hierarchy, it well behaves in the whole hierarchy.

Cache – oblivious model

The cache-oblivious algorithms outperform the old-fashioned RAM algorithms and quite often are close or equivalent to algorithms tuned to a specific memory hierarchy.

However, they well behave at all level of the hierarchy and are much more portable (with good performance) and much more robust against change in problem sizes than algorithms designed for a very specific architecture.

Cache – oblivious model

Several cache-oblivious algorithms and data structures have been found:

- Sorting
- Searching
- Static / dynamic B+-trees
- FFT
- Matrix transpose
- Matrix-Matrix multiply
- Computational geometry
- Priority queues
- Graph algorithms
- < ... >

We are not going into detail here, just listing few basic concepts

Some basic notes on data structures

Field reordering

Reorder fields in structures so that what is used together stays together

```
struct my_node {  
    double   key;  
    char     my_data[300];  
    my_node *next_node;  
}
```

Linked-list node

```
void myfunc(my_node *p, double key, <...>) {  
    while( p != NULL) {  
        if( p->key == key ) {  
            do_something( <...> );  
            break;  
        }  
        p = p → next_node;  
    }  
}
```

Linked-list traversal

Field reordering

Reorder fields in structures so that what is used together stays together

```
struct my_node
{
    double    key;
    char      my_data[300];
    my_node *next_node;
}
```



```
struct my_node
{
    double    key;
    my_node *next_node;
    char      my_data[300];
}
```

Field split: hot and cold

Split fields so that to keep consecutive the fields that are used sequentially

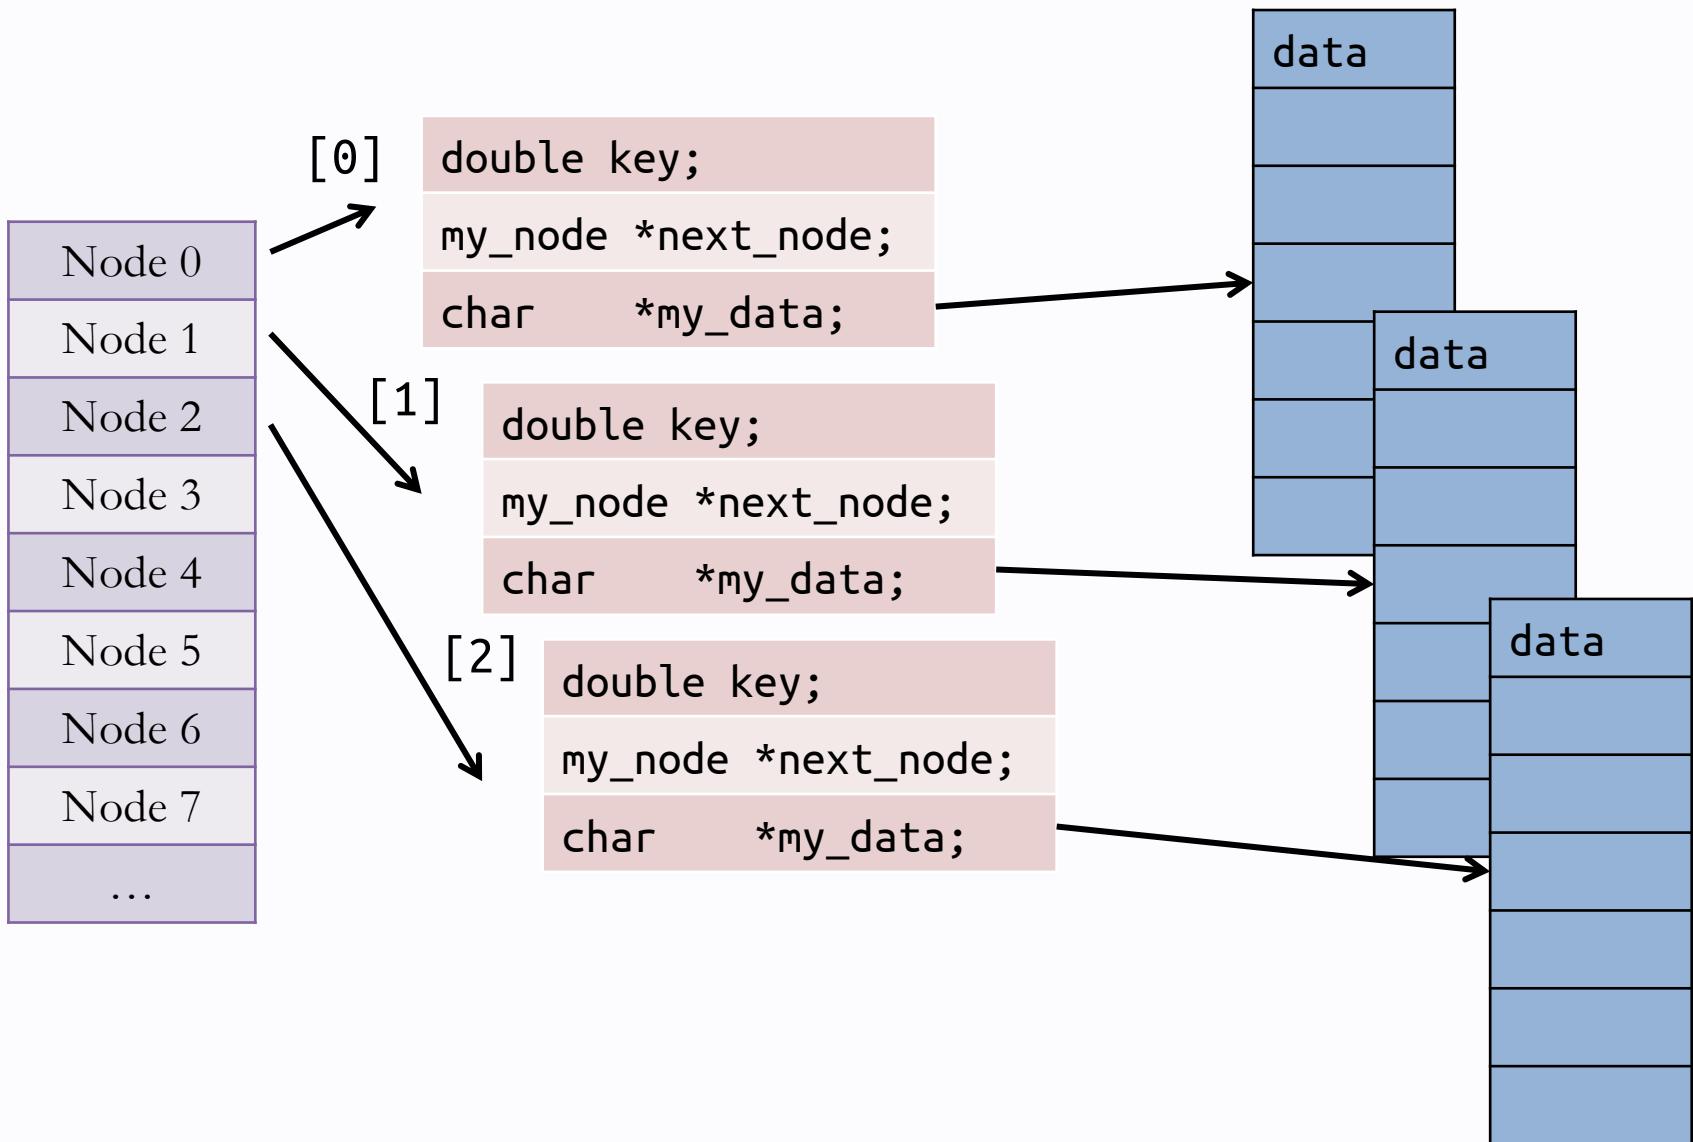
```
struct my_node
{
    double    key;
    my_node *next_node;
    char      my_data[300];
}
```



```
struct my_node
{
    double    key;
    my_node *next_node;
    void     *my_data;
}
```

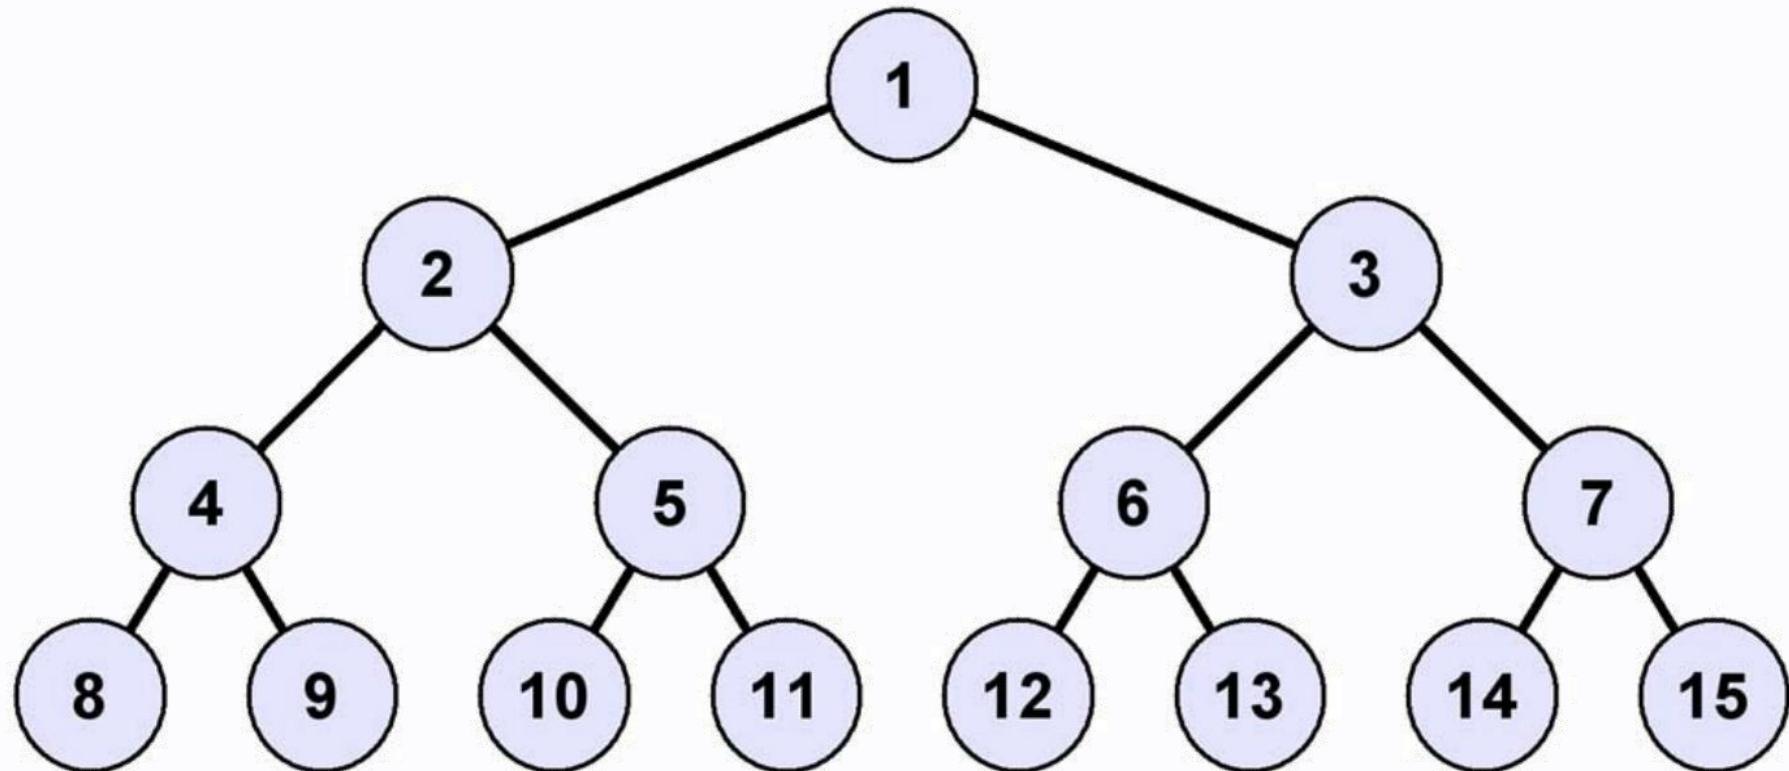
```
struct my_data
{
    char data[300];
}
```

Field split: hot and cold



Data structure reordering

Example: TREEs – depth-first order

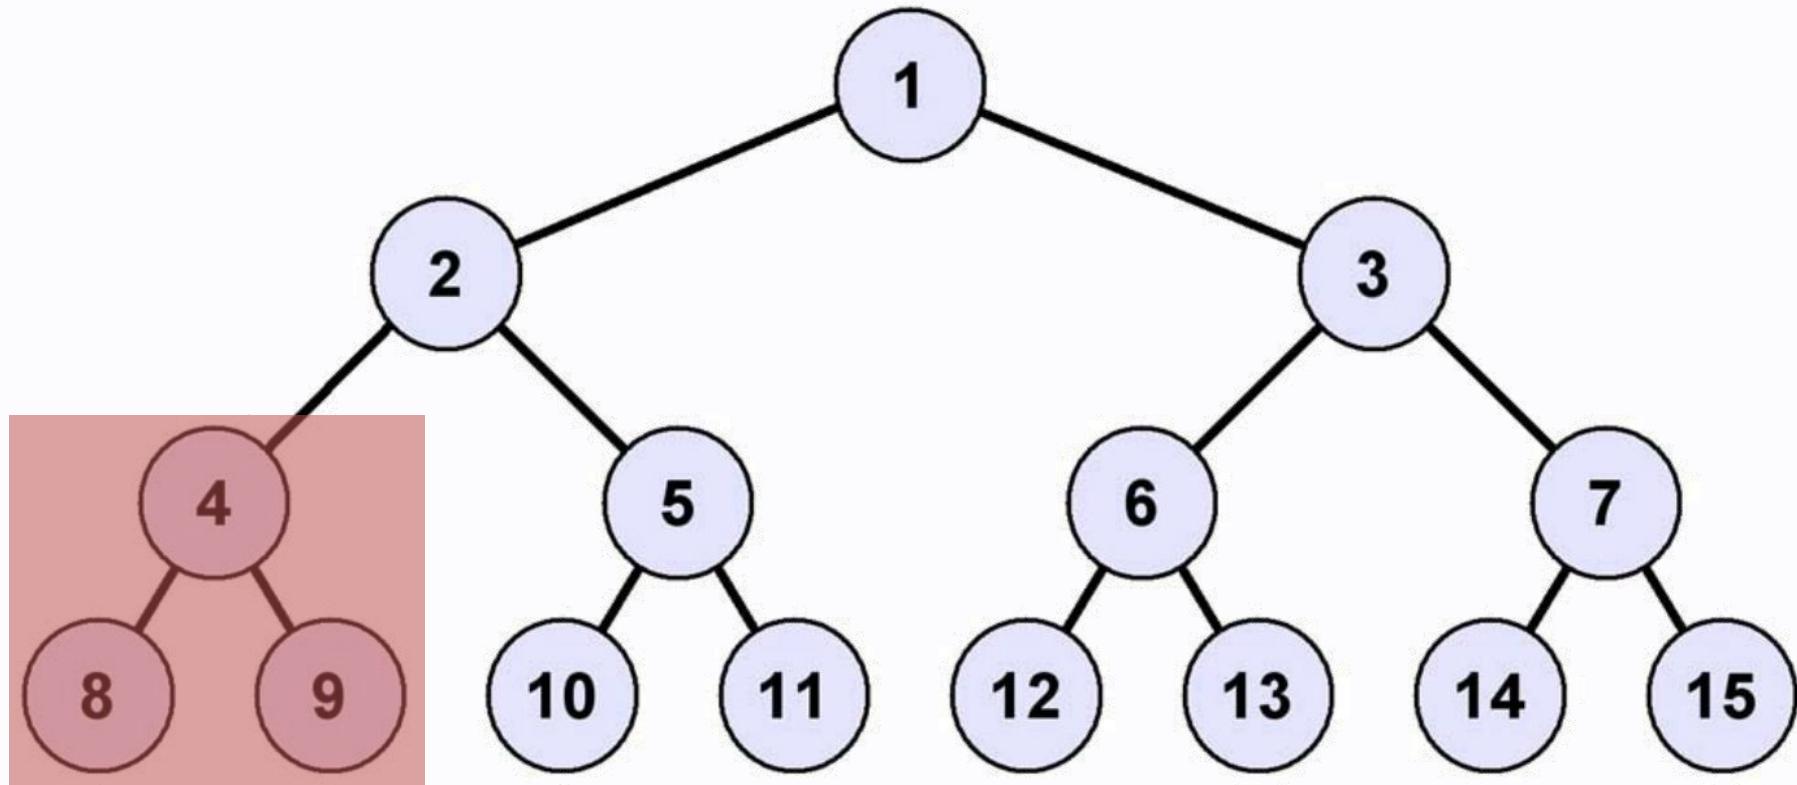


1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

$$\text{Left}[n] = 2n - \text{right}[n] = 2n+1$$

Data structure reordering

Example: TREEs – depth-first order

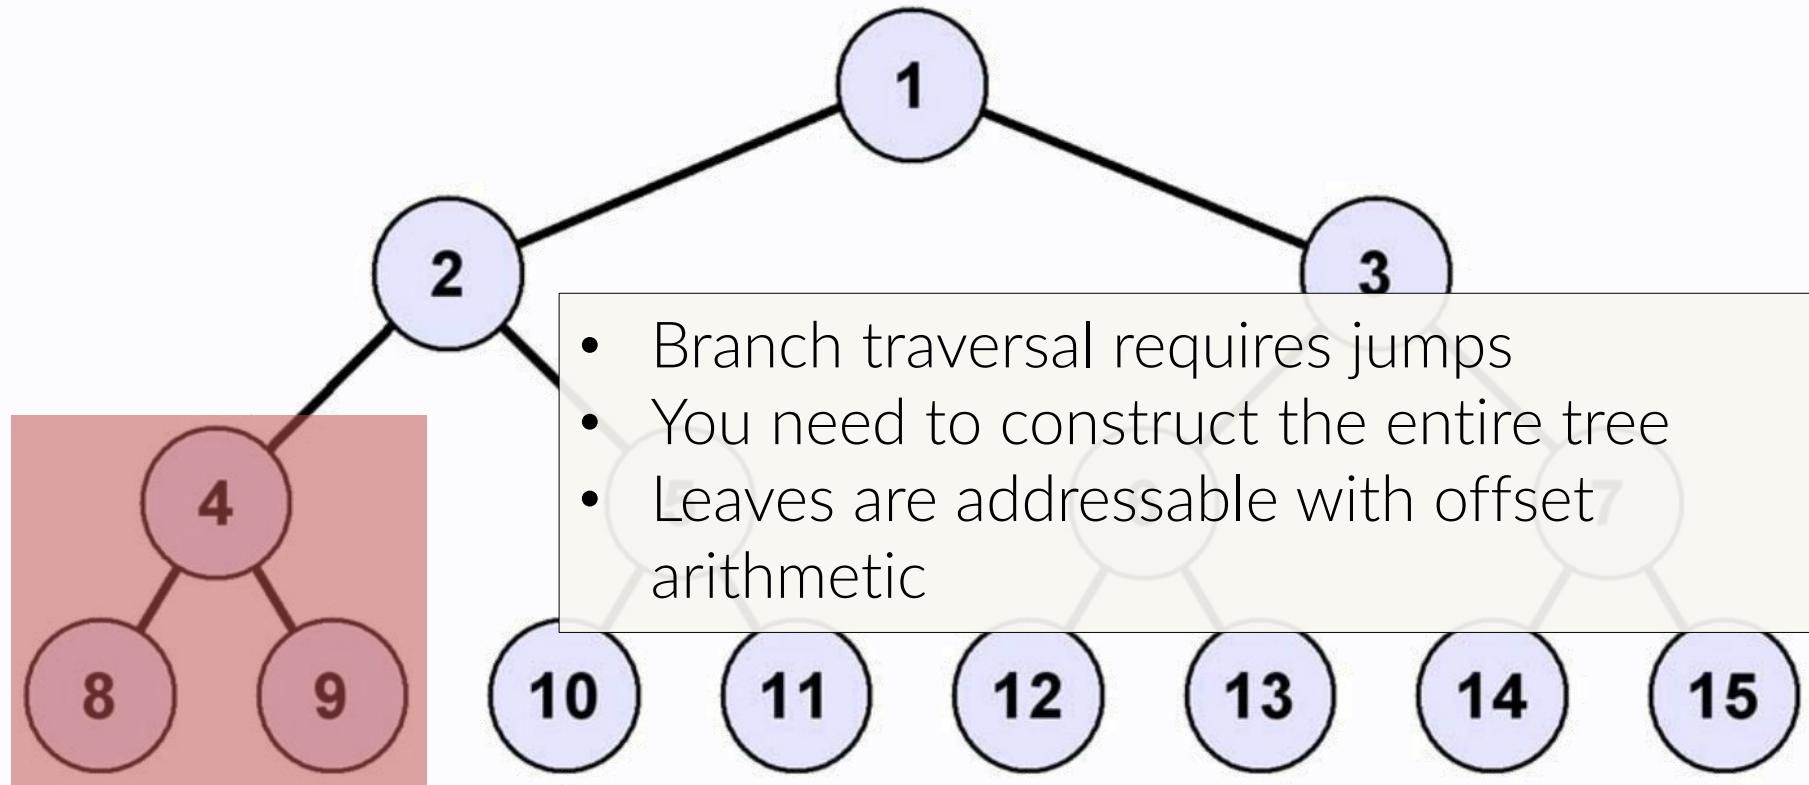


1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

$$\text{Left}[n] = 2n - \text{right}[n] = 2n+1$$

Data structure reordering

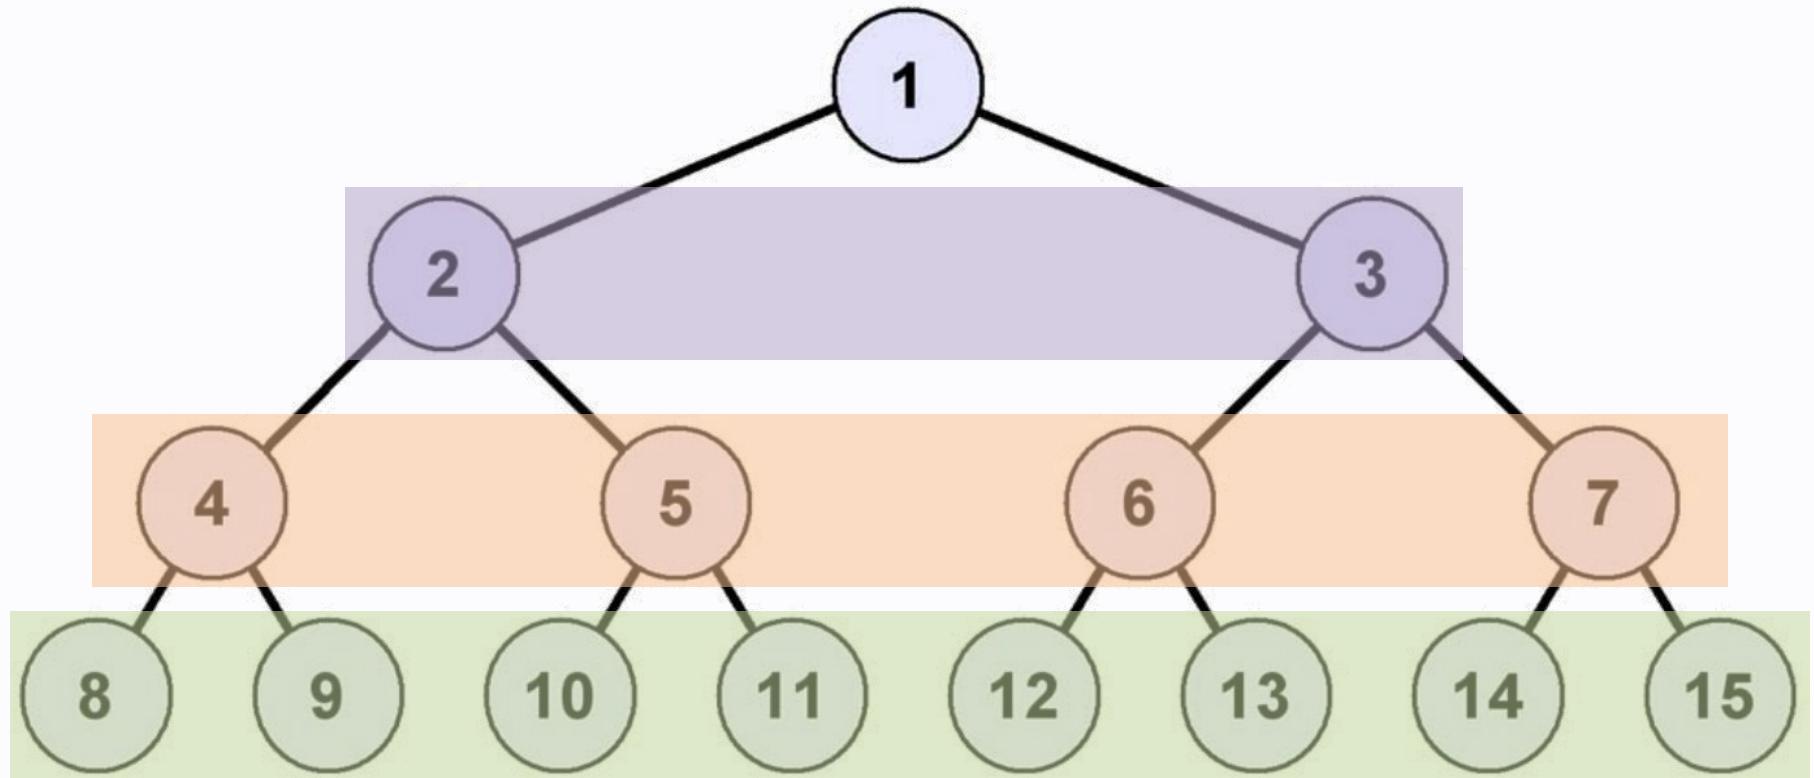
Example: TREEs – depth-first order



$$\text{Left}[n] = 2n - \text{right}[n] = 2n+1$$

Data structure reordering

Example: TREEs – depth-first order

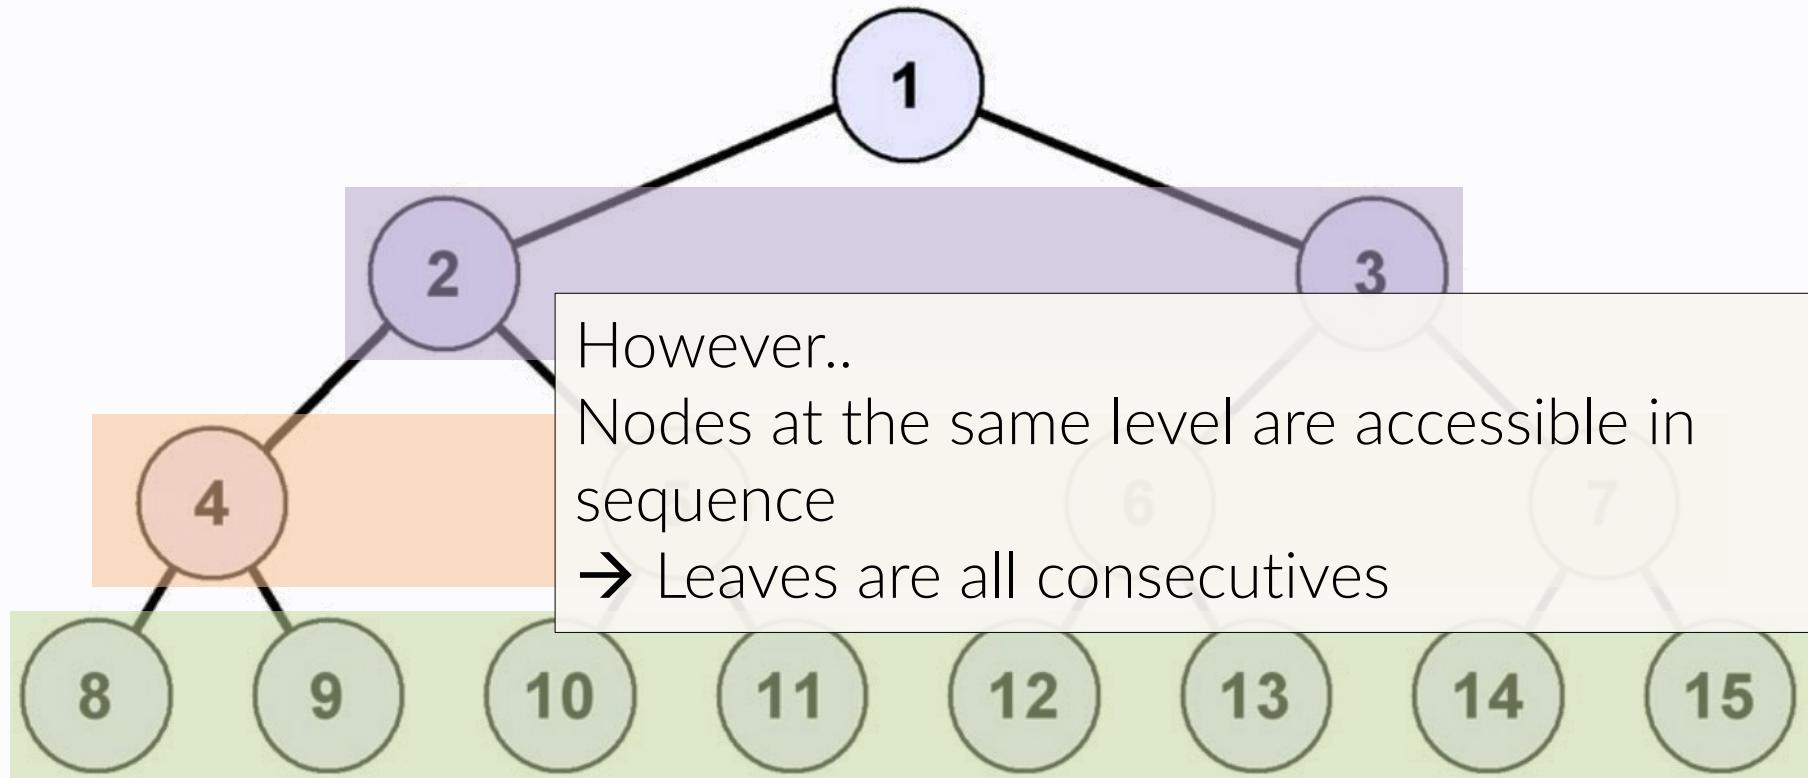


1	2	3	4	5	6	7	8	9	10	11	12	13	14	15
---	---	---	---	---	---	---	---	---	----	----	----	----	----	----

$$\text{Left}[n] = 2n - \text{right}[n] = 2n+1$$

Data structure reordering

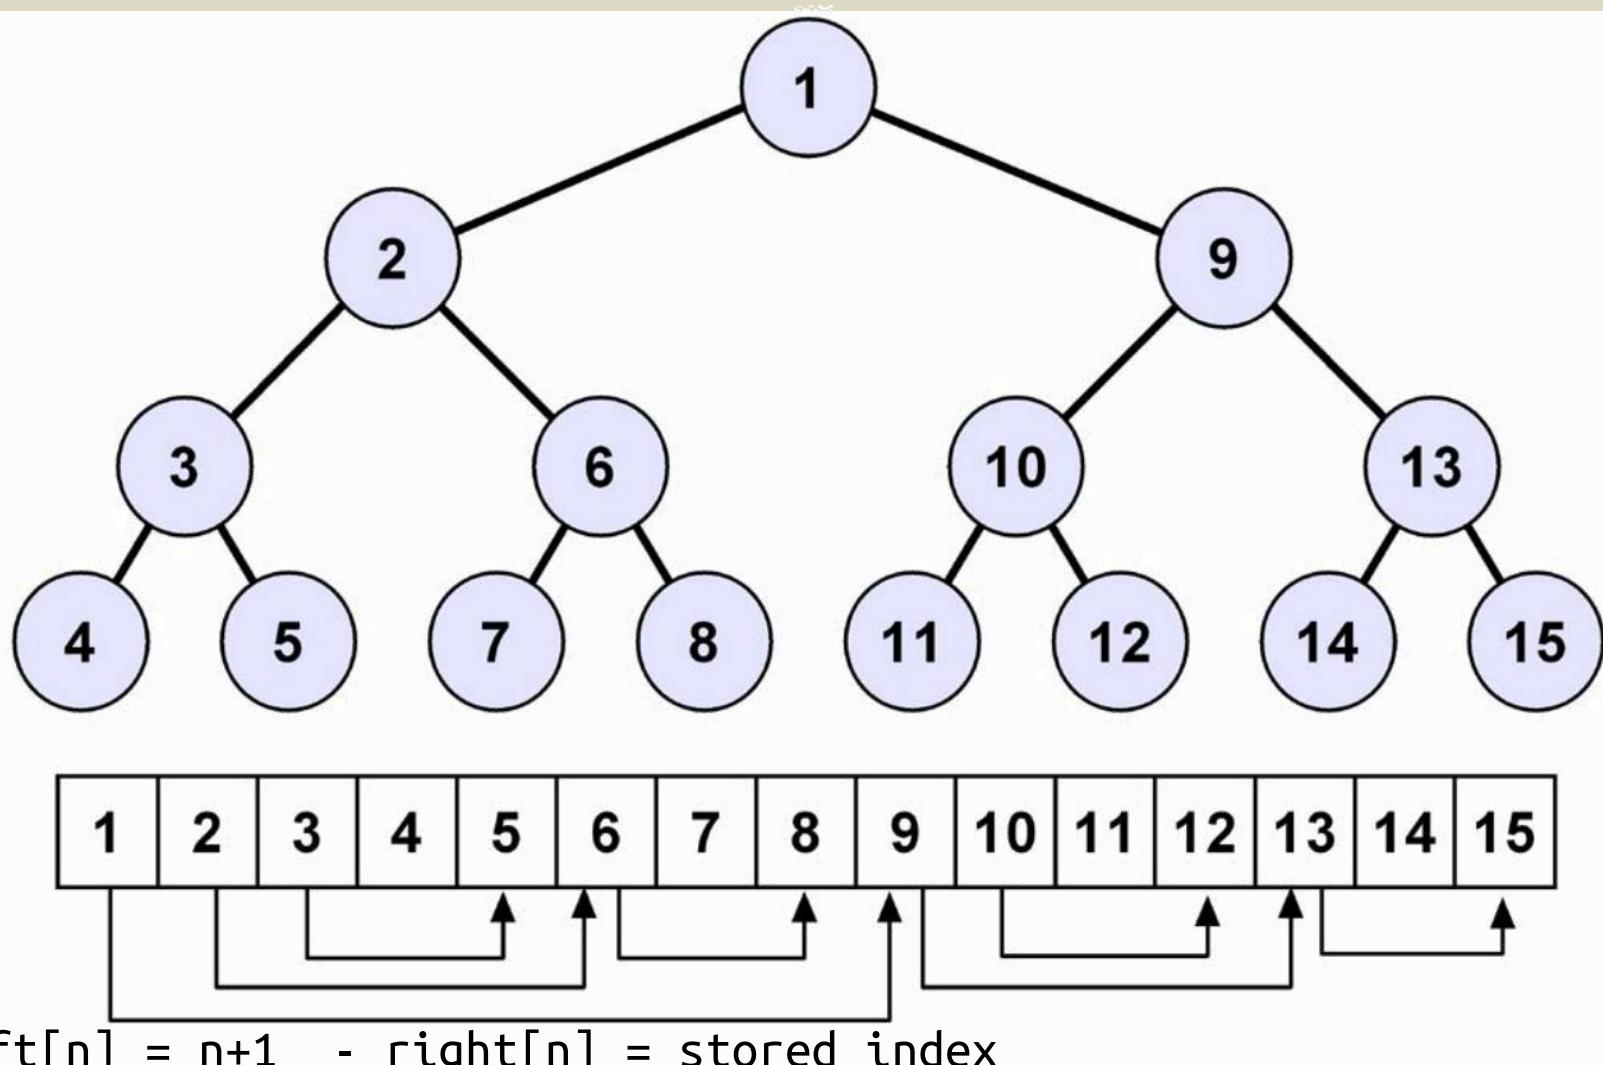
Example: TREEs – depth-first order



$$\text{Left}[n] = 2n - \text{right}[n] = 2n+1$$

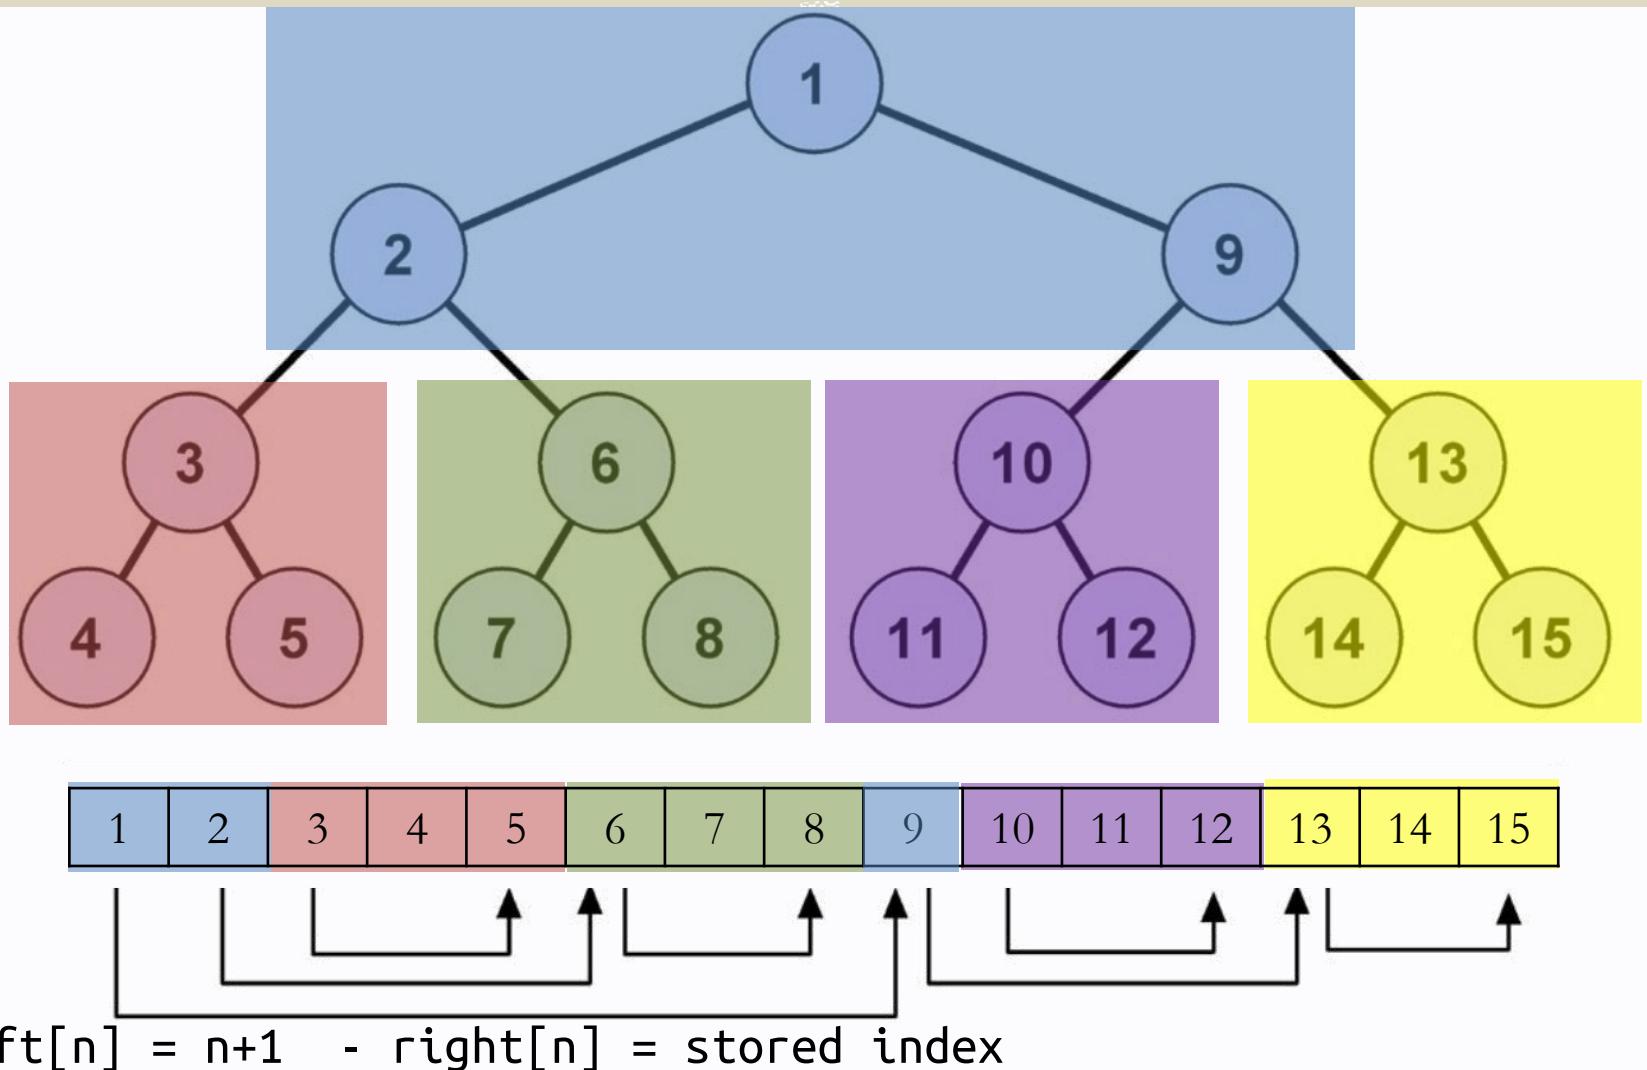
Data structure reordering

Example: TREEs – breadth-first order



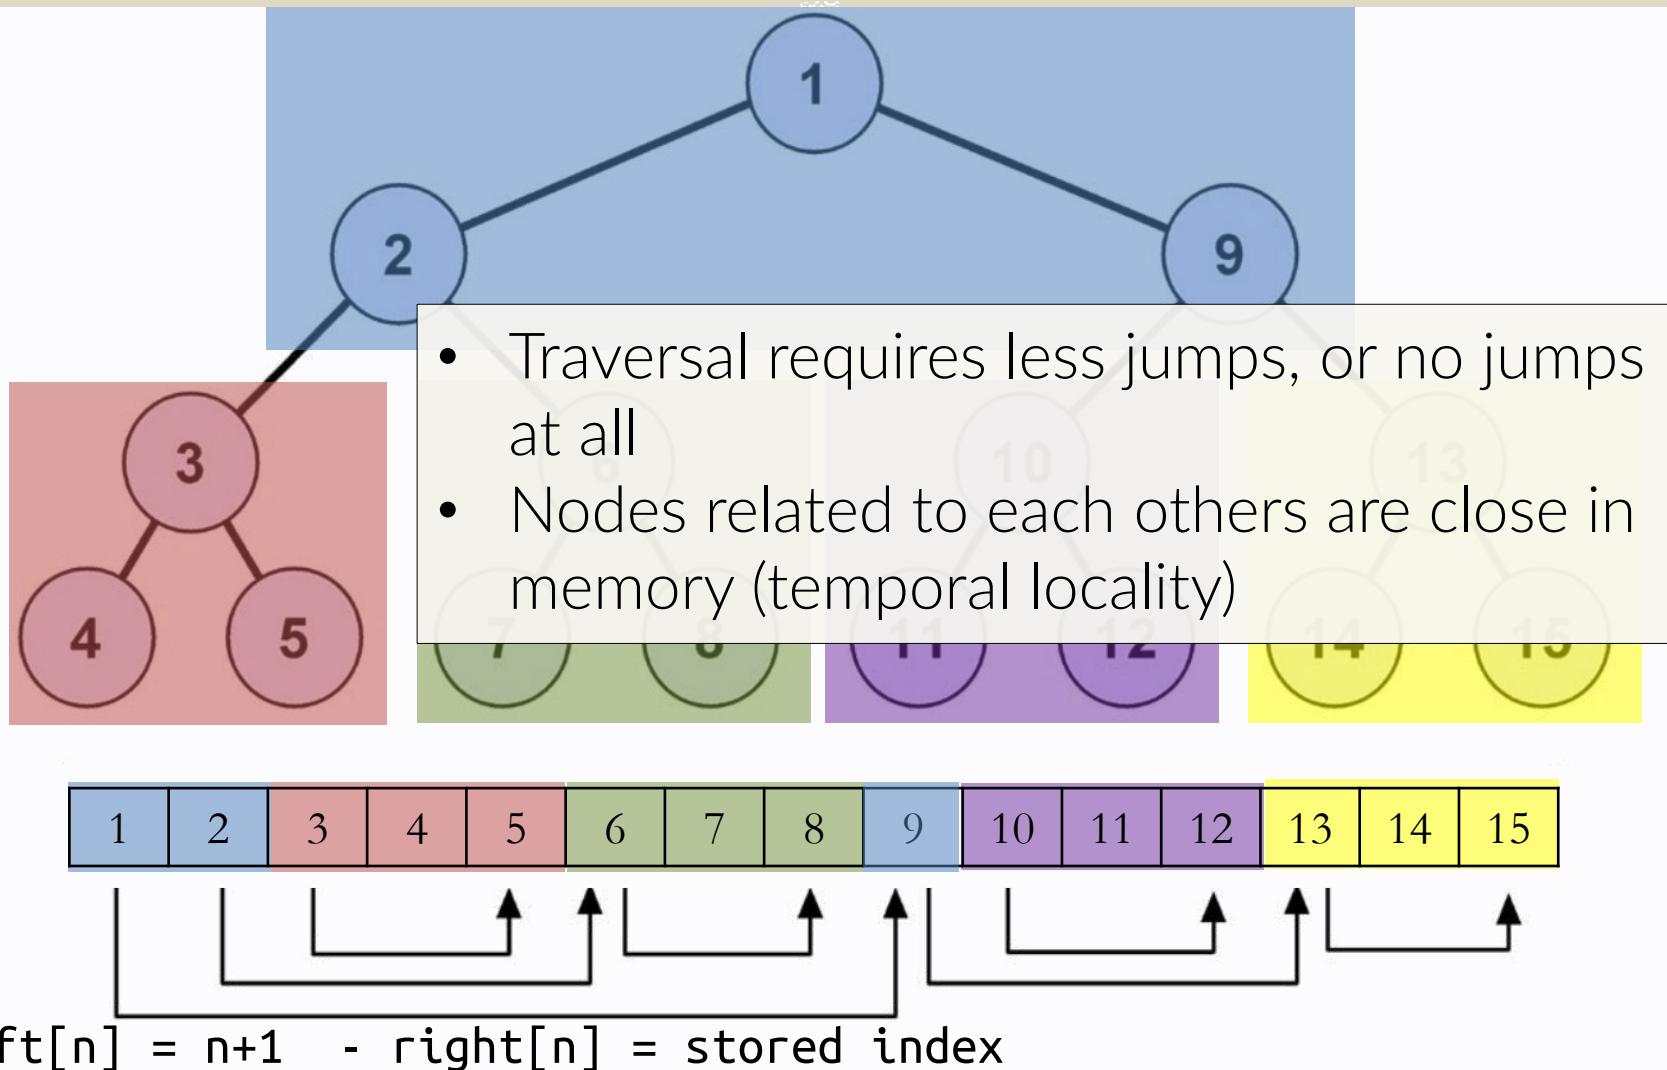
Data structure reordering

Example: TREEs – breadth-first order



Data structure reordering

Example: TREEs – breadth-first order



Organizing data for locality

Data reordering

As we have seen, there are several approaches to the cache-optimization problem:

- Cache -*aware* or -*oblivious* algorithm design
- Embedded heuristics and policies in the CPU
- Prefetching
 - hardware
 - software (explicit or “induced”)
- SO level (page mapping)
- Fields reordering in data structures

and ...

→ **DATA REORDERING**

Data reordering

When the memory bandwidth is “limited” – which may be the case for highly parallel + multicore systems with a strong NUMA hierarchy, **data locality optimization** can play a strong role.

Re-organizing data in “space” (whichever is their n -dimensional space) so that the access pattern is optimal for a given algorithm is related to such locality optimization.

Ex. 1: data pattern might be trivial, as in matrix transpose/mul

→ very specific ordering or pattern design

Ex. 2: data pattern may be spatially-coherent but unknown before it happens. For instance, in radiative transfer

→ optimization of data needed for a general case

Data reordering

Optimizing data locality for a general case

In the “space” the data live in – for instance our usual 3D space – there might be a metric that correlates with spatial coherence (in our 3D space there is the Euclidean metric, for instance).

Generally, it is more probable to access in a short time lapse points that are also spatially close.

Then, two popular measure to exploit this are

- Minimizing the distance distortion
- Preserve the locality

i.e. keeping close in 1D memory world points that are close in n -dimensions enhances the probability of using neighbouring memory locations while they are still in the cache.

Data reordering

What is the “minimum distance distortion” ?

In 1-D the answer is trivial.

In 2-D is less trivial, and it is increasingly less trivial while the number of dimensions rises.

Data reordering

Scanline
order
row-major

(a)	0 1 2 3 4 5 6 7
	8 9 10 11 12 13 14 15
	16 17 18 19 20 21 22 23
	24 25 26 27 28 29 30 31
	32 33 34 35 36 37 38 39
	40 41 42 43 44 45 46 47
	48 49 50 51 52 53 54 55
	56 57 58 59 60 61 62 63

Scanline
order
col-major

(b)	0 8 16 24 32 40 48 56
	1 9 17 25 33 41 49 57
	2 10 18 26 34 42 50 58
	3 11 19 27 35 43 51 59
	4 12 20 28 36 44 52 60
	5 13 21 29 37 45 53 61
	6 14 22 30 38 46 54 62
	7 15 23 31 39 47 55 63

Block order
+ scan sub-
order

(c)	0 1 2 3 16 17 18 19
	4 5 6 7 20 21 22 23
	8 9 10 11 24 25 26 27
	12 13 14 15 28 29 30 31
	32 33 34 35 48 49 50 51
	36 37 38 39 52 53 54 55
	40 41 42 43 56 57 58 59
	44 45 46 47 60 61 62 63

Z- order
or
Bit-
interleaved
or
Morton order

(d)	0 1 4 5 16 17 20 21
	2 3 6 7 18 19 22 23
	8 9 12 13 24 25 28 29
	10 11 14 15 26 27 30 31
	32 33 36 37 48 49 52 53
	34 35 38 39 50 51 54 55
	40 41 44 45 56 57 60 61
	42 43 46 47 58 59 62 63

Data reordering

Let's go in deeper detail about that “Z-order”

Let's say you want to map a linear access order

0, 1, 2, 3, 4, 5, ..., nth

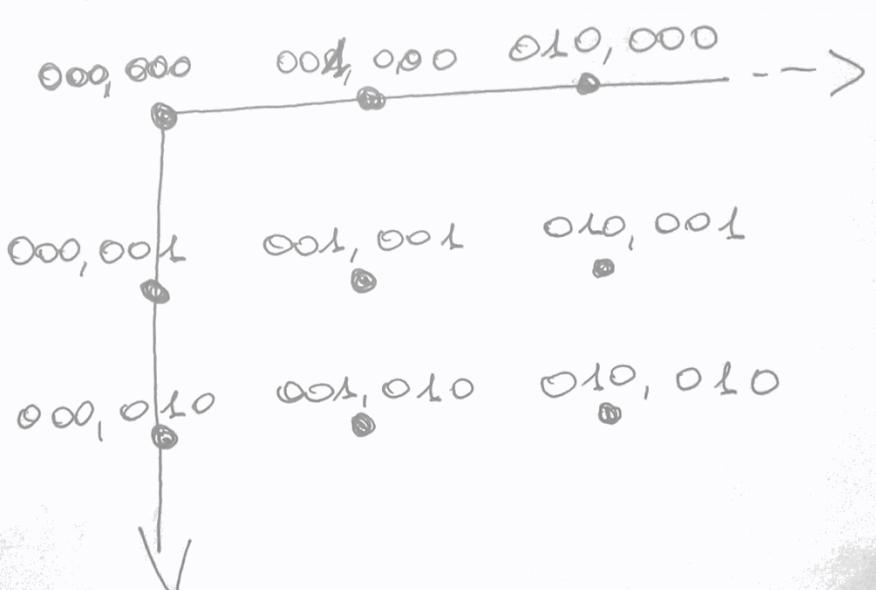
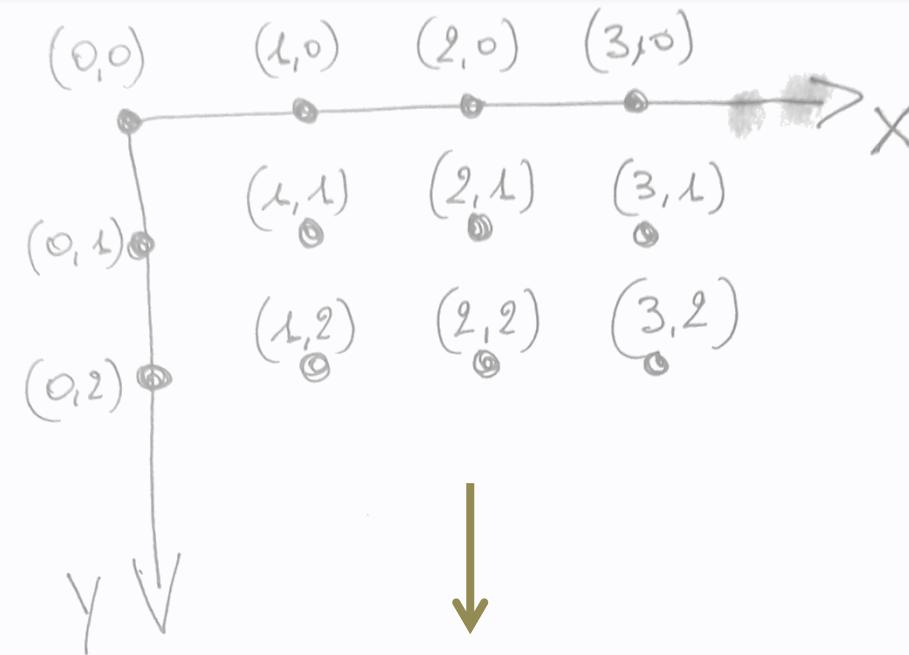
on some spatially-distributed data with integer coordinates.

Now, let's rewrite the linear access in base 2:

0000, 0001, 0010, 0011, 0100, 0101, ...

What happens if the *bits* of the indexes of our traversal order are taken from the *bits* of the spatial coordinates of our points with a peculiar reshuffling?

Data reordering



Let's define the binary representation of an integer number x :

$$x = x_i \dots x_2 x_1 x_0$$

so that a couple (x, y) reads as:

$$(x_i \dots x_2 x_1 x_0, y_i \dots y_2 y_1 y_0)$$

Then let's define the following reshuffle so to *interleave the bits*

$$(x, y) \rightarrow (y_i x_i \dots y_2 x_2 y_1 x_1 y_0 x_0)$$

$$(0,0) \rightarrow 00\ 00 = 0$$

$$(1,0) \rightarrow 00\ 01 = 1$$

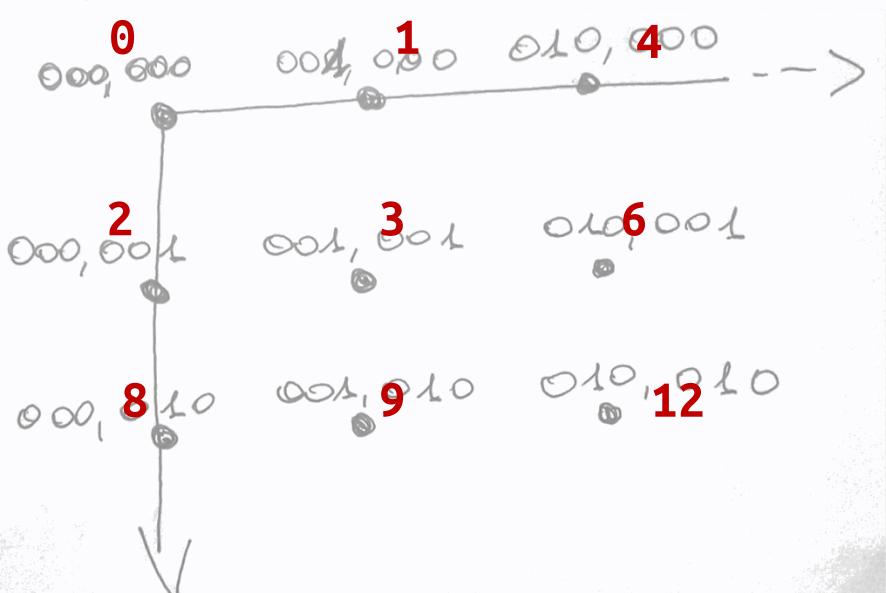
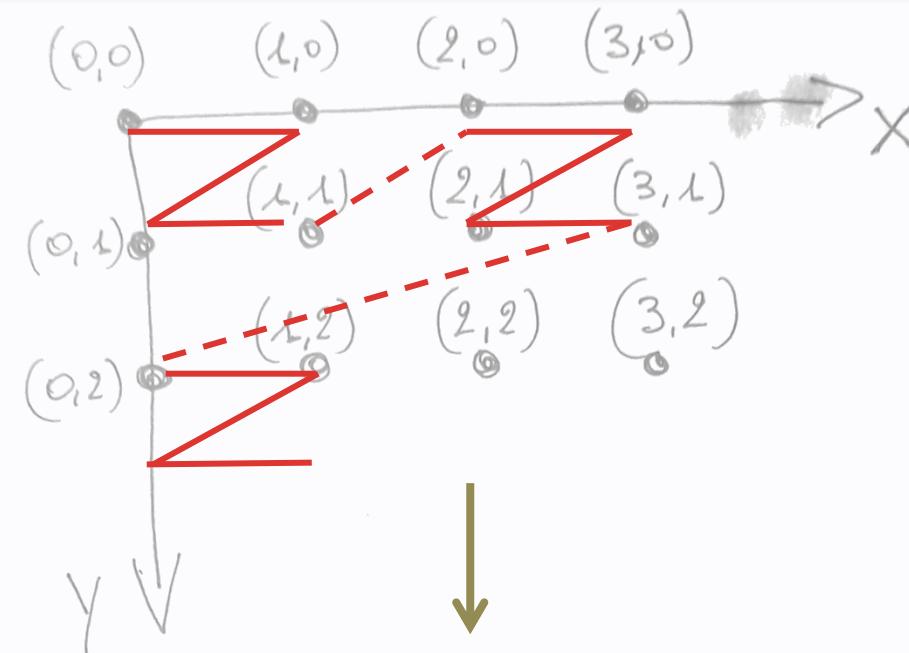
$$(2,0) \rightarrow 01\ 00 = 4$$

$$(0,1) \rightarrow 00\ 10 = 2$$

$$(1,1) \rightarrow 00\ 11 = 3$$

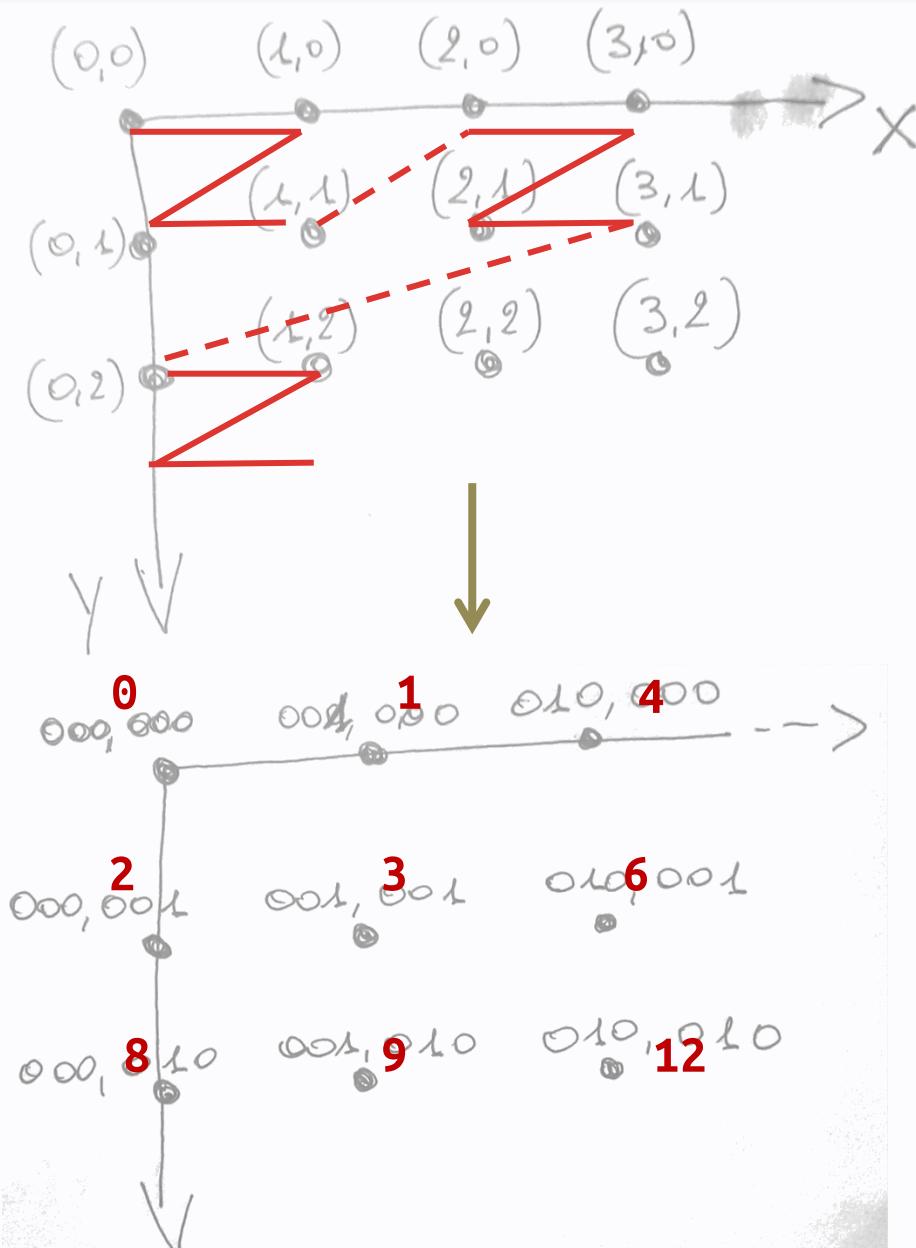
...

Data reordering

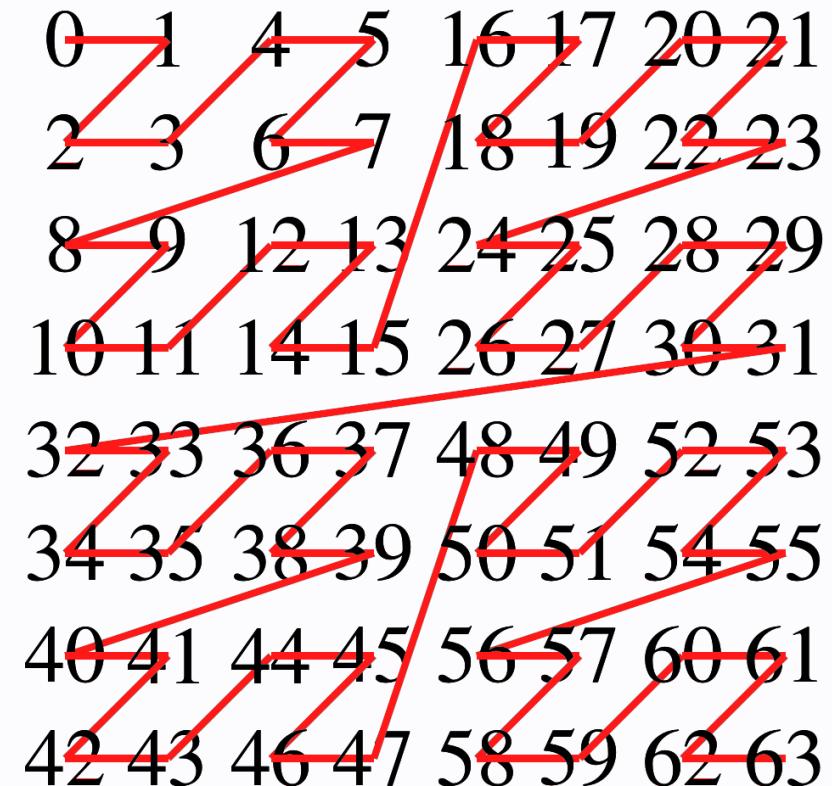


The result of this interleaving is a mapping from the 2-D plane to the 1-D line known as Z-line which is one of the **plane-filling curves** discovered by Peano

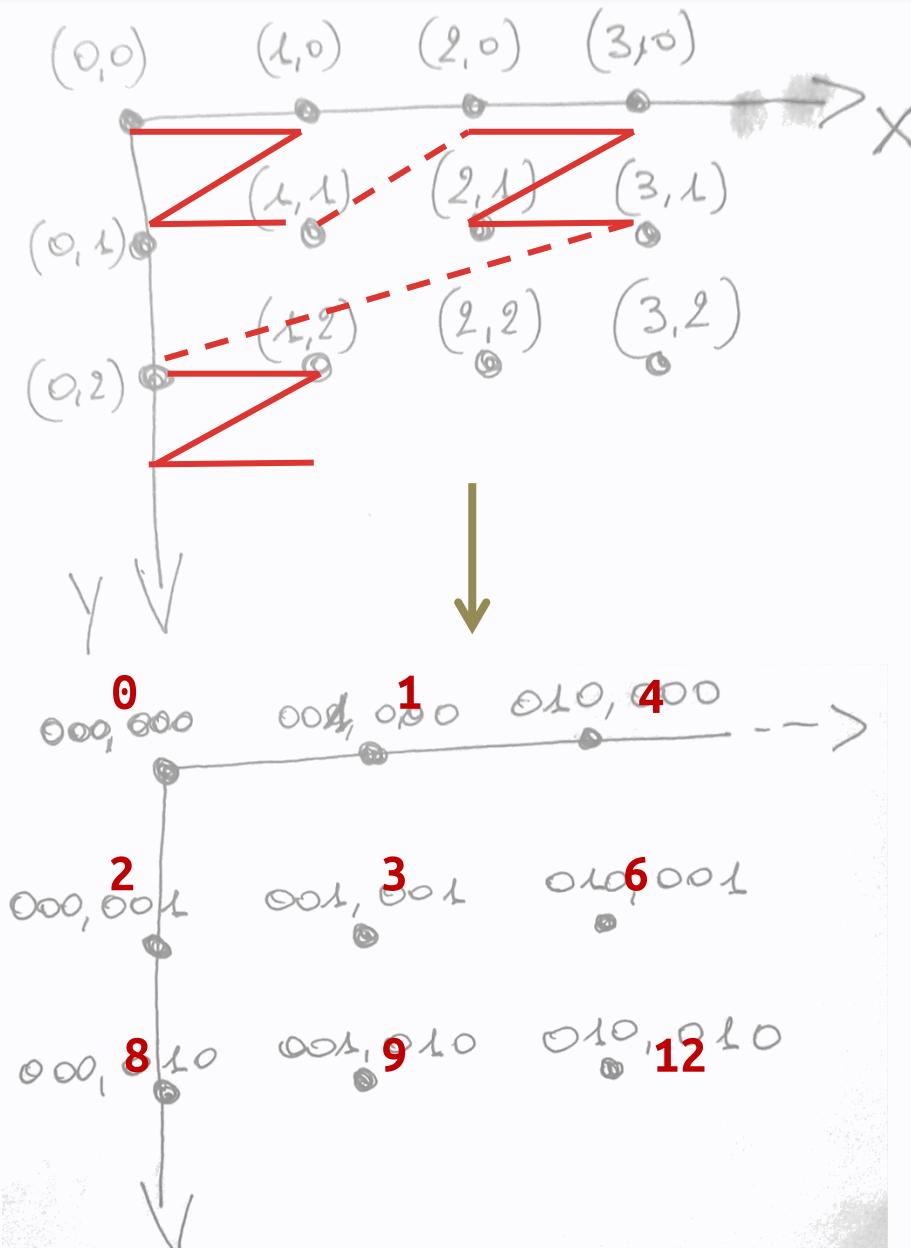
Data reordering



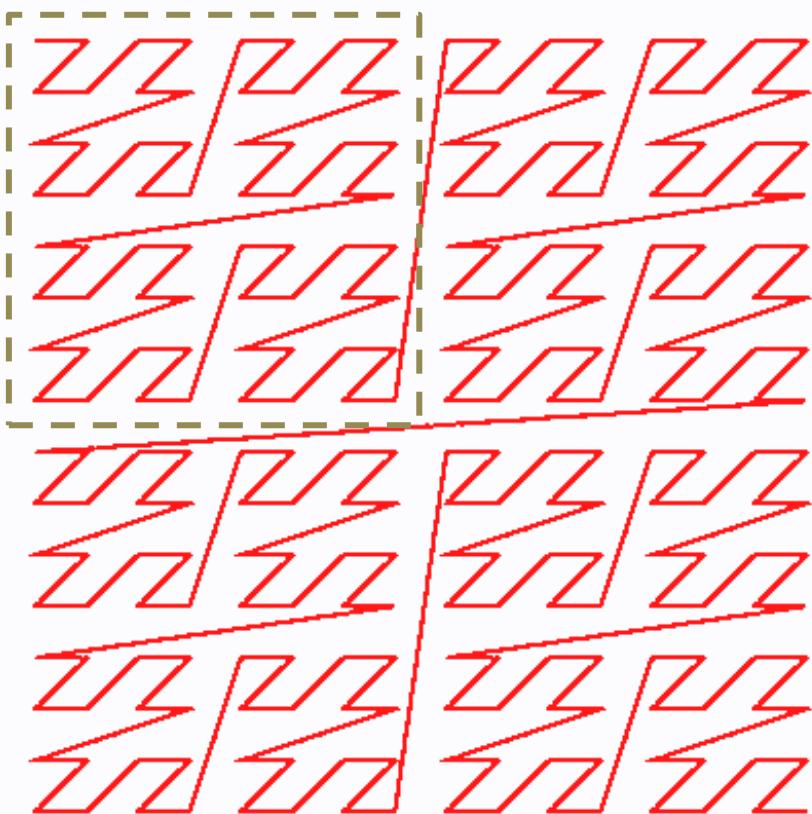
The result of this interleaving is a mapping from the 2-D plane to the 1-D line known as Z-line which is one of the **plane-filling curves** discovered by Peano



Data reordering



The result of this interleaving is a mapping from the 2-D plane to the 1-D line known as Z-line which is one of the **plane-filling curves** discovered by Peano

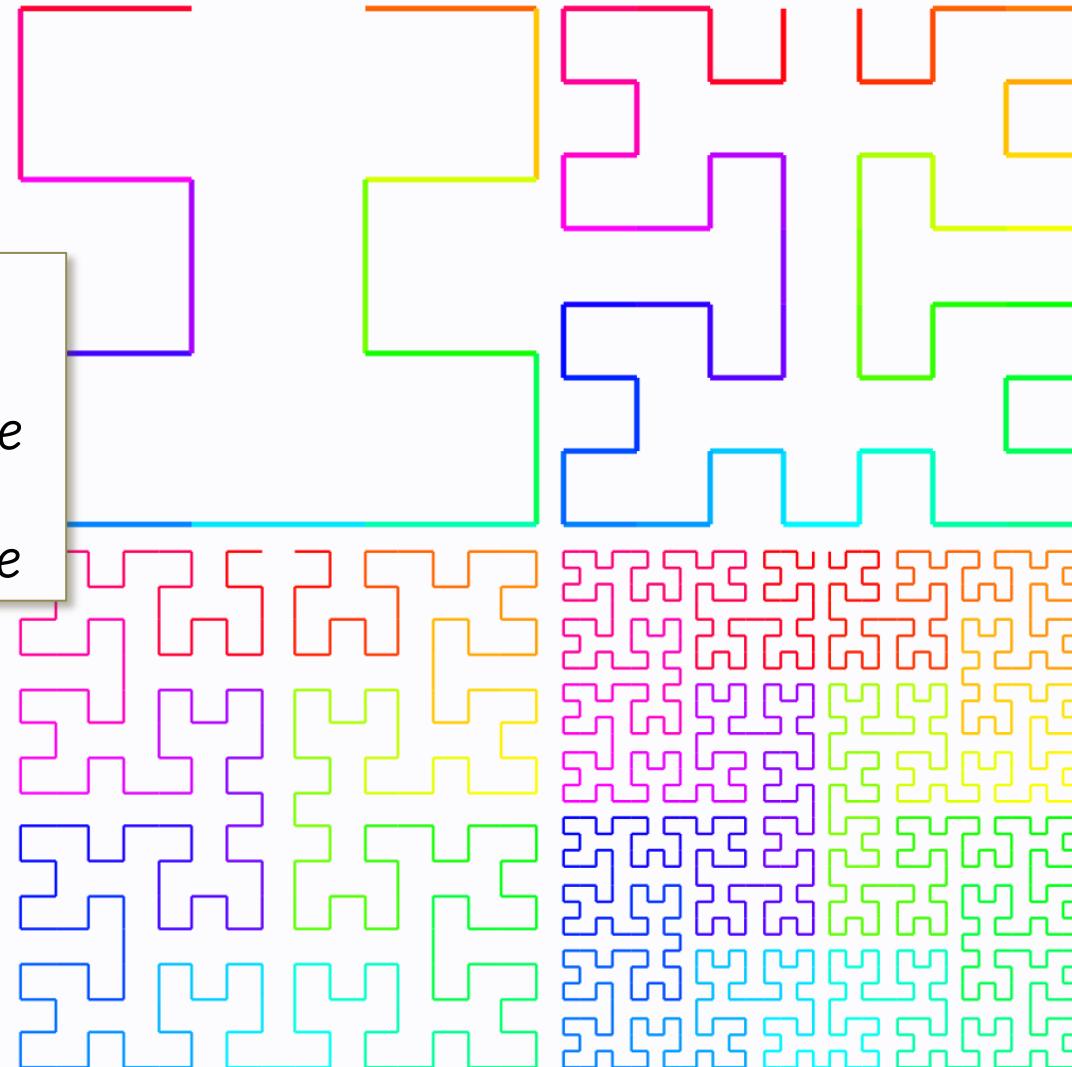


The wonderful land of *Peano*-curves

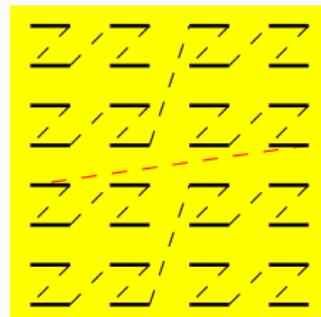
There are many of these curves

Hilbert curve

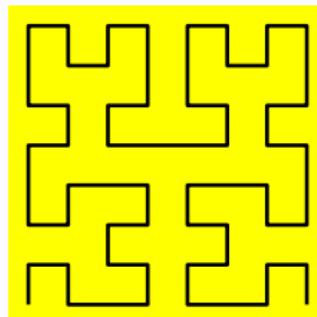
Note that there
are no jumps
along the curve



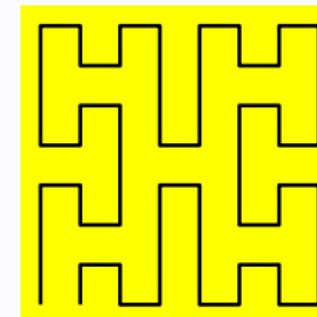
The wonderful land of *Peano*-curves



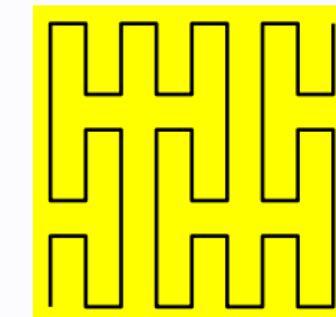
Z-order



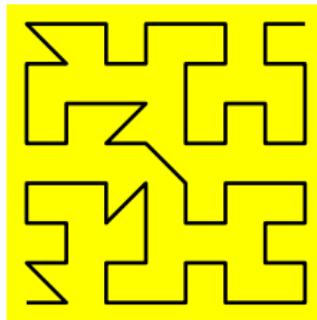
Hilbert curve



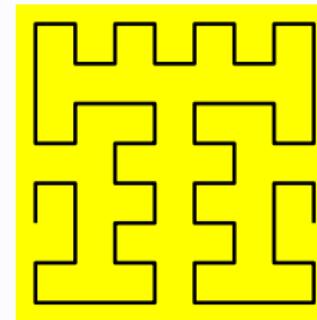
H-order



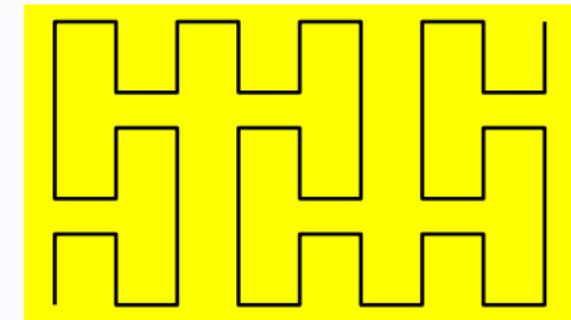
Peano's curve



AR²W²-curve



$\beta\Omega$ -curve

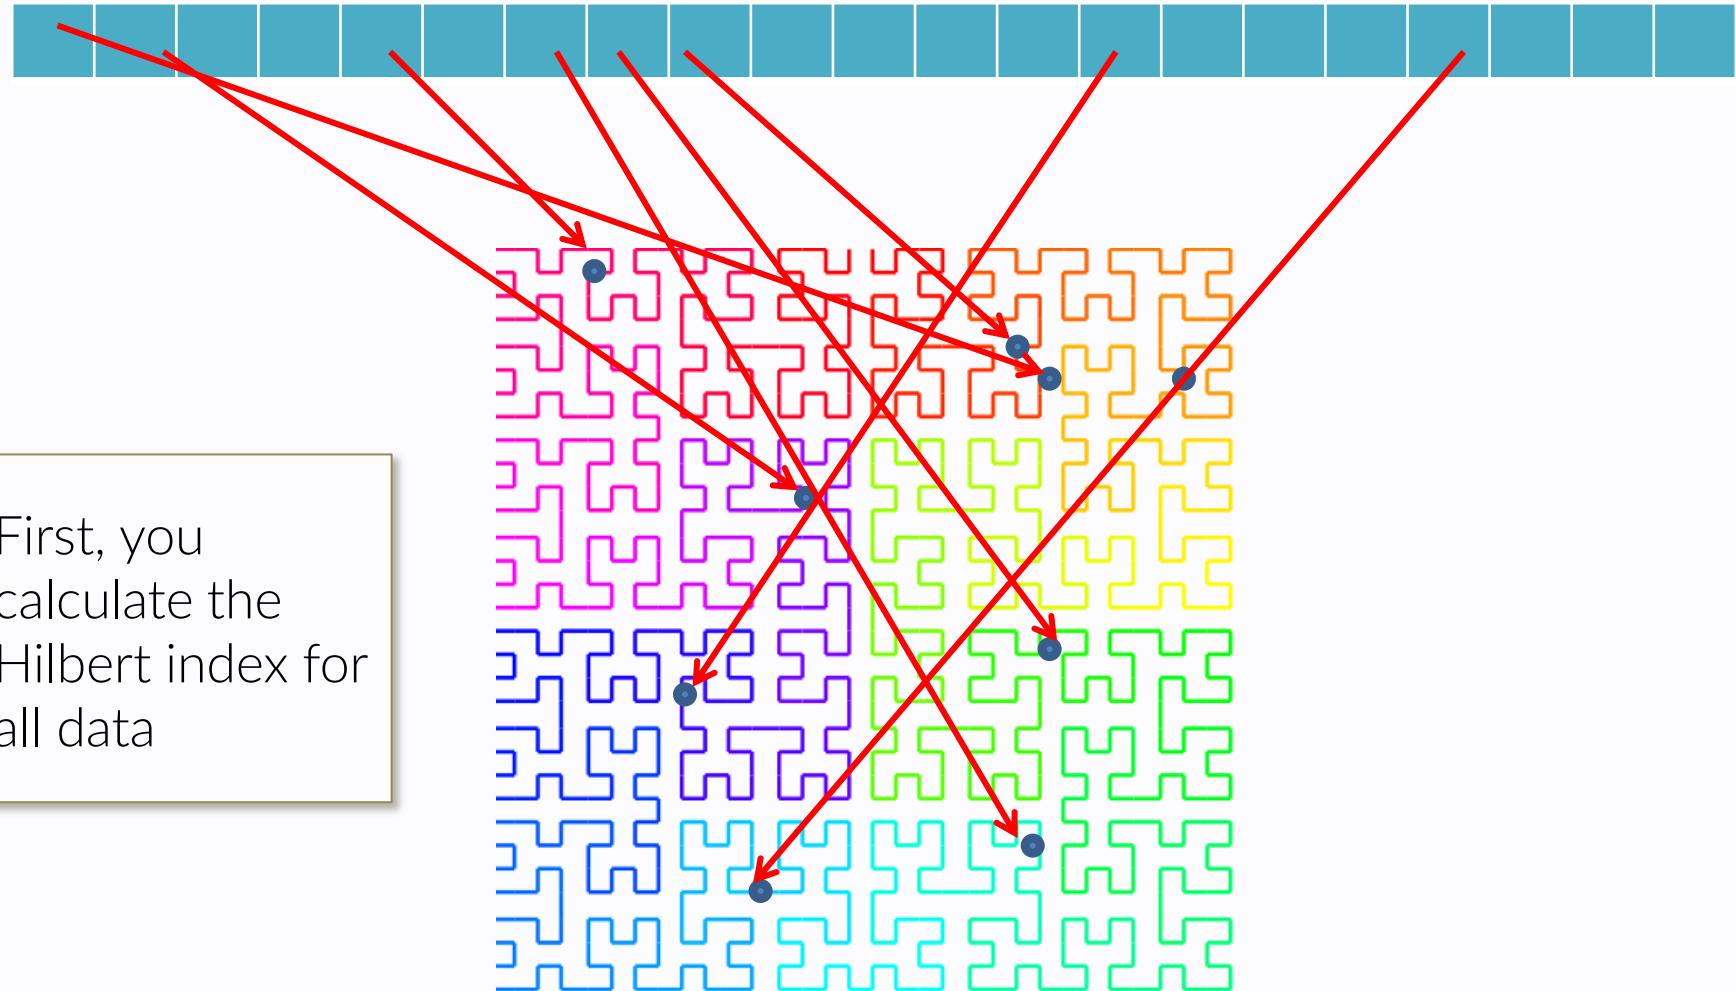


Balanced Peano

Each curve has some peculiar properties which reflects in the distance distortion they provide, i.e. on their performance in keeping “locality” in different situations

The wonderful land of Peano-curves

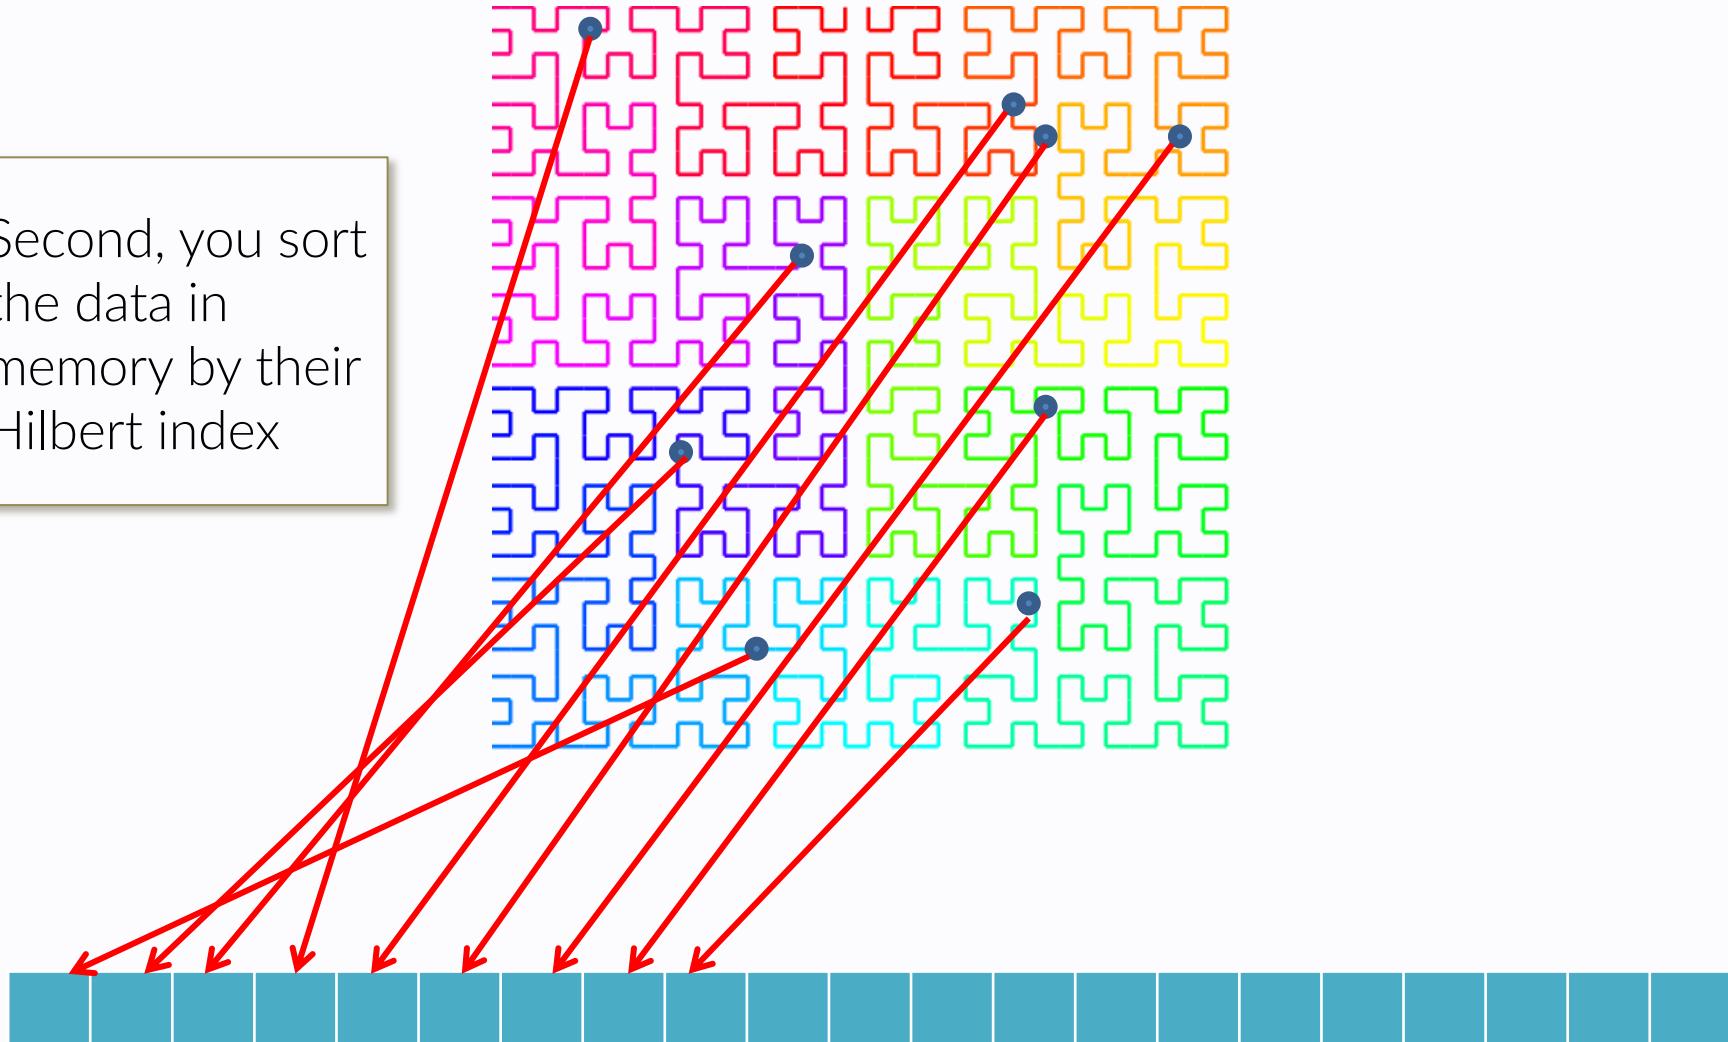
Re-ordering data in memory



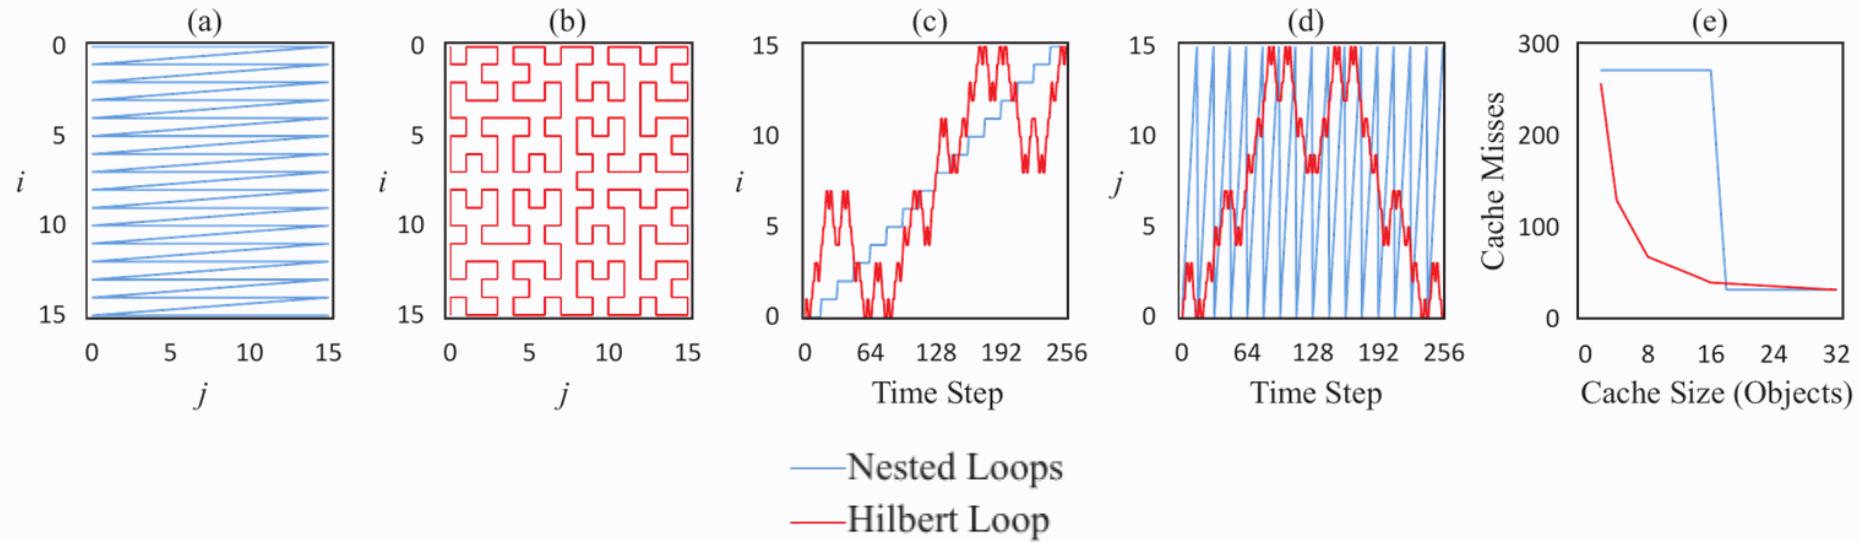
The wonderful land of Peano-curves

Re-ordering data in memory

Second, you sort the data in memory by their Hilbert index



The wonderful land of *Peano*-curves

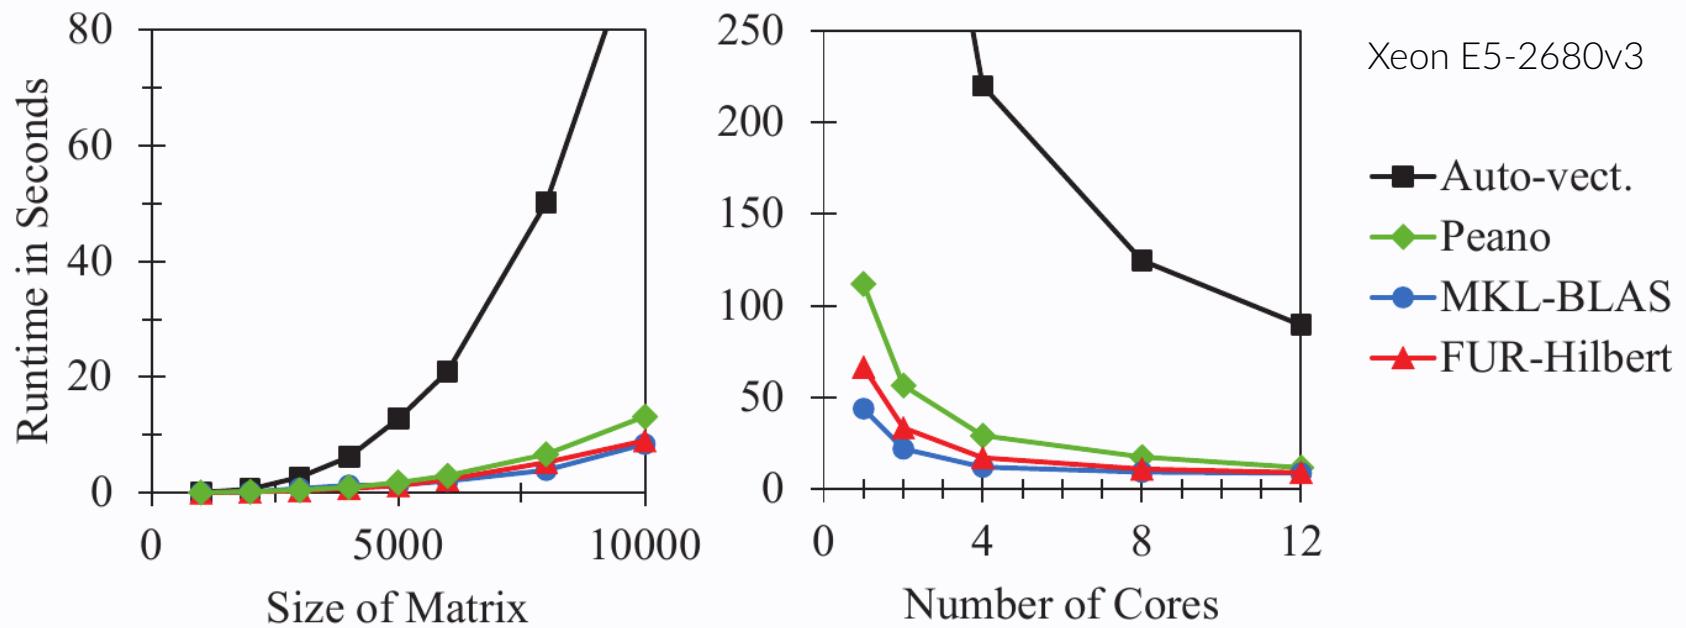


Cache-oblivious at work : traversal order

Comparison of the traversal order for classic nested-loops and Hilbert curve.

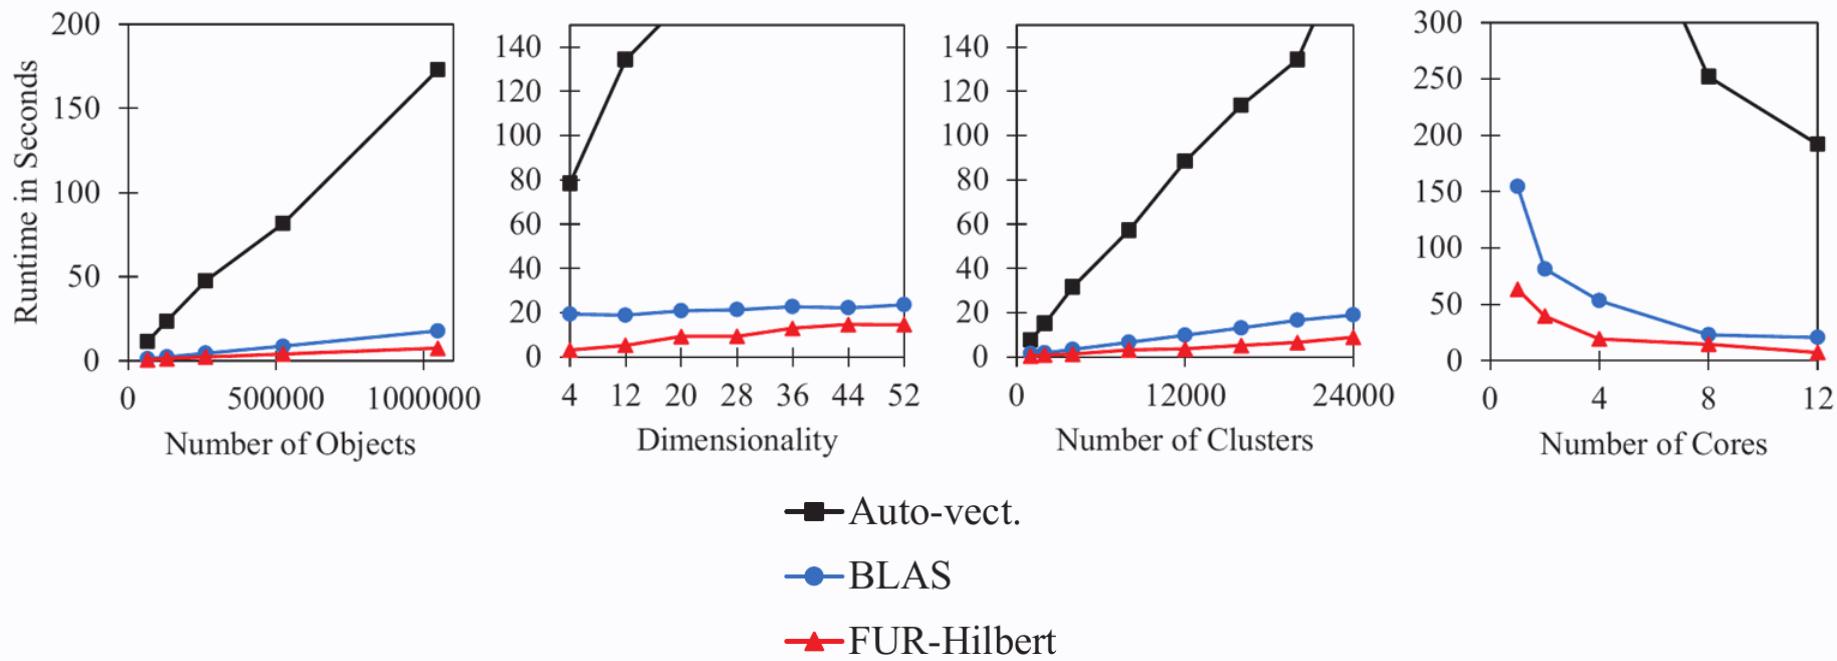
Note the increased locality in j and the highly reduced numbers of L1 miss

The wonderful land of *Peano*-curves



Cache-oblivious at work : matrix multiplication

The wonderful land of *Peano*-curves



Cache-oblivious at work : *k*-means

Cache recap
In two slides

RECAP – Foes: 3 **C**'s of cache misses

▶ Compulsory misses

Unavoidable misses when data are read for the first time

▶ Capacity misses

- Not enough space to hold all data
- Too much data accessed in between successive use

▶ Conflict misses

Cache trashing due to data mapping to same cache lines

RECAP – Friends: 3 **R**'s of cache hits

► Rearrange (code & data)

Design layout to improve temporal & spatial locality

► Reduce (size)

- Smaller data size – smaller chunks accessed
- Fewer instructions

► Reuse (cache lines)

Increase spatial & temporal locality – keep resident data for more operations

Outline

- Memory optimizations
 - > the importance of memory access
 - > the importance of the cache
 - > cache access optimization in loops
- Loops optimizations
- Other optimizations

Recap: a friend we've already seen
Avoid the avoidable..

Loops optimization

Is there something wrong in this simple nested loop ?

```
for(p = 0; p < Np; p++)
```

```
    for(i = 0; i < Ng; i++)
```

```
        for(j = 0; j < Ng; j++)
```

```
            for(k = 0; k < Ng; k++)
```

```
{
```

```
    dist = sqrt(
```

```
        pow(x[p] - (double)i/Ng - half_size, 2) +
```

```
        pow(y[p] - (double)j/Ng - half_size, 2) +
```

```
        pow(z[p] - (double)k/Ng - half_size, 2));
```

```
    if(dist < Rmax)
```

```
        dummy += dist;
```

```
}
```

...yes, I know that you know

Loops optimization

OPT 1: avoid expensive functions (like **sqrt()**)

```
for(p = 0; p < Np; p++)
```

```
    for(i = 0; i < Ng; i++)
```

```
        for(j = 0; j < Ng; j++)
```

```
            for(k = 0; k < Ng; k++)
```

```
{
```

```
    dist2 =
```

```
        pow(x[p] - (double)i/Ng - half_size, 2) +
```

```
        pow(y[p] - (double)j/Ng - half_size, 2) +
```

```
        pow(z[p] - (double)k/Ng - half_size, 2);
```

```
    if(dist2 < Rmax2)
```

```
        dummy += sqrt(dist);
```

```
}
```

Loops optimization

OPT 1: avoid expensive functions (like **pow()**)

```
for(p = 0; p < Np; p++)
    for(i = 0; i < Ng; i++)
        for(j = 0; j < Ng; j++)
            for(k = 0; k < Ng; k++)
            {
                dx = x[p] - (double)i/Ng - half_size;
                dy = y[p] - (double)j/Ng - half_size;
                dz = z[p] - (double)k/Ng - half_size;

                dist2 = dx*dx + dy*dy + dz*dz;
                if(dist2 < Rmax2)
                    dummy += sqrt(dist);
            }
```

Loops optimization

OPT 1: avoid expensive operations (like **floating point division**)

```
for(p = 0; p < Np; p++)
```

```
for(i = 0; i < Ng; i++)
```

```
for(j = 0; j < Ng; j++)
```

```
for(k = 0; k < Ng; k++)
```

{

```
dx = x[p] - (double)i * Ng_inv - half_size;  
dy = y[p] - (double)j * Ng_inv - half_size;  
dz = z[p] - (double)k * Ng_inv - half_size;
```

```
dist2 = dx*dx + dy*dy + dz*dz;
```

```
if(dist2 < Rmax2)
```

```
dummy += sqrt(dist);
```

}

Loops optimization

OPT 2: avoid the avoidable + manual hoisting

```
for(i = 0; i < Ng; i++) {  
    dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 = dx2*dx2;  
  
    for(j = 0; j < Ng; j++) {  
        dy2 = y[p] - (double)j * Ng_inv - half_size;  
        dy2 = dy2*dy2;  
        dist2_xy = dx2 + dy2;  
  
        for(k = 0; k < Ng; k++) {  
            dz = z[p] - (double)k * Ng_inv - half_size;  
            dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                dummy += sqrt(dist);        } } }
```

Loops optimization

OPT 3: be clear with the compiler, use *local variables*

```
for(int i = 0; i < Ng; i++) {  
    double dx2 = x[p] - (double)i * Ng_inv - half_size;  
    dx2 = dx2*dx2;  
  
    for(int j = 0; j < Ng; j++) {  
        double dy2 = y[p] - (double)j * Ng_inv <...>  
        dy2 = dy2*dy2;  
        double dist2_xy = dx2 + dy2;  
  
        for(int k = 0; k < Ng; k++) {  
            double dz = z[p] - (double)k * ...;  
            double dist2 = dist2_xy + dz*dz;  
            if(dist2 < Rmax2)  
                dummy += sqrt(dist);      } } }
```

Loops optimization

OPT 4: be clear with the compiler , also through **register** directive

```
double register Ng_inv = 1.0 / Ng;
for(i = 0; i < Ng; i++) {
    dx2 = x[p] - (double)i * Ng_inv - half_size;
    dx2 = dx2*dx2;
    for(j = 0; j < Ng; j++) {
        dy2 = y[p] - (double)j * Ng_inv - half_size;
        dy2 = dy2*dy2;
        register dist2_xy = dx2 + dy2;
        for(k = 0; k < Ng; k++) {
            register dz = z[p] - (double)k * ...;

            register dist2 = dist2_xy + dz*dz;
            if(dist2 < Rmax2)
                dummy += sqrt(dist);      } } }
```

Loops optimization

Compilers (at -O2,3) are very good in *loop hoisting*, that is to extract invariant operations/quantities from within the loops.

However, their speculation is limited by the obscurity of the code (for instance: aliasing of memory locations).

Ex. from previous code snippet:

```
(double)<i,j,k> * Ng_inv + half_size
```

is performed N^3+N^2+N times, always returning the same values.
Hoisting will save $N(N^2+N^1+1)$ **mul** and **add** (and **mem** accesses).

```
double ijk[Ng]
for(i = 0; i < Ng; i++) ijk[i] = i * Ng_inv + half_size
```

Loops optimization

Also, limit the use of **global variables** or variables with a **wide scope**, and use local variable with limited scope every time it is possible (and it often is).

That helps the compiler in optimizing the registers and cache use.

Trivial ex:

```
int i;  
double temp;  
<...>  
for(i = 0; i < Ng; i++){  
    temp = <...>  
}
```

→ <...>

```
for(int i = 0; i < Ng;i++){  
    double temp = <...>  
}
```

Loops and branches

Loops optimizations: branches

Conditional branches should be avoided as much as possible inside loops:

- moving them outside the loop and writing specialized loops
- performing variables/quantities set-up pree-emptively outside the loop
- using pointers to functions instead of selecting functions inside the loop
- substituting conditional branches with different operations

Loops optimizations: branches

ex 1: Taking decisions before and outside the loop

```
for(i = 1; i < top; i++)  
{  
    if(case1 == 0) {  
        if(case2 == 0) {  
            if(case3 == 0)  
                result += i;  
            else  
                result -= i;  
        }  
        else {  
            if(case3 == 0)  
                result *= i;  
            else  
                result /= i;  
        }  
    }  
    else {  
        if(case2 == 0) {  
            if(case3 == 0)  
                result += log10((double)i);  
            else  
                result -= log10((double)i);  
        }  
        else {  
            if(case3 == 0)  
                result *= sin((double)i);  
            else  
                result /= (sin((double)i) +  
                           cos((double)i));  
        }  
    }  
}
```

Loops optimizations: branches

ex. 1 : Taking decisions before and outside the loop

- define a specialized function for each case
- **before** and **outside** the loop set a function pointer to the right function

```
void (*func)(double *, int);  
<here make func pointing to the right place>
```

```
double temp    = 0;  
double result = 0;  
for(i = 1; i < top; i++)  
{  
    func( & temp, i);  
    result = temp;  
}
```

→ ...just have a look at the live results..

Loops optimizations: branches

However:

- Using function pointers you may incur in additional overhead due to function call.
If the code snippets in different if-branches (or at least the most executed ones) are large/expensive, it might well be pointless (in modern CPUs).
- “Unrolling” the if-tree outside the for – then having multiple for loops :

```
if (case1 == 0) {  
    if (case2 == 0) {  
        if (case3 == 0) {  
            for(i = 1; i < top; i++)  
                result += i;  
        } else  
            for(i = 1; i < top; i++)  
                result -= i;  
    }  
}
```

may be highly unpractical if the branches are big piece of code.
There's no really a Swiss-knife recipe.

→ ...just have a look at the live results..

Loops optimizations: branches

ex. 2 : restructuring code

Consider the following code snippet

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```

Loops optimizations: branches

ex. 2 : restructuring code

Consider the following code snippet (*)

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

qsort(data, SIZE, sizeof(int), compare);

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```

(*) of course, you are adding an overhead due to the sorting routine, so the total running time may be even larger. Moreover, you should have all the values available so that does not work for real-time streamings. However, the point here is to focus on how – in general – it is better to avoid conditionals inside loop, with any possible trick or change in workflow

→ ...just have a look at the live results..

Loops optimizations: branches

ex. 2 : restructuring code

We can do even better, without adding operations

```
// generate random numbers
for (cc = 0; cc < SIZE; cc++)
    data[cc] = rand() % TOP;

// take action depending on their value
for (ii = 0; ii < SIZE; ii++)
{
    t = (data[ii] - PIVOT -1) >> 31;
    sum += ~t & data[ii];
}
```

→ ...just have a look at the live results..

Loops optimizations: branches

gcc -O

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred
sum is 983597794767, elapsed seconds: 5.40445
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow
sum is 983597794767, elapsed seconds: 2.23186
(in total: 2.44473 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart
sum is 983597794767, elapsed seconds: 2.8878
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ■
```

gcc -O3

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03
sum is 983597794767, elapsed seconds: 0.660148
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03
sum is 983597794767, elapsed seconds: 0.650005
(in total: 0.795181 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03
sum is 983597794767, elapsed seconds: 0.679286
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ■
```

Loops optimizations: branches

gcc -O3

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03
sum is 983597794767, elapsed seconds: 0.660148
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03
sum is 983597794767, elapsed seconds: 0.650005
(in total: 0.795181 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03
sum is 983597794767, elapsed seconds: 0.679286
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% █
```

gcc -O3 -march=native

```
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.03n
sum is 983597794767, elapsed seconds: 0.217864
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.wow.03n
sum is 983597794767, elapsed seconds: 0.215645
(in total: 0.355377 seconds)
luca@GGG:~/code/HPC_LECTURES/branch_prediction/branch_prediction_2% ./branchpred.smart.03n
sum is 983597794767, elapsed seconds: 0.224288
```

Loops optimizations: branches

What changes in the base version with -O3 ? → conditional move

Modern CPUs have the capability of performing *conditional move*, i.e to execute **concurrently** both branches of a conditional – if they are “simple enough” – and to select the right result upon the evaluation of the conditional

perform op1 → res in AX

perform op2 → res in BX

compare

if flag → mov BX in AX

HOWEVER: loops with conditionals can not be fully vectorized !!

Loops optimizations: branches

Why the difference in the base v. between **-O3** and **-O3 -march=native** ?

gcc -O3

.L8:

```
    movdqu  xmm0, XMMWORD PTR [rax]
    movdqu  xmm6, XMMWORD PTR [rax]
    movdqa  xmm2, xmm4
    add     rax, 16
    → pcmpgt d xmm0, xmm5
    pand   xmm0, xmm6
    pcmpgt d xmm2, xmm0
    movdqa  xmm3, xmm0
    punpckldq  xmm3, xmm2
    punpckhdq  xmm0, xmm2
    paddq  xmm1, xmm3
    paddq  xmm1, xmm0
    cmp     rax, rcx
    jne     .L8
```

- compare **4 integers** at a time using xmmX registers, that are common to x86_64 architectures.

- increase the counter by **4 int**

12 instructions to process **4 int** (3 ins/element)

gcc -O3 -march=native

.L8:

```
    vmovdqu  ymm2, YMMWORD PTR [rax]
    add     rax, 32
    → vpcmpgt d ymm0, ymm2, ymm3
    vpand   ymm0, ymm0, ymm2
    vpmovsxdq ymm2, xmm0
    vextracti128  xmm0, ymm0, 0x1
    vpaddq  ymm1, ymm2, ymm1
    vpmovsxdq ymm0, xmm0
    vpaddq  ymm1, ymm0, ymm1
    cmp     rax, rcx
    jne     .L8
```

bytes reshuffling inside regs to add one int at a time
The v prefix tells you these are AVX2 256-bits instr.

compare **8 integers** at a time using ymmX registers. This requires AVX2 that is set on by **-march=native** for this CPU

increase the counter by **8 int**

9 instructions to process **8 int** (1.12 ins/element)

Comments on the previous slides

Your branch predictor is quite good – as good as being right about ~95% of times in real codes – but still can not be right all the times. And penalties in this cases are quite severe (*).

Compiler's optimization in this simple case is able to convert the conditional in a `cmove`-style *when it is available*, which may makes the code as efficient as with the non-branching versions (or even more, because it is simpler).

Letting the compiler using the best instructions for the CPU also translate in partial vectorization, that further push performance.

(*) refer to the section on pipelines to understand why in more detail.

The example before was purposely built..

Normally it does not happen that you have cases so perfectly suited for some kind of bitwise trickery.

Then, provided that you should always catch the opportunity, there is another, more general, way.

That amounts to re-define your if condition in a different way: instead of having a branch that depends on the condition evaluation, express it in two different values to be used in the subsequent lines of code.

Loops optimizations: branches

```
for (ii = 0; ii < SIZE; ii++)
{
    if (data[ii] > PIVOT)
        sum += data[ii];
}
```

Can be rephrased as

```
for (ii = 0; ii < SIZE; ii++)
{
    acc = ( data[ii]>PIVOT )? data[ii] : 0;
    sum += acc;
}
```

Loops optimizations: branches

Let's look to another practical example

You have 2 arrays, A and B, and you want to swap their elements so that

$$A[i] \geq B[i]$$

for all i .

A straightforward implementation would be:

```
for (i = 0; i < SIZE; i++)
{
    if ( A[i] < B[i] )
    {
        t = B[i];
        B[i] = A[i];
        A[i] = t;
    }
}
```

Loops optimizations: branches

However, that implementation suffers exactly of the same problem we have just discussed.

An alternative way to write the same code, but in a more effective style is:

```
for (i = 0; i < SIZE; i++)
{
    int min = A[i] > B[i] ? B[i] : A[i];
    int max = A[i] >= B[i] ? A[i] : B[i];

    A[i] = max;
    B[i] = min;
}
```

Loops optimizations: branches

Different forms of the conditional statement we are considering. They may seem awkward, but let us look at them in more detail in the following slides.

standard

```
for (uint ii = 0; ii < SIZE; ii++)
{
    if ( B[ii] > A[ii] )
    {
        int t = A[ii];
        A[ii] = B[ii];
        B[ii] = t;
    }
}
```

smart2

```
for (uint ii = 0; ii < SIZE; ii++)
{
    int register t = -(A[ii]<B[ii]);
    int register x = A[ii]^B[ii];
    A[ii] = A[ii]^(x & t);
    B[ii] = B[ii]^(x & t);
}
```

smart

```
for (uint ii = 0; ii < SIZE; ii++)
{
    int max = (A[ii]>B[ii])? A[ii]:B[ii];
    int min = (A[ii]>B[ii])? B[ii]:A[ii];
    A[ii] = max;
    B[ii] = min;
}
```

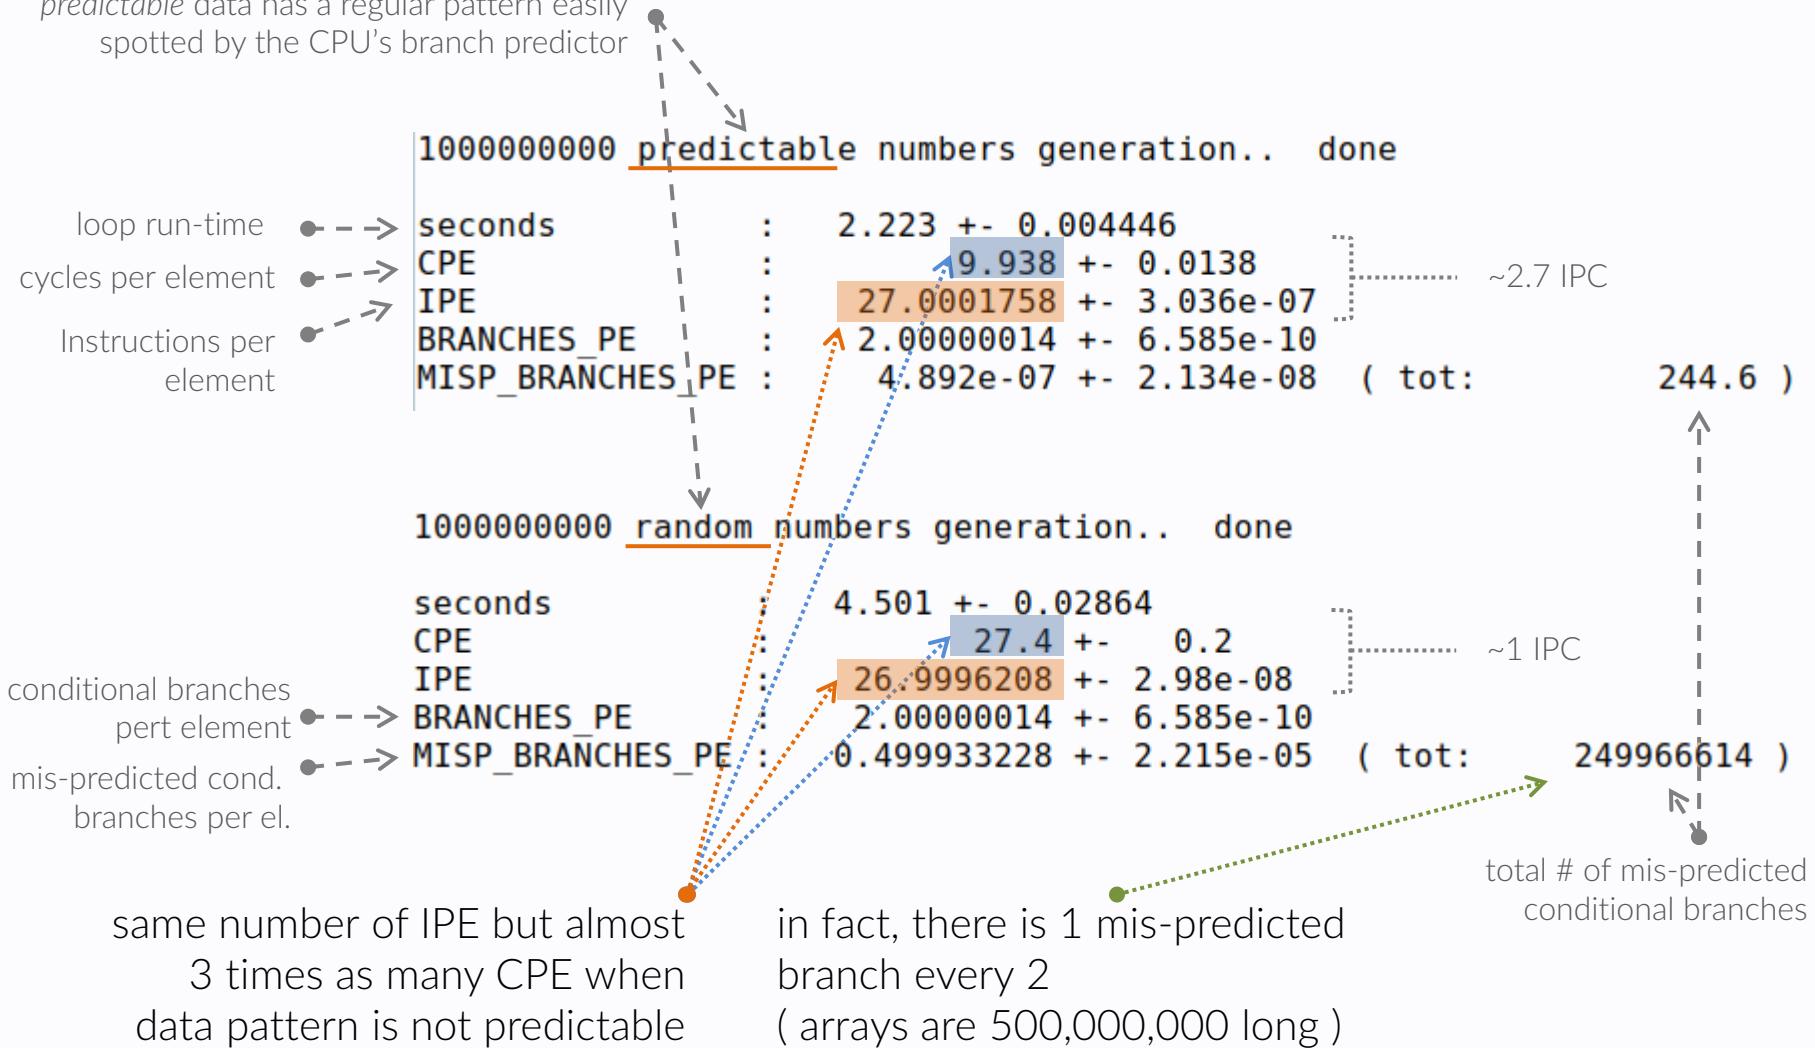
smart3

```
for (uint ii = 0; ii < SIZE; ii++)
{
    int d = A[ii]-B[ii];
    d &= (d >> 31);
    A[ii] = A[ii] - d;
    B[ii] = B[ii] + d;
}
```

Loops optimizations: branches

Standard implementation, -O0

predictable data has a regular pattern easily spotted by the CPU's branch predictor

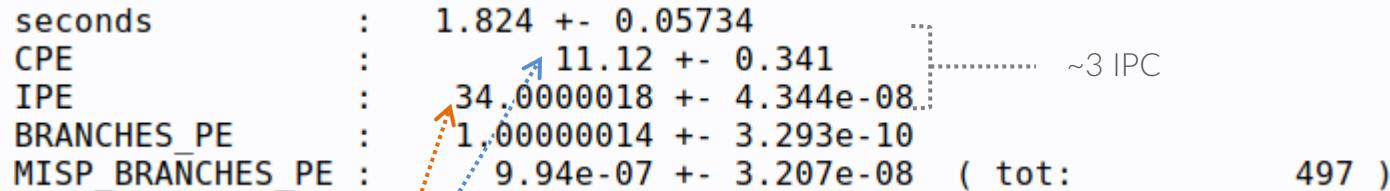


Loops optimizations: branches

smart implementation, -O0

predictable

```
seconds : 1.824 +- 0.05734
CPE      : 11.12 +- 0.341
IPE      : 34.0000018 +- 4.344e-08
BRANCHES_PE : 1.00000014 +- 3.293e-10
MISP_BRANCHES_PE : 9.94e-07 +- 3.207e-08 ( tot: 497 )
```

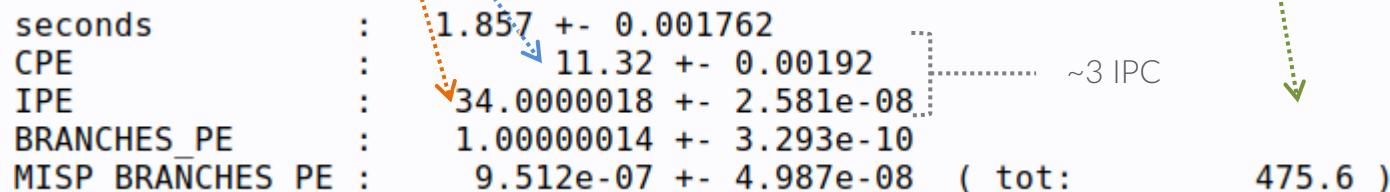


number of IPE is larger than for standard case, but the CPE is stable !

mis-predicted branches are very few and comparable in both cases

random

```
seconds : 1.857 +- 0.001762
CPE      : 11.32 +- 0.00192
IPE      : 34.0000018 +- 2.581e-08
BRANCHES_PE : 1.00000014 +- 3.293e-10
MISP_BRANCHES_PE : 9.512e-07 +- 4.987e-08 ( tot: 475.6 )
```



Loops optimizations: branches

smart2 implementation, -O0

predictable

seconds	:	1.811 +- 0.002471
CPE	:	11.04 +- 0.00246
IPE	:	37.0000018 +- 3.799e-08
BRANCHES_PE	:	1.00000014 +- 3.293e-10
MISP_BRANCHES_PE	:	1.0236e-06 +- 7.194e-08 (tot: 511.8)

number of IPE is [larger|smaller] than
for [stdrd|smart] case, and the CPE is
stable !

~3.4 IPC

511.8

mis-predicted branches are still very
few and comparable in both cases

random

seconds	:	1.745 +- 0.05743
CPE	:	10.63 +- 0.344
IPE	:	37.0000018 +- 4.344e-08
BRANCHES_PE	:	1.00000014 +- 3.293e-10
MISP_BRANCHES_PE	:	1.074e-06 +- 5.584e-08 (tot: 537)

~3.5 IPC

537

Loops optimizations: branches

smart3 implementation, -O0

predictable

seconds	:	1.628 +- 0.01015	
CPE	:	9.993 +- 0.0464	
IPE	:	32.0000017 +- 2.356e-08 ~3.2 IPC
BRANCHES_PE	:	1.00000014 +- 3.293e-10	
MISP_BRANCHES_PE	:	3.728e-07 +- 8.898e-08	(tot: 186.4)

number of IPE as smart case, larger than for stdrd, the CPE is stable

mis-predicted branches are still very few and comparable in both cases

random

seconds	:	1.688 +- 0.03554	
CPE	:	10.36 +- 0.211	
IPE	:	32.0000017 +- 2.98e-08 ~3.1 IPC
BRANCHES_PE	:	1.00000014 +- 3.293e-10	
MISP_BRANCHES_PE	:	3.1e-07 +- 5.183e-08	(tot: 155)

Comments on the previous slides

The standard implementation relies on the ability of your CPU's branch predictor to guess the correct data pattern.

When it is successful, it is **really** so (the lowest CPE and IPE).

However, whenever the data pattern is not guessable by the branch predictor things quickly become really weird.

Writing the code differently may make you loosing something in terms of CPE/IPE (with respect to the best possible standard case) but not really in terms of time-to-solution.

And, above all, the code behaviour is **stable** with both predictable and unpredictable patterns.

std implementation, -O0

```
.L20:
    mov    eax, DWORD PTR -88[rbp] # ii
    lea    rdx, 0[0+rax*4]
    mov    rax, QWORD PTR -56[rbp] # B
    add    rax, rdx
    mov    edx, DWORD PTR [rax]
    mov    eax, DWORD PTR -88[rbp] # ii
    lea    rcx, 0[0+rax*4]
    mov    rax, QWORD PTR -64[rbp] # A
    add    rax, rcx
    mov    eax, DWORD PTR [rax]
# branchpred2.c:116:      if ( B[ii] > A[ii] )
    cmp    edx, eax
    jle    .L19
# branchpred2.c:118:      int t = A[ii];
    mov    eax, DWORD PTR -88[rbp] # ii
    lea    rdx, 0[0+rax*4]
    mov    rax, QWORD PTR -64[rbp] # A
    add    rax, rdx
# branchpred2.c:118:      int t = A[ii];
    mov    eax, DWORD PTR [rax]
    mov    DWORD PTR -84[rbp], eax # t
# branchpred2.c:119:      A[ii] = B[ii];
    mov    eax, DWORD PTR -88[rbp] # ii
    lea    rdx, 0[0+rax*4]
    mov    rax, QWORD PTR -56[rbp] # B
    add    rax, rdx #
# branchpred2.c:119:      A[ii] = B[ii];
    mov    edx, DWORD PTR -88[rbp] # ii
    lea    rcx, 0[0+rdx*4]
    mov    rdx, QWORD PTR -64[rbp] # A
    add    rdx, rcx
# branchpred2.c:119:      A[ii] = B[ii];
    mov    eax, DWORD PTR [rax]
# branchpred2.c:119:      A[ii] = B[ii];
    mov    DWORD PTR [rdx], eax
# branchpred2.c:120:      B[ii] = t;
    mov    eax, DWORD PTR -88[rbp] # ii
    lea    rdx, 0[0+rax*4],
    mov    rax, QWORD PTR -56[rbp] # B
    add    rdx, rax #
# branchpred2.c:120:      B[ii] = t;
    mov    eax, DWORD PTR -84[rbp] # t
    mov    DWORD PTR [rdx], eax
.L19:
# branchpred2.c:112:      for (uint ii = 0; ii < SIZE; ii++)
    add    DWORD PTR -88[rbp], 1 # ii,
.L18:
# branchpred2.c:112:      for (uint ii = 0; ii < SIZE; ii++)
    mov    eax, DWORD PTR -88[rbp] # ii
    cmp    eax, DWORD PTR -120[rbp]      # SIZE
    jb    .L20
```

1 cmp instr. with
a following jump



2 cmov instr.
can use 2
pipelines


smart implementation, -O0

```
.L19:
# branchpred2.c:122:      max = A[ii] > B[ii] ? A[ii] : B[ii];
    mov    eax, DWORD PTR -92[rbp] # tmp229, ii
    cdqe
    lea    rdx, 0[0+rax*4]
    mov    rax, QWORD PTR -56[rbp] # B
    add    rax, rdx
    mov    edx, DWORD PTR [rax]
    mov    eax, DWORD PTR -92[rbp] # ii
    cdqe
    lea    rcx, 0[0+rax*4] # _51,
    mov    rax, QWORD PTR -64[rbp] # A
    add    rax, rcx
    mov    eax, DWORD PTR [rax]
# branchpred2.c:122:      max = A[ii] > B[ii] ? A[ii] : B[ii];
    cmp    edx, eax
    cmovge eax, edx
    mov    DWORD PTR -88[rbp], eax # max
# branchpred2.c:123:      min = A[ii] < B[ii] ? A[ii] : B[ii];
    mov    eax, DWORD PTR -92[rbp] # ii
    cdqe
    lea    rdx, 0[0+rax*4] # _55,
    mov    rax, QWORD PTR -56[rbp] # B
    add    rax, rdx # _56, _55
    mov    edx, DWORD PTR [rax]
# branchpred2.c:123:      min = A[ii] < B[ii] ? A[ii] : B[ii];
    mov    eax, DWORD PTR -92[rbp] # ii
    cdqe
    lea    rcx, 0[0+rax*4] # _59,
    mov    rax, QWORD PTR -64[rbp] # A
    add    rax, rcx
    mov    eax, DWORD PTR [rax]
# branchpred2.c:123:      min = A[ii] < B[ii] ? A[ii] : B[ii];
    cmp    edx, eax
    cmovle eax, edx
    mov    DWORD PTR -84[rbp], eax # min
# branchpred2.c:131:      A[ii] = max;
    mov    eax, DWORD PTR -92[rbp] # ii
    cdqe
    lea    rdx, 0[0+rax*4]
    mov    rax, QWORD PTR -64[rbp] # A
    add    rdx, rax
# branchpred2.c:131:      A[ii] = max;
    mov    eax, DWORD PTR -88[rbp] # max
    mov    DWORD PTR [rdx], eax
# branchpred2.c:132:      B[ii] = min;
    mov    eax, DWORD PTR -92[rbp] # ii
    cdqe
    lea    rdx, 0[0+rax*4] # _66,
    mov    rax, QWORD PTR -56[rbp] # B
    add    rdx, rax
# branchpred2.c:132:      B[ii] = min;
    mov    eax, DWORD PTR -84[rbp] # min
    mov    DWORD PTR [rdx], eax
# branchpred2.c:107:      for (ii = 0; ii < SIZE; ii++)
    add    DWORD PTR -92[rbp], 1 # ii,
.L18:
# branchpred2.c:107:      for (ii = 0; ii < SIZE; ii++)
    mov    eax, DWORD PTR -92[rbp] # ii
    cmp    eax, DWORD PTR -104[rbp] # SIZE
    jl    .L19
```

Loops optimizations: branches

All optimizations at a glance, with *predictable* data

opt	smart	data	time	+-	err	CPE	+-	err	INS	+-	err	BRE
00	N	P	1.697e+00	+-	2.876e-03	1.041e+01	+-	6.740e-04	2.700e+01	+-	3.650e-08	2.000e+00
00	smart	P	1.859e+00	+-	5.434e-02	1.118e+01	+-	3.320e-01	3.400e+01	+-	3.406e-07	1.000e+00
00	smart2	P	1.815e+00	+-	2.680e-03	1.114e+01	+-	9.030e-03	3.700e+01	+-	3.942e-08	1.000e+00
00	smart3	P	1.628e+00	+-	1.015e-02	9.993e+00	+-	4.640e-02	3.200e+01	+-	2.356e-08	1.000e+00
03	N	P	6.448e-01	+-	3.248e-03	3.578e+00	+-	2.550e-03	8.000e+00	+-	2.849e-08	2.000e+00
03	smart	P	4.548e-01	+-	7.833e-03	2.064e+00	+-	2.040e-02	4.250e+00	+-	3.953e-08	2.500e-01
03	smart2	P	4.506e-01	+-	1.419e-03	2.109e+00	+-	1.650e-02	4.500e+00	+-	3.723e-08	2.500e-01
03	smart3	P	4.332e-01	+-	2.937e-03	1.905e+00	+-	6.180e-03	3.500e+00	+-	2.542e-08	2.500e-01
03n	N	P	3.518e-01	+-	1.096e-02	1.357e+00	+-	1.410e-02	1.125e+00	+-	4.850e-09	2.500e-01
03n	smart	P	3.905e-01	+-	2.258e-03	1.367e+00	+-	1.030e-02	1.125e+00	+-	2.643e-08	1.250e-01
03n	smart2	P	4.114e-01	+-	1.941e-02	1.671e+00	+-	5.650e-02	1.750e+00	+-	3.118e-08	1.250e-01
03n	smart3	P	4.119e-01	+-	3.449e-03	1.523e+00	+-	1.500e-02	1.500e+00	+-	1.613e-09	1.250e-01

The standard implementation (label “N” in the table) exhibits a better behaviour at -O0 considering CPE and above all IPE (label “INS” in the table), whereas run times are comparable among different variants (std, smart, smart2 and smart3).

However, at -O3 its behaviour is the worst one, with CPE larger by ~75% and IPE larger by a factor of ~2.

Only with very aggressive optimization the compiler can generate a code comparable with the smartX ones

Loops optimizations: branches

All optimizations at a glance, with *non predictable* data

opt	smart	data	time	+ -	err	CPE	+ -	err	INS	+ -	err	BRE
00	N	R	4.426e+00	+ - 2.966e-02	2.725e+01	+ - 1.830e-01	2.700e+01	+ - 2.788e-08	2.000e+00			
00	smart	R	1.855e+00	+ - 1.822e-03	1.139e+01	+ - 1.860e-03	3.400e+01	+ - 3.161e-08	1.000e+00			
00	smart2	R	1.718e+00	+ - 5.001e-02	1.055e+01	+ - 2.940e-01	3.700e+01	+ - 3.942e-08	1.000e+00			
00	smart3	R	1.688e+00	+ - 3.554e-02	1.036e+01	+ - 2.110e-01	3.200e+01	+ - 2.980e-08	1.000e+00			
03	N	R	2.306e+00	+ - 1.549e-02	1.418e+01	+ - 8.650e-02	8.000e+00	+ - 4.292e-08	2.000e+00			
03	smart	R	4.178e-01	+ - 3.808e-02	2.138e+00	+ - 7.510e-02	4.250e+00	+ - 3.838e-08	2.500e-01			
03	smart2	R	4.517e-01	+ - 1.640e-03	2.098e+00	+ - 1.450e-02	4.500e+00	+ - 2.644e-08	2.500e-01			
03	smart3	R	4.321e-01	+ - 2.602e-03	1.910e+00	+ - 1.260e-02	3.500e+00	+ - 2.710e-08	2.500e-01			
03n	N	R	4.178e-01	+ - 3.130e-03	1.600e+00	+ - 6.140e-03	1.249e+00	+ - 2.661e-08	2.500e-01			
03n	smart	R	3.918e-01	+ - 1.255e-03	1.363e+00	+ - 1.260e-02	1.125e+00	+ - 2.602e-08	1.250e-01			
03n	smart2	R	4.107e-01	+ - 1.860e-02	1.653e+00	+ - 3.830e-02	1.750e+00	+ - 9.429e-09	1.250e-01			
03n	smart3	R	4.131e-01	+ - 1.929e-03	1.516e+00	+ - 9.290e-03	1.500e+00	+ - 1.397e-09	1.250e-01			

When dealing with random data, the difference between std implementation and the other ones is even more obvious up to -O3, with the CPE being larger by a factor of ~3 and ~4 than in predictable data case at -O0 and -O3 respectively.

With very aggressive optimization the compiler can generate a code comparable with the smartX ones (for this very simple code snippet).

Loops optimizations: branches

Using the PGI compiler..

opt	smart	data	time	+-	err	CPE	+-	err	INS	+-	err	BRE
00	N	P	9.380e-01	+-	5.053e-03	5.443e+00	+-	4.480e-03	1.700e+01	+-	3.373e-08	2.000e+00
00	N	R	3.263e+00	+-	1.724e-02	1.982e+01	+-	2.920e-02	1.700e+01	+-	7.580e-08	2.000e+00
00	smart	P	1.671e+00	+-	5.070e-02	1.016e+01	+-	3.010e-01	3.300e+01	+-	2.788e-08	3.000e+00
00	smart	R	4.132e+00	+-	4.019e-02	2.525e+01	+-	2.350e-01	3.300e+01	+-	4.344e-08	3.000e+00
00	smart2	P	1.848e+00	+-	3.021e-03	1.126e+01	+-	1.530e-02	3.000e+01	+-	2.471e-08	1.000e+00
00	smart2	R	1.775e+00	+-	6.319e-02	1.082e+01	+-	3.790e-01	3.000e+01	+-	6.909e-08	1.000e+00
00	smart3	P	1.362e+00	+-	6.870e-02	8.283e+00	+-	4.140e-01	2.300e+01	+-	3.573e-08	1.000e+00
00	smart3	R	1.462e+00	+-	2.043e-03	8.873e+00	+-	8.040e-03	2.300e+01	+-	4.012e-08	1.000e+00
03	N	P	5.530e-01	+-	1.608e-03	3.010e+00	+-	6.270e-03	7.500e+00	+-	5.952e-08	2.000e+00
03	N	R	2.343e+00	+-	1.076e-02	1.428e+01	+-	6.810e-02	7.500e+00	+-	2.998e-08	2.000e+00
03	smart	P	3.788e-01	+-	2.156e-03	1.772e+00	+-	4.020e-02	1.875e+00	+-	1.867e-08	2.500e-01
03	smart	R	3.780e-01	+-	1.939e-03	1.769e+00	+-	2.440e-02	1.875e+00	+-	1.073e-08	2.500e-01
03	smart2	P	4.013e-01	+-	1.917e-03	2.210e+00	+-	4.540e-03	3.125e+00	+-	4.165e-09	2.500e-01
03	smart2	R	4.011e-01	+-	1.736e-03	2.212e+00	+-	6.520e-03	3.125e+00	+-	1.863e-09	2.500e-01
03	smart3	P	3.862e-01	+-	4.548e-03	2.080e+00	+-	2.630e-02	2.375e+00	+-	9.701e-09	2.500e-01
03	smart3	R	3.873e-01	+-	2.408e-03	2.067e+00	+-	9.120e-03	2.375e+00	+-	2.734e-08	2.500e-01
03n	N	P	5.538e-01	+-	2.179e-03	3.012e+00	+-	5.910e-03	7.500e+00	+-	1.070e-08	2.000e+00
03n	N	R	2.403e+00	+-	1.996e-02	1.464e+01	+-	1.010e-01	7.500e+00	+-	3.822e-08	2.000e+00
03n	smart	P	3.759e-01	+-	2.942e-03	1.776e+00	+-	2.770e-02	1.875e+00	+-	3.132e-08	2.500e-01
03n	smart	R	3.831e-01	+-	2.355e-03	1.746e+00	+-	1.320e-02	1.875e+00	+-	7.552e-09	2.500e-01
03n	smart2	P	4.047e-01	+-	4.673e-03	2.225e+00	+-	9.760e-03	3.125e+00	+-	1.050e-08	2.500e-01
03n	smart2	R	4.019e-01	+-	1.069e-03	2.214e+00	+-	7.920e-03	3.125e+00	+-	2.281e-09	2.500e-01
03n	smart3	P	3.848e-01	+-	3.987e-03	2.054e+00	+-	1.290e-02	2.375e+00	+-	3.029e-08	2.500e-01
03n	smart3	R	3.836e-01	+-	1.914e-03	2.048e+00	+-	3.770e-03	2.375e+00	+-	2.684e-08	2.500e-01

Comments on the previous slides

As we have seen, the gcc compiler can generate a code comparable to what it does with smartX variants only with very aggressive optimization and using AVX 256-bits instructions.

When random data are used, the standard implementation exhibits a behaviour that is much worse than with data with a predictable pattern, whereas trying not to use conditional branches generates a more stable code.

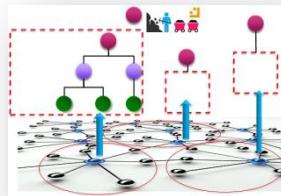
In this case, pgi compiler proves to be less able than gcc to generate optimal code, with -O3n level (*) still having a pronounced spike in CPE for random data.

Bottom-line is: **do not take it for granted that you lit up the compiler's optimization and everything will go seamlessly towards a triumph.**

(*) actually it is -O4 -fast -tp haswell -Mvect=simd:256

It may be really easy to get lost in “optimization”, in hunting every single line wondering why some incredible trick that – you’re convinced – should work, actually does not.

“Optimization” includes also optimizing your effort and your time, so always remember that the far most important ingredients are:



The algorithms that you choose



The data model you design



The overall quality, cleanliness and robustness of your code

That said, being really conscious of what happens at low-level helps you in writing an high-quality, clean and robust code.

Loops optimizations: branches

A last example.. about the fact that design and simplicity are the best move

Just changing point of view sometimes may help:

```
for ( j = 0; j < N; j++ )
    for ( i = 0; i < M; i++ )
    {
        if ( i > j )
            matrix[j][i] = 1.0;
        else if ( i < j )
            matrix[j][i] = -1.0;
        else
            matrix[j][i] = 0.0;
    }
```

Loops optimizations: branches

Can be easily re-written with no conditional evaluations

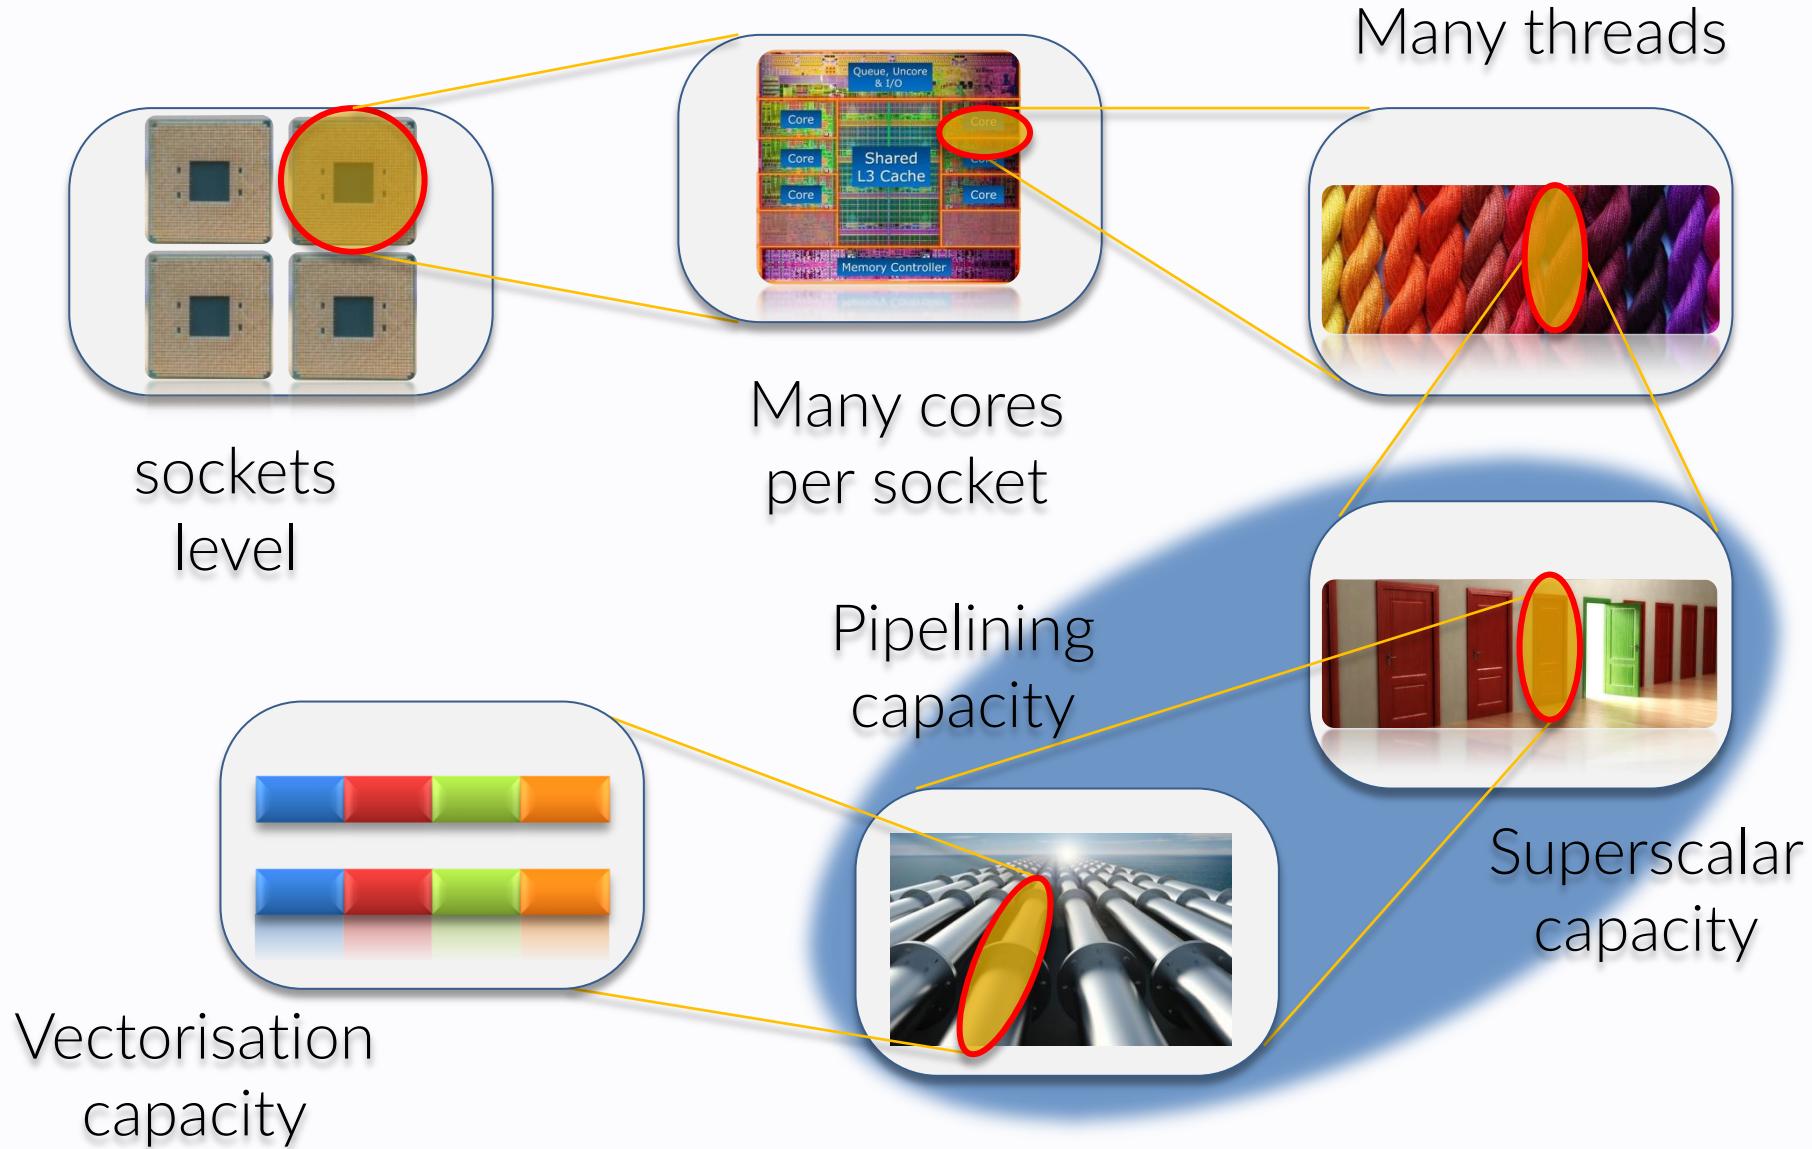
```
for ( j = 0; j < N; j++ )
{
    for ( i = 0; i < j; i++ )
        matrix[j][i] = -1.0;

    matrix[j][i] = 0.0;

    for ( i = j+1; i < M; i++ )
        matrix[j][i] = 1.0;
}
```

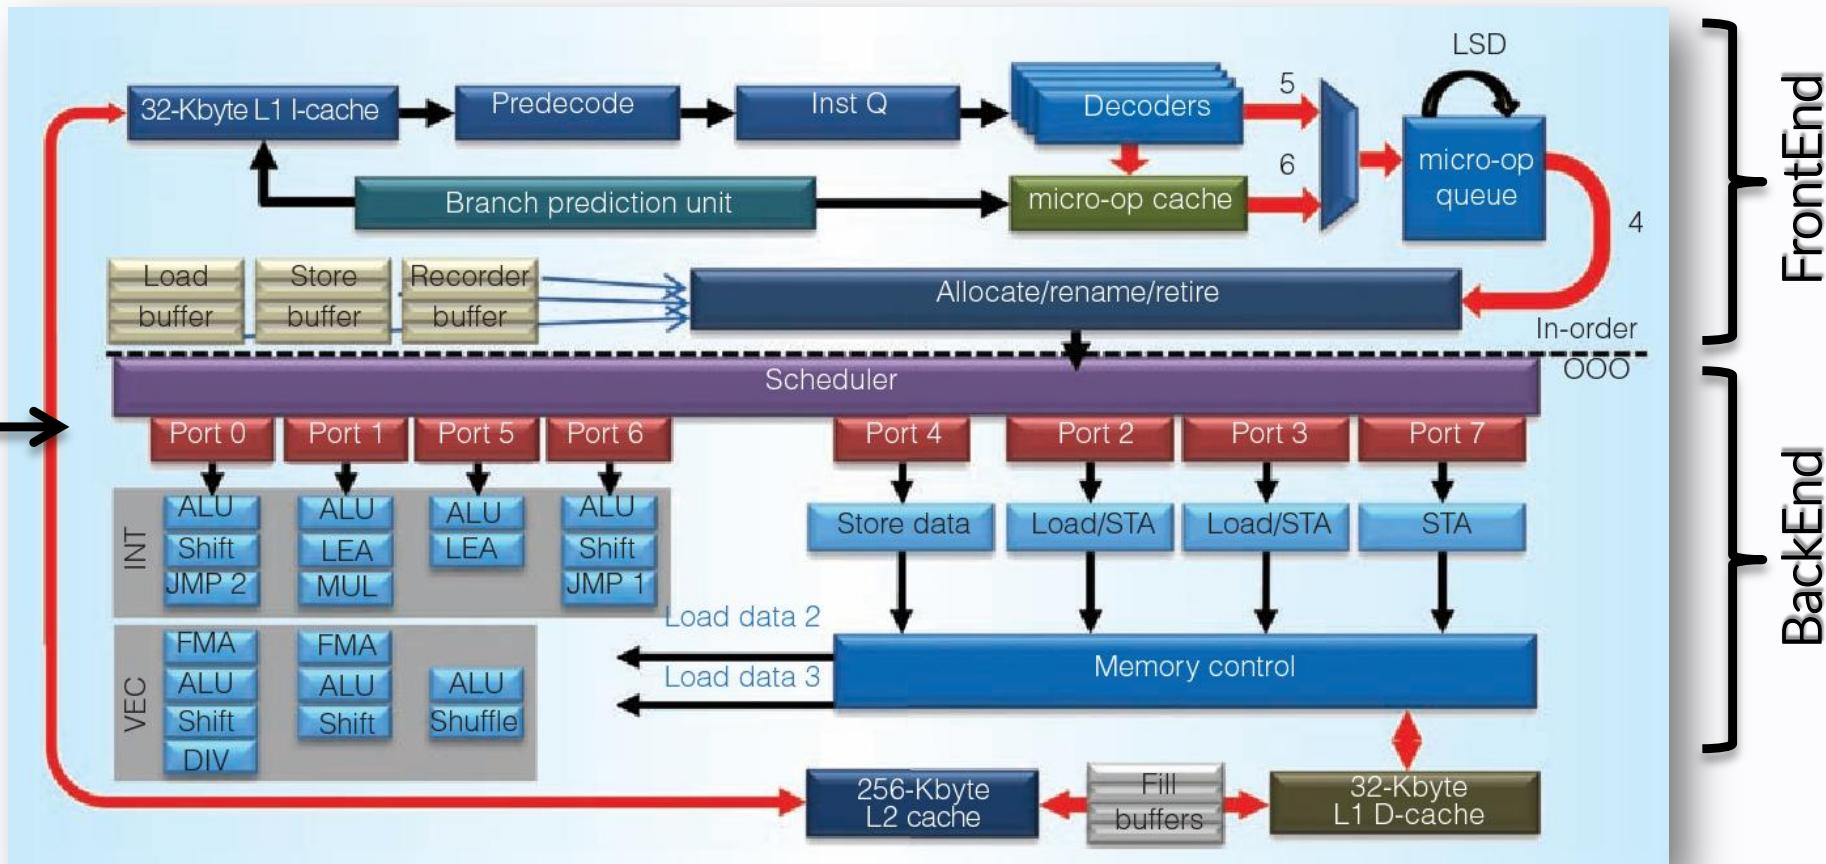
Pipelining

Architectural hierarchy of modern CPUs



Architectural hierarchy of modern Cores

More than 1 port is available to execute CPU instructions, although different units have different specializations (ALU, LEA, SHIFT, FMA, ...) : that is superscalar capacity, i.e. the capacity of executing more than 1 instructions per cycle.



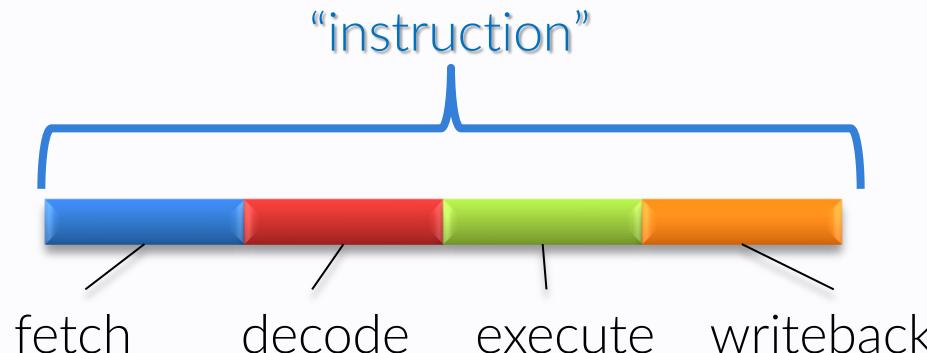
6th generation SkyLake core micro-arch.

Pipelines

It would be obvious to think that an “instruction” is a kind of *atomic* operation that the CPU perform as a whole. Indeed that was true until the mid of the 80s.

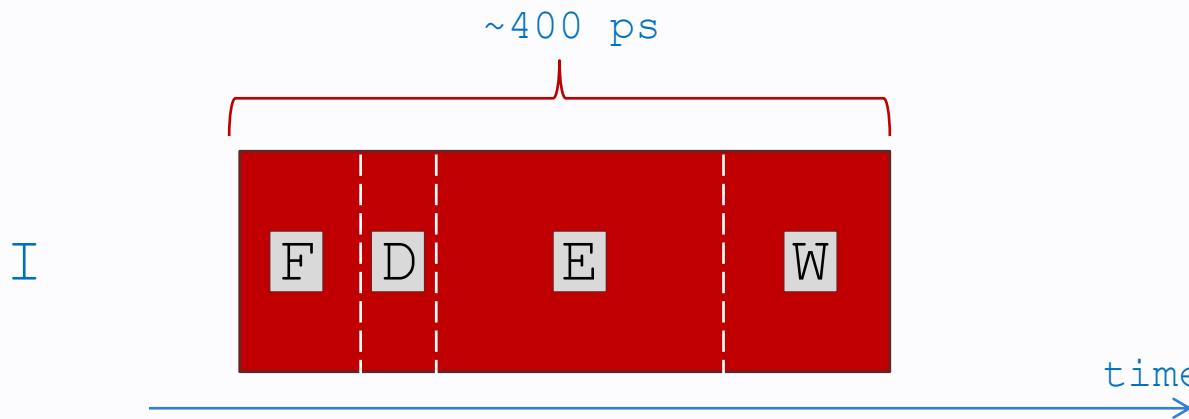
If you think carefully about it, it is easy to understand that actually an “instruction” involves at least the following *independent* steps:

1. It must be recalled from memory/Icache (**fetching**)
2. It must be “understood and interpreted” by the CPU (**decoding**)
3. It must be **executed**
4. The result must be accounted in memory (**writeback**)



Pipelines

Let us consider the instruction as a whole:



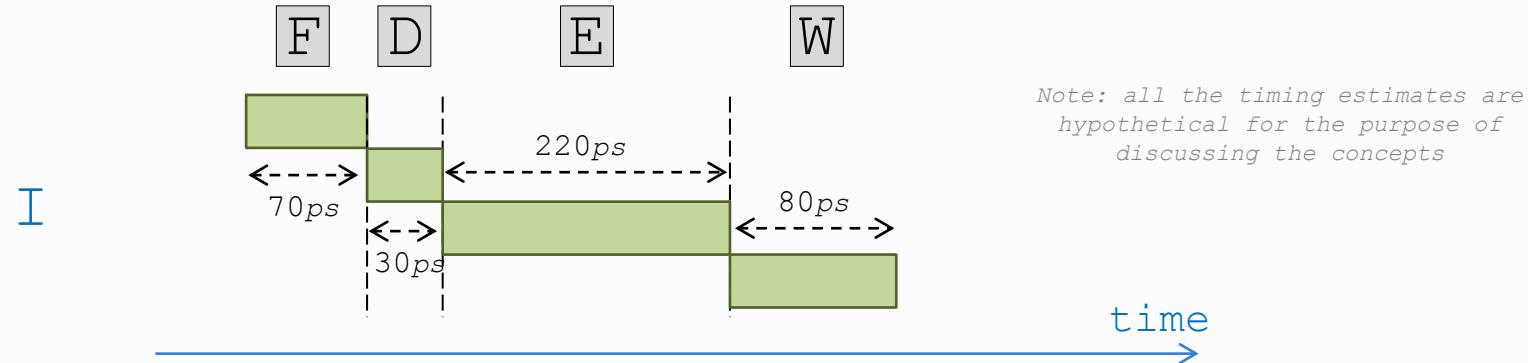
If all the four stages take $\sim 400\text{ps}$, we then would obtain a **throughput** of 2.5GIPS (giga-instructions per second).

400ps is also the total time required to get a result from an instruction, and so it is the **latency** of the instruction.

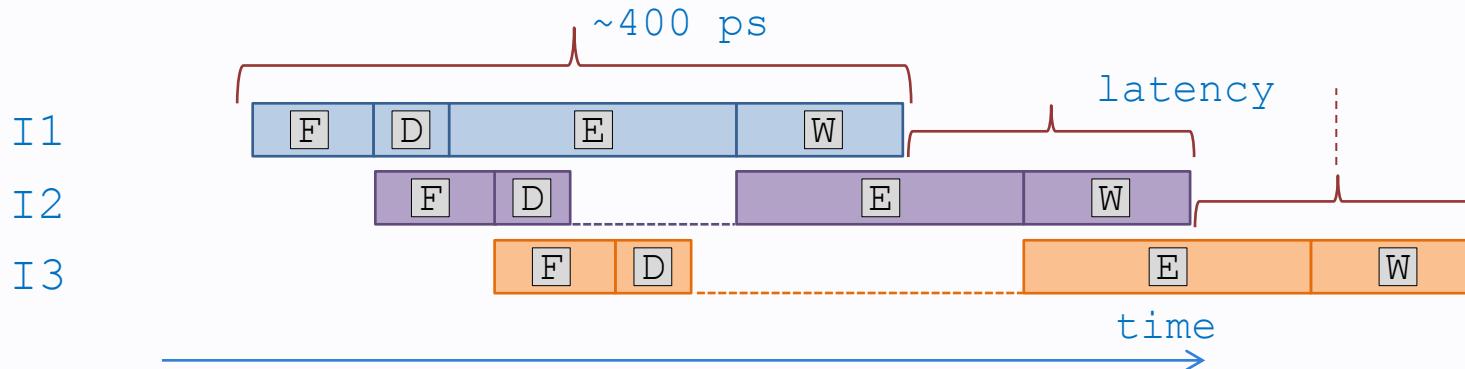
However, if we were able to “detach” the four stages, we could organize things differently, like in a car-building chain, or even at the mensa of the university.

Pipelines

If different and independent logical units of the CPU are in charge of each step:



they actually could operate subsequently on **different instructions**:

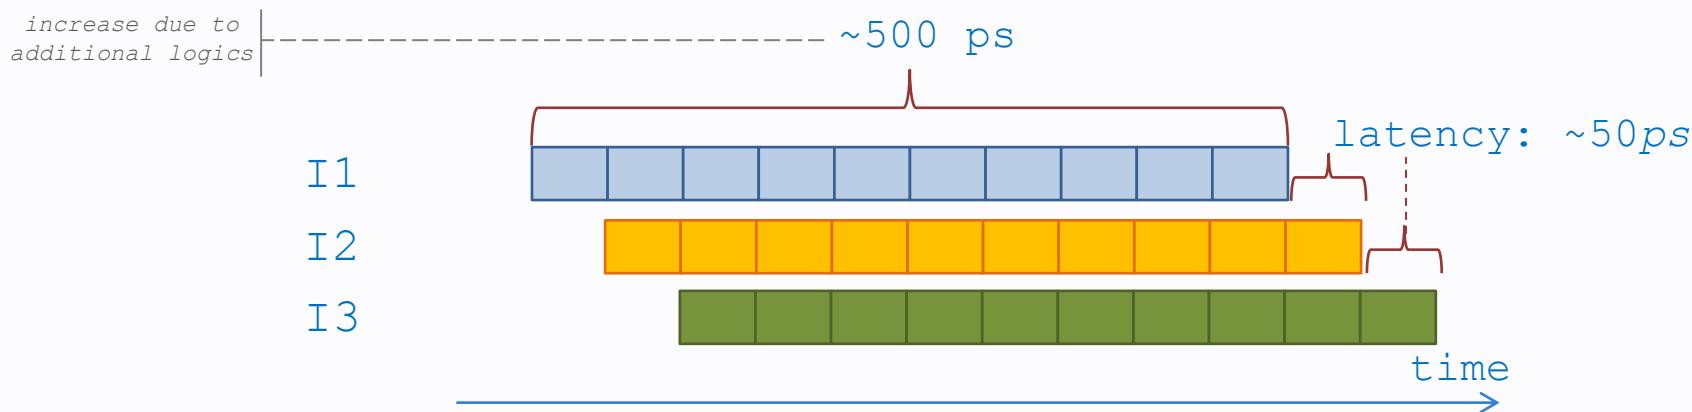


If the stage delays are not uniform, the throughput is limited by the latency $F + (D+E) - (F+D) = E \sim 220\text{ps}$, which means we have a throughput of $\sim 4.5\text{GIPS}$ just because of logic units separation.

Pipelines

Therefore, introducing the instructions pipelining, we can increase the **throughput** of our system by a large factor.

However, the efficiency of the pipelines is limited by its longest stage: the better option would be to have all equal stages, for instance further subdividing each stage – especially the most demanding ones (*) :



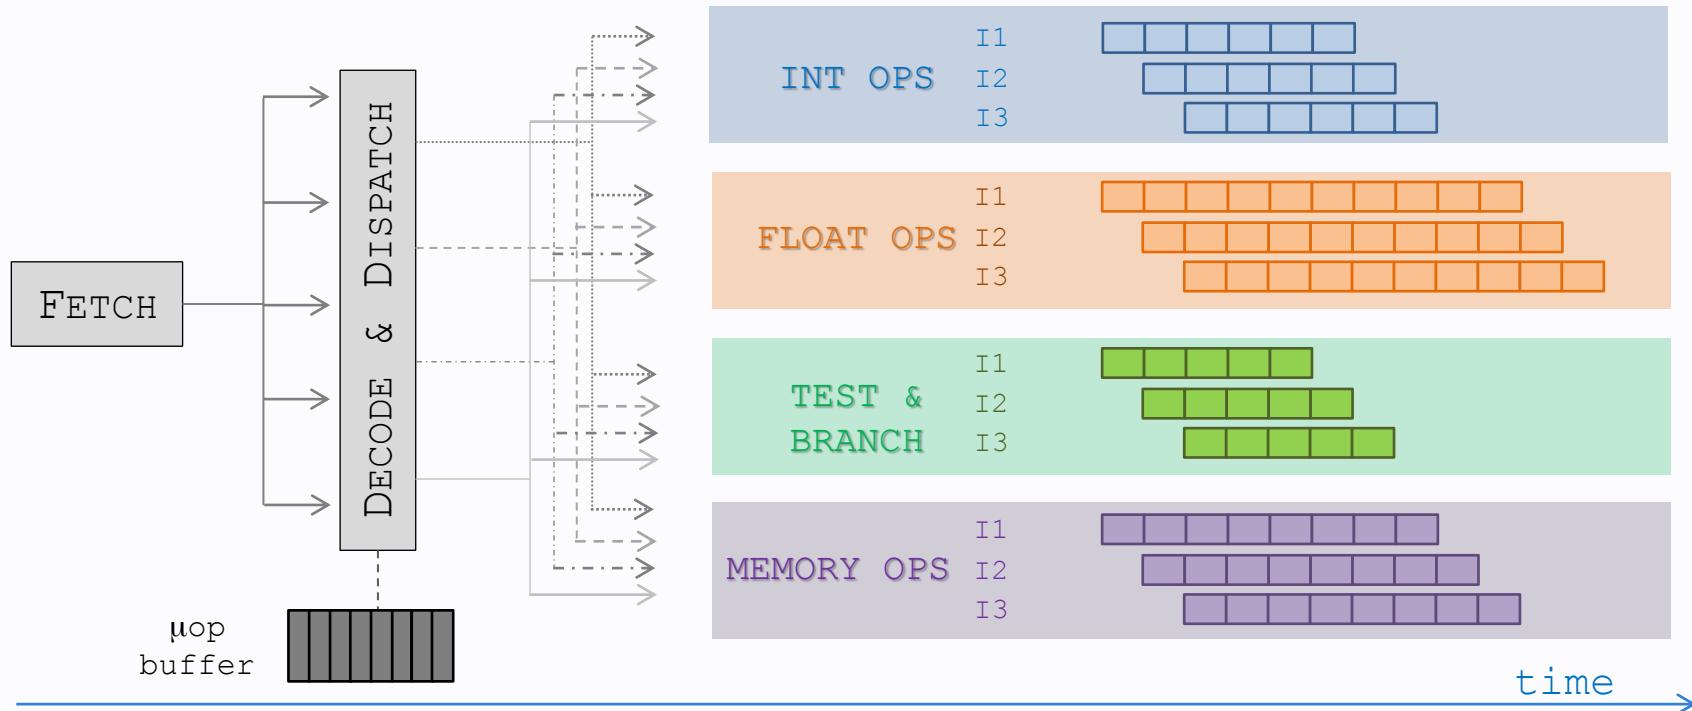
Now the throughput of our system has increased to 1 instruction retired every 50ps, i.e. **20GIPS** (note: these figures are not real, they have only an illustrative purpose).

(*) this is called *superpipelining*: modern CPUs may have 10-20 stages per pipeline.

Pipelines and superscalar processors

Pipelining is then about increasing the throughput of a system, and as we have seen it could be really effective. However, there is more that can be done working on the **execution stage** that, of course, encompasses a large number of different *functional units*, performing a **different and independent tasks**.

With an enhancement of the decode/dispatch stage, we can address multiple pipelines that can be active at the same time:



Pipeline hazards

Pipelining a system with *feedbacks* (i.e.: in which a result from an instruction can be the input of a following instruction, or can change the workflow) can expose the system to *hazards* when there are dependencies between successive instructions.

There are two forms of dependency:

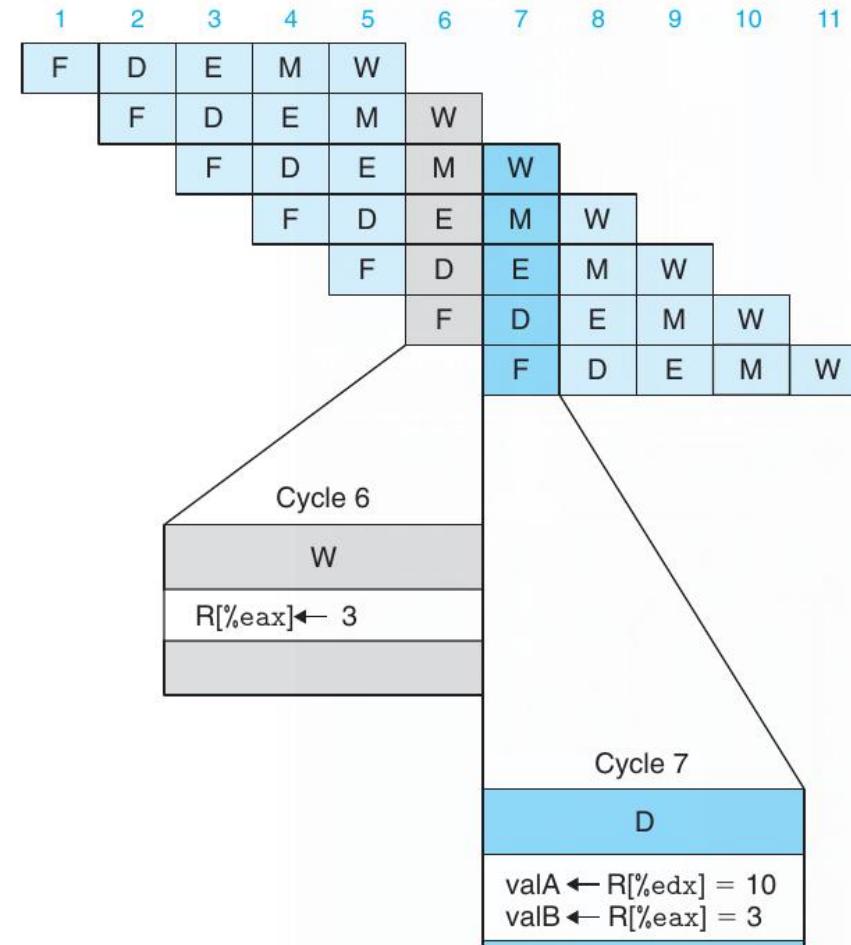
1. *data dependency*: the result computed by the i th instruction are used by one or more following instructions;
2. *control dependency*: the result of an instruction determines the value of the pc, i.e. the location of the next instruction to be executed.

When a dependency can lead to a wrong result, it is called an *hazard*.

Pipeline hazards

It is not of your interest in this course to discuss the details of how to implement the pipelines and logics in CPUs.
However, let's just discuss an example (*)

```
# prog1
0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: nop
0x00f: addl %edx,%eax
0x011: halt
```



(*) taken from computer systems: a
programmers perspective, Pearson, 2015

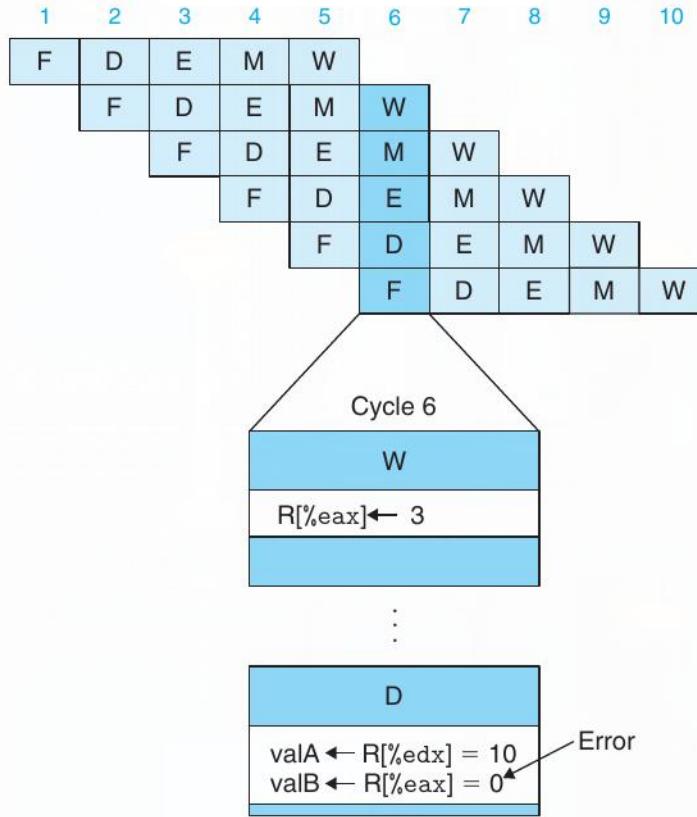
Pipeline hazards

prog2

```

0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: nop
0x00e: addl %edx,%eax
0x010: halt

```

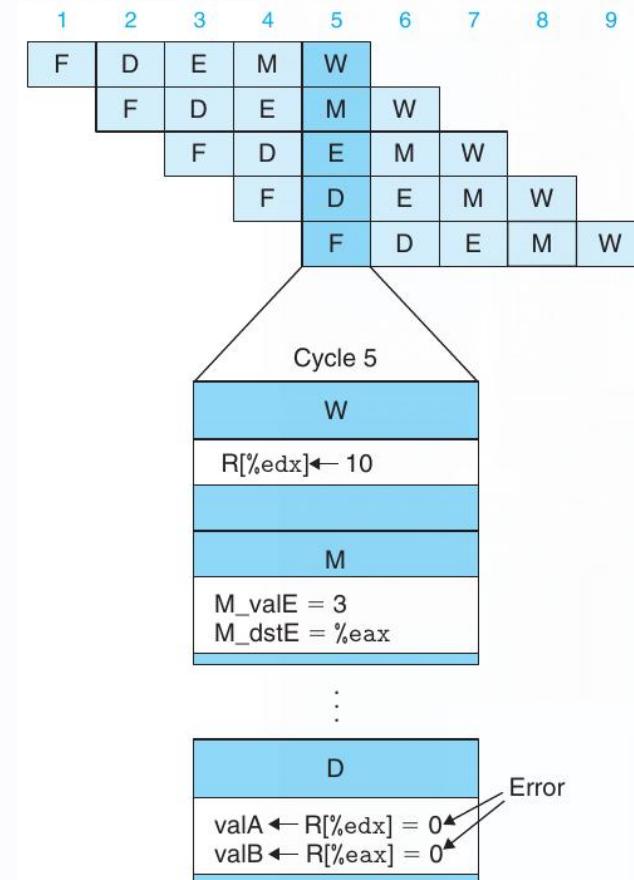


prog3

```

0x000: irmovl $10,%edx
0x006: irmovl $3,%eax
0x00c: nop
0x00d: addl %edx,%eax
0x00f: halt

```



Pipeline hazards

So, whenever there is a data dependency hazard, the processor injects *bubbles* in the pipeline, during which no work gets done.

The number of cycles between when an instruction's execute stage begins and when its result is available for other instructions' usage, is the *latency* of the instruction and grows with the pipeline deepness (the number of stages).

Typical latencies in modern processors (read the manual of yours) range from 1cyc for integer ops, to 3-6cyc for add and mul flop to $\geq 10-20$ for div flop, ~ 50 for trig flop.

Latency for memory loads can be severely troublesome, because they are highly unpredictable and they block all other operations making it difficult to fill the delay with some other ops.

Pipeline hazards

If a control hazard arise, the entire pipeline content has to be flushed away, and a new instruction has to start from the beginning, then wasting a lot of CPU cycles.

This is one of the main cause of performance loss in modern superscalar-superpipeline-out-of-order CPUs. The logics for the runtime branch predictor occupies a large space on processor chips but it definitely is worth it.

Best branch predictors are as good as 95% of accuracy: nonetheless, the branch mis-prediction, or branch miss, determines a huge performance loss: typical values for penalty are 10-20 cycles!

That is because the longer the pipeline, the further in the future you have to scrutinize the flow, the more difficult it is and the larger will be the misprediction penalty.

Pipeline hazards

A very long pipeline, then, is not much more effective than a shorter one due to the real intrinsic nature of the codes that run on the CPUs.

The hazards we have just described are actually very common (very rarely a program is a stream of totally independent instructions with no jumps) and so the full exploitation of superscalarity+superpipelining is never reached.

Basically, that's the reason why we do not have 100-stage deep pipelines.

Loops optimizations: pipelining

As we have seen, modern processor may be able to “perform more than one operations at a time”, through super-pipelining operations.

As for what concerns the programmer’s perspective, codes must be written so to make the pipelines saturation as effective as possible.

```
for (int i = 0; i < N; i++)
    S += a[i] * b[i];
```

This way S is both read and written (the S as input in iteration i depends on S as output of iteration $i-1$) and the FP pipeline is difficult to be exploited (iteration i must unavoidably wait for iteration $i-1$ to end, which takes ~3-6 cycles at least)

Loops optimizations: pipelining

Make things easier for the compiler and the CPU

```
for (i = 0; i < N; i += 2) {  
    sum0 += a[i] * b[i];  
    sum1 += a[i+1] * b[i+1]; }
```

v. A

Unrolling make it easier for the CPU to saturate the pipelines.
Just try to separate load / multiply from addition:

```
for (i = 0; i < N; i += 2) {  
    double tmp0 = a[i] * b[i];  
    double tmp1 = a[i+1] * b[i+1];  
    sum0 += tmp0;  
    sum1 += tmp1; }
```

v. B

Loops optimizations: pipelining

Make things easier for the compiler and the CPU

More unrolling may work better, it depends on the CPU

```
for (i = 0; i < N; i += 4) {  
    sum0 += a[i] * b[i];  
    sum1 += a[i+1] * b[i+1];  
    sum2 += a[i+2] * b[i+2];  
    sum3 += a[i+3] * b[i+3]; }
```

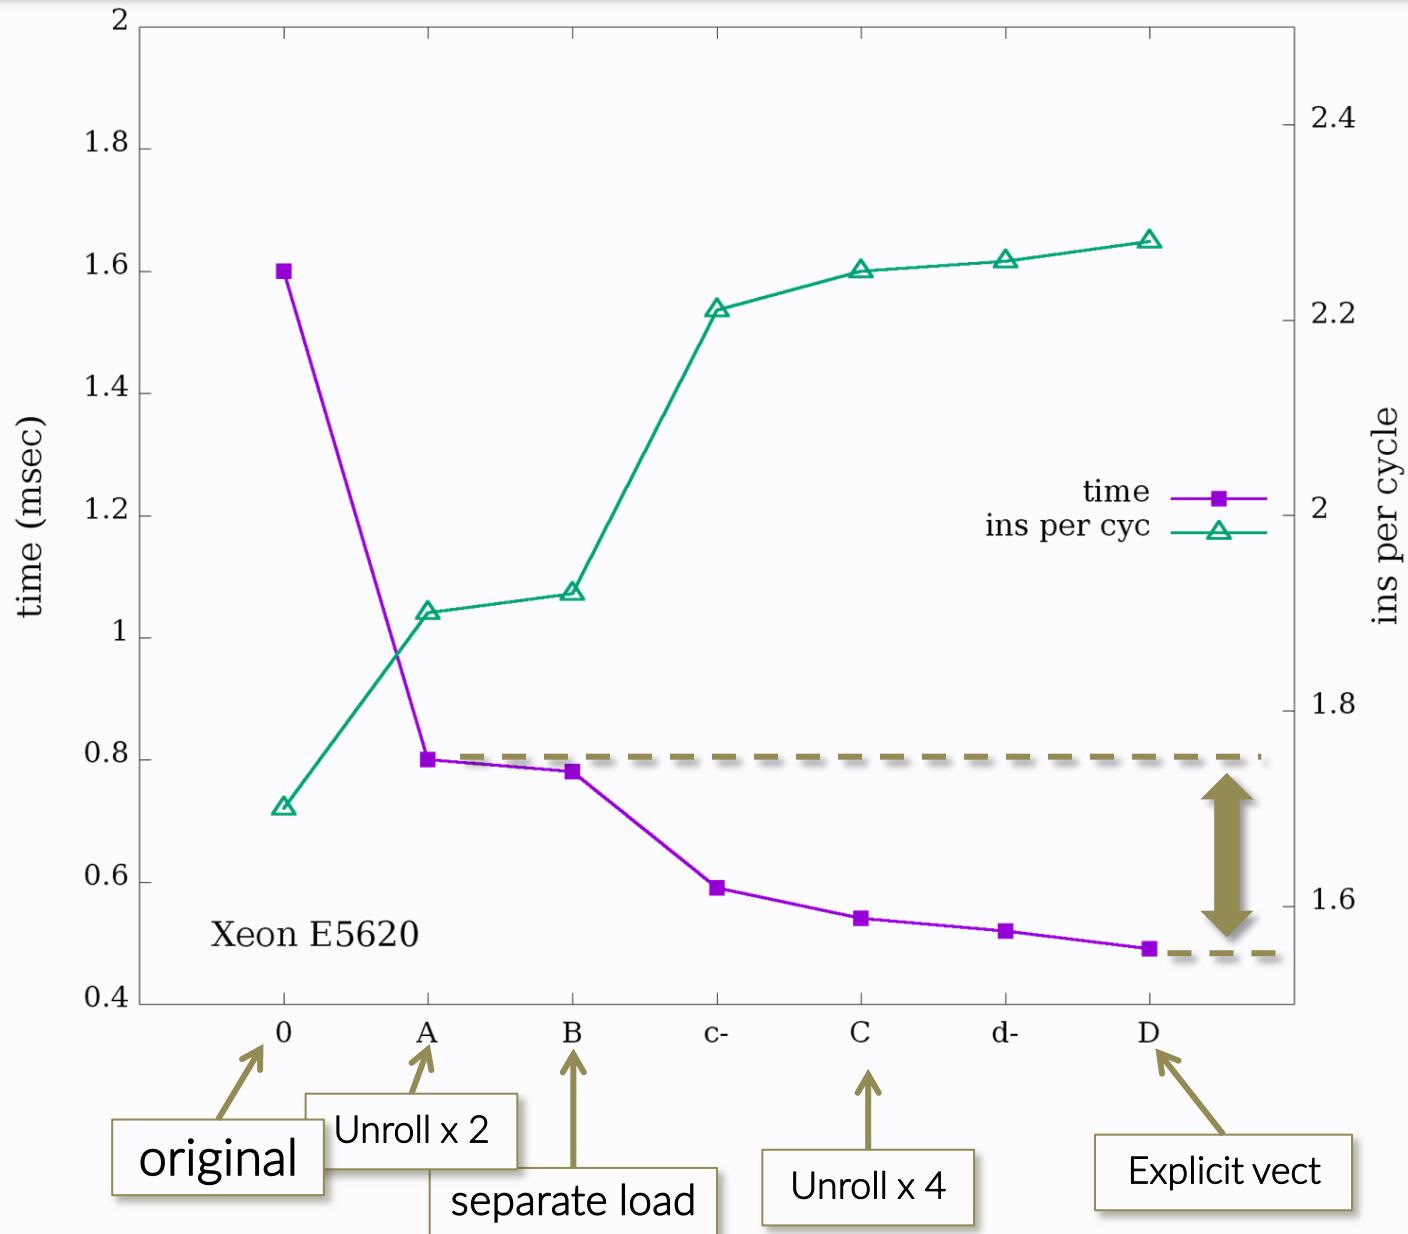
v. C

Or, eventually, let the explicit vectorization to step in

```
typedef double v4df __attribute__((vector_size (4*sizeof(double))));  
v4df array1, array2;  
v4df sum;  
for (i = 0; i < N/4; i++)  
    sum += array1[i] * array2[i];
```

v. D

Loops optimizations: pipelining



Pipeline: hands-on

Let us now look in more detail to what happens at assembler level, so to understand in deeper detail why some implementations are more effective than others.

We will use `-O0`, for the usual reason: with such simple kernels, compilers are very good (not all equally good..) in optimizing the code, and differences among improved version are more vagues.

...that **does NOT** relieve you of knowing how to write not-that-bad code..

Pipeline: hands-on

v0

```
for (int i = 0; i < N; i++)
    S += a[i] * b[i];
```

differences btw 2 compilers

v0

```
for (int i = 0; i < N; i++)
    S += a[i] * b[i];
```

```
.LB3365:
##      for ( int i = 0; i < N; i++ )
    cmpl    %ebx, %r14d
    jge     .LB3366
##  lineno: 185
..LN20:

    movslq  %r14d, %rax
    movq    -72(%rbp), %rcx
    vmovsd  (%rcx,%rax,8), %xmm0

    vmovsd  -32(%rbp), %xmm1

    vfmadd132sd    (%r12,%rax,8), %xmm1, %xmm0
    vmovsd  %xmm0, -32(%rbp)

    addl    $1, %r14d
    jmp     .LB3365
```

```
.L8:
# v0.c:185:    sum += array1[i] * array2[i];
    mov    eax, DWORD PTR -136[rbp]          # tmp158, i
    cdqe
    lea    rdx, 0[0+rax*8]
    mov    rax, QWORD PTR -96[rbp] # tmp159, array1
    add    rax, rdx
    vmovsd xmm1, QWORD PTR [rax]
    mov    eax, DWORD PTR -136[rbp]
    cdqe
    lea    rdx, 0[0+rax*8] # _23,
    mov    rax, QWORD PTR -88[rbp] # tmp161, array2
    add    rax, rdx
    vmovsd xmm0, QWORD PTR [rax]
    vmulsd xmm0, xmm1, xmm0
    vmovsd xmm1, QWORD PTR -104[rbp]          # tmp163, sum
    vaddsd xmm0, xmm1, xmm0
    vmovsd QWORD PTR -104[rbp], xmm0          # sum, tmp162
    inc    DWORD PTR -136[rbp]    # i
    # v0.c:184:    for ( int i = 0; i < N; i++ )
    mov    eax, DWORD PTR -136[rbp]          # tmp164, i
    cmp    eax, DWORD PTR -148[rbp]          # tmp164, N
    jl     .L8
```

differences btw 2 compilers

```
.L83365:
##    for ( int i = 0; i < N; i++ )
        cmpl    %rbx,%r14d
        jge     .L83366
##  line#: 185
..LN20:
        movslq  %r14d,%rax
        movq    -72(%rbp),%rcx
        vnmovsd (%rcx,%rax),%xmm0
        vfmadd132sd (%r12,%rax,8),%xmm1,%xmm0
        vnmovsd %xmm0,-32(%rbp)
        addl    $1,%r14d
        jmp     .L83365
```

Note: no optimization has been switched on in both cases

luca@GGG:~/code/HPC_lectures/pipeline/2018% pgi/v0_papi_00 50000000
generating 100000000 numbers..done
sum is 1.24993e+07
time is 0.176009 (min 0.175734, std dev 0.00054682, all 1.76643)
transfer rate was 4.23 GB/sec (28.34% of theoretical max that is 15 GB/sec)
IPC: 1
[time: 0.176sec ins: 5e+08]

```
.L8:
# v0.c:185: sum += array1[i] * array2[i];
    nov    eax,DWORD PTR -136(%rbp)      # tmp158, i
    cdqe
    lea    rdx,0[%rax*8]
    nov    eax,QWORD PTR -96(%rbp) # tmp159, array1
    add    rax,rdx
    vnmovsd xmm1,QWORD PTR [rax]
    nov    eax,DWORD PTR -136(%rbp)
    cdqe
    lea    rdx,0[%rax*8] # _23,
    nov    rax,QWORD PTR -88(%rbp) # tmp161, array2
    add    rax,rdx
    vnmovsd xmm0,xmm1,xmm0
    vnmulsd xmm0,xmm0,xmm0
    vnmovsd xmm0,QWORD PTR -104(%rbp) # tmp163, sum
    vaddsd xmm0,QWORD PTR -104(%rbp),xmm0 # sum, tmp162
    vnmovsd xmm0,QWORD PTR -104(%rbp),xmm0 # sum, tmp162
    inc    DWORD PTR -136(%bp),%i
    # v0.c:184: for ( int i = 0; i < N; i++ )
    nov    eax,DWORD PTR -136(%bp)
    cmp    eax,DWORD PTR -148(%rbp) # tmp164, N
    jl     .L8
```

luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v0_papi_00 50000000
generating 100000000 numbers..done
sum is 1.24993e+07
time is 0.190236 (min 0.189567, std dev 0.000538453, all 1.90762)
transfer rate was 3.92 GB/sec (26.22% of theoretical max that is 15 GB/sec)
IPC: 1.9
[time: 0.1902sec ins: 1.00158e+09]

Short comment on the previous slides

The PGI compiler has opted for using specialised fused multiply-addition instruction (`fmaadd`) that allows for a more compact code, but an IPC of 1 (which basically means that the code occupies 1 pipeline at a time).

The gcc compiler, instead, opts for a redundant code (the 2 subsequent blocks `mov | cdqe | lea | mov | add | vmosd`) that can involve 2 pipelines, and separate multiplication and addition.

Both compiler could be induced to generate different and more efficient code playing with their options (the default behaviour can differ among different compilers).

The result is more or less equivalent from the point of view of the run-time due to the fact that the code generated by gcc is longer (more instructions) but it is dispatched to more pipelines (larger IPC).

NOTE: the IPC is a good metric, but must be understood in a larger context. I.e. a larger IPC is good, but how good is the code depends on what instructions are being executed.

turning on the optimization...

Let's have a preview of how the generated code changes..

```
for (int i = 0; i < N; i++)
    S += a[i] * b[i];
```

-00

```
.L8:
# v0.c:185:    sum += array1[i] * array2[i];
    mov    eax, DWORD PTR -136[rbp]          # tmp158, i
    cdqe
    lea    rdx, 0[0+rax*8]
    mov    rax, QWORD PTR -96[rbp] # tmp159, array1
    add    rax, rdx
    vmovsd xmm1, QWORD PTR [rax]
    mov    eax, DWORD PTR -136[rbp]
    cdqe
    lea    rdx, 0[0+rax*8] # _23,
    mov    rax, QWORD PTR -88[rbp] # tmp161, array2
    add    rax, rdx
    vmovsd xmm0, QWORD PTR [rax]
    vmulsd xmm0, xmm1, xmm0
    vmovsd xmm1, QWORD PTR -104[rbp]          # tmp163,
    vaddsd xmm0, xmm1, xmm0
    vmovsd QWORD PTR -104[rbp], xmm0          # sum, tmp
    inc    DWORD PTR -136[rbp]      # i
# v0.c:184:    for ( int i = 0; i < N; i++ )
    mov    eax, DWORD PTR -136[rbp]          # tmp164,
    cmp    eax, DWORD PTR -148[rbp]          # tmp164,
    jl     .L8
```

-O3 -march=native

```
.L8:
# v0.c:55:    sum += array1[i] * array2[i];
    vmovupd ymm5, YMMWORD PTR 0[r13+rax]
    vmulpd  ymm0, ymm5, YMMWORD PTR [r15+rax]
    add    rax, 32 # ivtmp.18,
    vaddsd  xmm4, xmm0, xmm4
    vunpckhpd   xmm1, xmm0, xmm0
    vextractf128  xmm0, ymm0, 0x1
    vaddsd  xmm1, xmm1, xmm4
# v0.c:55:    sum += array1[i] * array2[i];
    vaddsd  xmm1, xmm0, xmm1
    vunpckhpd   xmm0, xmm0, xmm0
    vaddsd  xmm4, xmm1, xmm0
    cmp    rax, r12
    jne     .L8    #,
```

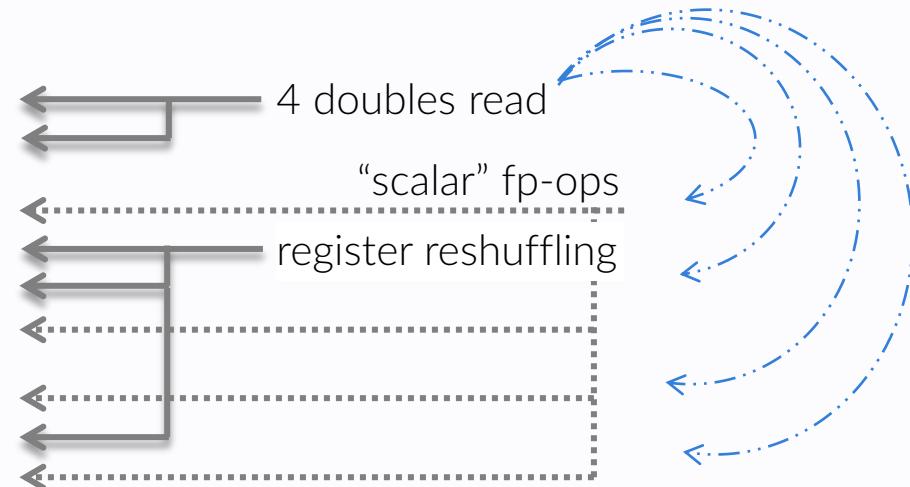
Let's have a preview of how the generated code changes..

v0

```
for (int i = 0; i < N; i++)
    S += a[i] * b[i];
```

```
.L8:
# v0.c:55:    sum += array1[i] * array2[i];
vmovupd    ymm5, YMMWORD PTR 0[r13+rax]
vmulpd    ymm0, ymm5, YMMWORD PTR [r15+rax]
add        rax, 32 # ivtmp.18,
vaddsd    xmm4, xmm0, xmm4
vunpckhpd  xmm1, xmm0, xmm0
vextractf128  xmm0, ymm0, 0x1
vaddsd    xmm1, xmm1, xmm4
# v0.c:55:    sum += array1[i] * array2[i];
vaddsd    xmm1, xmm0, xmm1
vunpckhpd  xmm0, xmm0, xmm0
vaddsd    xmm4, xmm1, xmm0
cmp        rax, r12
jne        .L8    #,
```

*gcc 8.2 on SkyLake
-O3 -march=native*



HINT:

data are fetched with vector instructions,
then processed in a mixed way

Short comment on the previous slide

Looking at the assembler generated by the compiler is often useful to understand where to direct optimization efforts.

In the previous slide, the compiler (with `-O3 -march=native`) chooses to fetch data from memory with a vector instruction, loading 4 doubles at a time (note that it uses the instruction for *unaligned memory*^(*): `vmovupd`).

But however can not really exploit the ILP (Instruction-Level parallelism) due to the data dependency intrinsic in how we wrote the loop.

That's why the best first step is to exhibit the possible parallelism, for instance either by unrolling or by a different accumulation (v1 and v3 in the following slides).

(*) you should remember what alignment is from the lecture about memory allocation

v* profile - 00

time is :0.217358 (min 0.216092, std dev 0.000667496)
transfer rate was 3.43 GB/sec (22.95% of theoretical max that is 15 GB/sec) v0

503,469,659	cycles..	#	0.37	insn per cycle	(+- 12.99%)
185,027,735	instructions..				(+- 17.44%)

time is :0.139794 (min 0.139124, std dev 0.000733397)
transfer rate was 5.33 GB/sec (35.68% of theoretical max that is 15 GB/sec) v1

394,225,210	cycles..	#	0.52	insn per cycle	(+- 9.57%)
204,046,712	instructions..				(+- 27.62%)

time is :0.119491 (min 0.117614, std dev 0.00116417)
transfer rate was 6.24 GB/sec (41.75% of theoretical max that is 15 GB/sec) v3

1,183,361,839	cycles..	#	0.70	insn per cycle	(+- 6.59%)
824,904,000	instructions..				(+- 8.66%)

time is :0.0805418 (min 0.0798115, std dev 0.00128884)
transfer rate was 9.25 GB/sec (61.93% of theoretical max that is 15 GB/sec) v6

288,211,800	cycles..	#	0.32	insn per cycle	(+- 1.80%)
92,540,296	instructions..				(+- 3.80%)

V* profile – 03

```
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v0_papi_03n 50000000
generating 100000000 numbers..done
sum is 1.24993e+07
time is :0.104058 (min 0.103508, std dev 0.000375562, all 1.04728)
transfer rate was 7.16 GB/sec ( 47.94% of theoretical max that is 15 GB/sec)
    IPC: 0.65
        [ time: 0.1041sec - ins: 1.50004e+08 ]
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v1_papi_03n 50000000
generating 100000000 numbers..done ( 1.27)
pipeline demonstrator, step 1:
- unroll 2 times

sum is 1.24993e+07
time is :0.0755591 (min 0.0750261, std dev 0.000291917, all 0.762576)
transfer rate was 9.86 GB/sec ( 66.02% of theoretical max that is 15 GB/sec)
    IPC: 1.2
        [ time: 0.07556sec - ins: 1.93753e+08 ]
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v3_papi_03n 50000000
generating 100000000 numbers..done (in 1.08sec)
pipeline demonstrator, step 3:

sum is 1.24993e+07
time is :0.0545231 (min 0.0544507, std dev 8.46002e-05, all 0.545244)
transfer rate was 13.7 GB/sec ( 91.49% of theoretical max that is 15 GB/sec)
    IPC: 1.1
        [ time: 0.0545sec - ins: 1.6e+08 ]%
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v3c_papi_03n 50000000
generating 100000000 numbers..done (in 1.27sec)
pipeline demonstrator, step 3:
- unroll 2 times +
- separate mul and sum
- separate accumulations

sum is 1.24993e+07
time is :0.0527944 (min 0.0526688, std dev 0.000107074, all 0.53314)
transfer rate was 14.1 GB/sec ( 94.49% of theoretical max that is 15 GB/sec)
    IPC: 1.2
        [ time: 0.0528sec - ins: 1.75e+08 ]
```

```
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v6_papi_03n 50000000
generating 100000000 numbers..done (in 1.27sec)
sum is 1.24993e+07
time is: 0.0647406 (min 0.0642672, std_dev 0.00017916, all 0.654759)
transfer rate was 11.5 GB/sec ( 77.05% of theoretical max that is 15 GB/sec)
    IPC: 0.8
        [ time: 0.06475sec - ins: 1.00003e+08 ]
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v6_intrinsics_papi_03n 50000000
generating 100000000 numbers..done (in 1.1sec)
sum is 1.24993e+07
time is: 0.0504149 (min 0.0468492, std_dev 0.00132376, all 0.509429)
transfer rate was 14.8 GB/sec ( 98.95% of theoretical max that is 15 GB/sec)
    IPC: 0.67
        [ time: 0.05042sec - ins: 7.50003e+07 ]
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v8_papi_03n 50000000
generating 100000000 numbers..done (in 1.11sec)
sum is 1.24993e+07
time is: 0.0499751 (min 0.0466076, std_dev 0.00124928, all 0.499767)
transfer rate was 14.9 GB/sec ( 99.82% of theoretical max that is 15 GB/sec)
    IPC: 0.56
        [ time: 0.04998sec - ins: 6.25002e+07 ]
```

There's a lot of improvement in such a simple kernel, but why IPC even decreased? Hint: the code becomes memory-bound (look at the memory transfer rate).

comparison: v0 – v1

v0

```
for ( int i = 0; i < N; i++ )
    sum += array1[i] * array2[i];
```

v1

```
for ( int i = 0; i < N-1; i+=2 )
    // simply unrolling 2 times, exposes the fact that at least
    // 2 elements of the array can be processed independently
{
    sum1 += array1[ i ] * array2[ i ];
    sum2 += array1[ i+1 ] * array2[ i+1 ];
}
if ( N % 2 )
    sum = array1[ N-1 ] * array2[ N-1 ];
```

comparison: v0 – v1 [-00]

v0

```
.L8:
# v0.c:185:    sum += array1[i] * array2[i];
    mov    eax, DWORD PTR -136[rbp]      # tmp158, i
    cdqe
    lea    rdx, 0[0+rax*8]
    mov    rax, QWORD PTR -96[rbp] # tmp159, array1
    add    rax, rdx
    vmovsd xmm1, QWORD PTR [rax]
    mov    eax, DWORD PTR -136[rbp]
    cdqe
    lea    rdx, 0[0+rax*8] # _23,
    mov    rax, QWORD PTR -88[rbp] # tmp161, array2
    add    rax, rdx
    vmovsd xmm0, QWORD PTR [rax]
    vmulsd xmm0, xmm1, xmm0
    vmovsd xmm1, QWORD PTR -104[rbp]      # tmp163, sum
    vaddsd xmm0, xmm1, xmm0
    vmovsd QWORD PTR -104[rbp], xmm0      # sum, tmp162
    inc    DWORD PTR -136[rbp]      # i
# v0.c:184:    for ( int i = 0; i < N; i++ )
    mov    eax, DWORD PTR -136[rbp]      # tmp164, i
    cmp    eax, DWORD PTR -148[rbp]      # tmp164, N
    jl     .L8
```

v1

```
.L10:
# v1_aligned.c:91:    sum1 += array1[ i ] * array2[ i ];
    movl   -148(%rbp), %eax      # i, tmp186
    cltq
    leaq   0(%rax,8), %rdx
    movq   -96(%rbp), %rax # array1, tmp187
    addq   %rdx, %rax
    vmovsd (%rax), %xmm1
    movl   -148(%rbp), %eax      # i, tmp188
    cltq
    leaq   0(%rax,8), %rdx
    movq   -88(%rbp), %rax # array2, tmp189
    addq   %rdx, %rax
    vmovsd (%rax), %xmm0
    vmulsd %xmm0, %xmm1, %xmm0
    vmovsd -112(%rbp), %xmm1      # sum1, tmp191
    vaddsd %xmm0, %xmm1, %xmm0
    vmovsd %xmm0, -112(%rbp)      # tmp190, sum1
    movl   -148(%rbp), %eax      # i, tmp192
    cltq
    incq   %rax
    leaq   0(%rax,8), %rdx
    movq   -96(%rbp), %rax # array1, tmp193
    addq   %rdx, %rax
    vmovsd (%rax), %xmm1
    movl   -148(%rbp), %eax      # i, tmp194
    cltq
    incq   %rax
    leaq   0(%rax,8), %rdx
    movq   -88(%rbp), %rax # array2, tmp195
    addq   %rdx, %rax
    vmovsd (%rax), %xmm0
    vmulsd %xmm0, %xmm1, %xmm0
    vmovsd -104(%rbp), %xmm1      # sum2, tmp197
    vaddsd %xmm0, %xmm1, %xmm0
    vmovsd %xmm0, -104(%rbp)
    addl   $2, -148(%rbp) #, i
.L9:
# v1_aligned.c:87:    for ( int i = 0; i < N-1; i+=2 )
    movl   -164(%rbp), %eax      # N, tmp198
    decl   %eax      # _42
    cmpl   %eax, -148(%rbp)      # i
    jl     .L10
```

comparison: v0 – v1 [-OO]

v0

```
.L8:
# v0.c:185:    sum += array1[i] * array2[i];
    mov    eax, DWORD PTR -136[rbp]          # tmp158, i
    cdqe
    lea    rdx, 0[0+rax*8]
    mov    rax, QWORD PTR -96[rbp] # tmp159, array1
    add    rax, rdx
    vmovsd xmm1, QWORD PTR [rax]
    mov    eax, DWORD PTR -136[rbp]
    cdqe
    lea    rdx, 0[0+rax*8] # _23,
    mov    rax, QWORD PTR -88[rbp] # tmp161, array2
    add    rax, rdx
    vmovsd xmm0, QWORD PTR [rax]
    vmulsd xmm0, xmm1, xmm0
    vmovsd xmm1, QWORD PTR -104[rbp]      # tmp163, sum
    vaddsd xmm0, xmm1, xmm0
    vmovsd QWORD PTR -104[rbp], xmm0      # sum, tmp162
    inc    DWORD PTR -136[rbp]          # i
    # v0.c:184:    for ( int i = 0; i < N; i++ )
    mov    eax, DWORD PTR -136[rbp]          # tmp164, i
    cmp    eax, DWORD PTR -148[rbp]          # tmp164, N
    jl     .L8
```

v1

```
.L10:
# v1_aligned.c:91:    sum1 += array1[ i ] * array2[ i ];
    movl  -148(%rbp), %eax          # i, tmp186
    cltq
    leaq  0(%rax,8), %rdx
    movq  -96(%rbp), %rax # array1, tmp187
    addq  %rdx, %rax
    vmovsd (%rax), %xmm1
    movl  -148(%rbp), %eax          # i, tmp188
    cltq
    leaq  0(%rax,8), %rdx
    movq  -88(%rbp), %rax # array2, tmp189
    addq  %rdx, %rax
    vmovsd (%rax), %xmm0
    vmulsd %xmm0, %xmm1, %xmm0
    vmovsd -112(%rbp), %xmm1      # sum1, tmp191
    vaddsd %xmm0, %xmm1, %xmm0
    vmovsd %xmm0, -112(%rbp)        # tmp190, sum1
    movl  -148(%rbp), %eax          # i, tmp192
    cltq
    incq
    leaq  0(%rax,8), %rdx
    movq  -96(%rbp), %rax # array1, tmp193
    addq  %rdx, %rax
    vmovsd (%rax), %xmm1
    movl  -148(%rbp), %eax          # i, tmp194
    cltq
    incq
    leaq  0(%rax,8), %rdx
    movq  -88(%rbp), %rax # array2, tmp195
    addq  %rdx, %rax
    vmovsd (%rax), %xmm0
    vmulsd %xmm0, %xmm1, %xmm0
    vmovsd -104(%rbp), %xmm1      # sum2, tmp197
    vaddsd %xmm0, %xmm1, %xmm0
    vmovsd %xmm0, -104(%rbp)
    addl  $2, -148(%rbp) #, i
.L9:
# v1_aligned.c:87:    for ( int i = 0; i < N-1; i+=2 )
    movl  -164(%rbp), %eax          # N, tmp198
    decl  %eax          # _42
    cmpl  %eax, -148(%rbp)        # i
    jl     .L10
```

The v1 version looks twice as longer (i.e.: more instructions to be executed), actually it really looks like the same code was copied & pasted below the original segment.

How does it come that the result is a faster execution?

comparison: v0 – v1 [-00]

It is clearer considering the output of the PAPI instrumented code, that returns the IPC.

Clearly, the 2fold unroll makes the CPU able to better exploit more than 1 logical units at a time.

```
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v0_papi_00 50000000
generating 100000000 numbers..done
sum is 1.24993e+07
time is :0.191208 (min 0.190616, std dev 0.00053432, all 1.91787)
transfer rate was 3.9 GB/sec ( 26.09% of theoretical max that is 15 GB/sec)
    IPC: 1.9
    [ time: 0.1912sec - ins: 1.00158e+09 ]
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v1_papi_00 50000000
generating 100000000 numbers..done ( 0.965)
pipeline demonstrator, step 1:
- unroll 2 times

sum is 1.24993e+07
time is :0.125749 (min 0.12396, std dev 0.00142263, all 1.26349)
transfer rate was 5.92 GB/sec ( 39.67% of theoretical max that is 15 GB/sec)
    IPC: 2.8
    [ time: 0.1258sec - ins: 9.76578e+08 ]
```

v1

```
for ( int i = 0; i < N-1; i+=2 )
    // simply unrolling 2 times, exposes the fact that at least
    // 2 elements of the array can be processed independently
{
    sum1 += array1[ i ] * array2[ i ];
    sum2 += array1[ i+1 ] * array2[ i+1 ];
}
if ( N % 2 )
    sum = array1[ N-1 ] * array2[ N-1 ];
```

v1b

```
double register sum0 = 0;
double register sum1 = 0;

double register * restrict a1 = __builtin_assume_aligned(array1, 32);
double register * restrict a2 = __builtin_assume_aligned(array2, 32);

tstart = CPU_TIME;
#pragma ivdep
for( int i = 0 ; i < N-1; i+=2 )
{
    sum0 += *(a1++) * *(a2++);
    sum1 += *(a1++) * *(a2++);
}
if( N%2 )
    sum = *a1 * *a2;
```

refining v1

```
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v1_papi_00 500000000
generating 100000000 numbers..done ( 0.973)
pipeline demonstrator, step 1:
- unroll 2 times

sum is 1.24993e+07
time is 0.125593 (min 0.125287, std dev 0.000476174, all 1.26231)
transfer rate was 5.93 GB/sec ( 39.72% of theoretical max that is 15 GB/sec)
    IPC: 2.8
    [ time: 0.1256sec - ins: 9.76578e+08 ]
```

```
luca@GGG:~/code/HPC_lectures/pipeline/2018% ./v1b_papi_00 500000000
generating 100000000 numbers.. done (in 0.973sec)
pipeline demonstrator, step 1b:
- unroll 2 times +
- align memory +
- minimize ptr arithmetic

sum is 1.24993e+07
time is 0.0977905 (min 0.0972945, std dev 0.000712555, all 0.977921)
transfer rate was 7.62 GB/sec ( 51.01% of theoretical max that is 15 GB/sec)
    IPC: 2.1
    [ time: 0.0978sec - ins: 5.76578e+08 ]
```

Let's go further in trying to make it clear for the CPU about the possible ILP

```
#pragma ivdep
for ( int i = 0; i < N_4; i+=4 )
{
    sum += array1[i] * array2[i] +
        array1[i+1] * array2[i+1] +
        array1[i+2] * array2[i+2] +
        array1[i+3] * array2[i+3];
}

for ( int i = N_4; i < N; i++ )
    sum += array1[i] * array2[i];
```

v3

prefetching

```
double register a;
double register b;
a = array1[0] * array2[0];
#pragma ivdep
for ( int i = 0; i < N_4; i+=4 )
{
    DO_NOT_OPTIMIZE;
    b = array1[i+4] * array2[i+4];
    DO_NOT_OPTIMIZE;
    sum += a +
        array1[i+1] * array2[i+1] +
        array1[i+2] * array2[i+2] +
        array1[i+3] * array2[i+3];
    a = b;
}

for ( int i = N_4; i < N; i++ )
    sum += array1[i] * array2[i];
```

v3b

refining: vectorization

```
typedef double v4df __attribute__ ((vector_size (4*sizeof(double))));  
typedef union {  
    v4df V;  
    double v[4];  
}v4df_u;  
  
#ifdef __GNUC_SOURCE  
    v4df sum_ = {0, 0, 0, 0};  
#else  
    v4df sum_ = {0};  
#endif  
v4df register mytmp;  
v4df register tmp = *((v4df*)&array1[0]) * *((v4df*)&array2[0]);  
  
int N_4 = N/4;  
int N_4 = N_4*4;  
tstart = CPU_TIME;  
#pragma ivdep  
for( int i = 1; i <= N_4; i++)  
{  
    _DO_NOT_OPTIMIZE_BEGIN;  
    mytmp = *((v4df*)array1 + i) * *((v4df*)array2 + i);  
    _DO_NOT_OPTIMIZE_END;  
    sum_ += tmp;  
    tmp = mytmp;  
}  
  
for ( int i = N_4; i < N; i++ )  
    sum += array1[ i ] * array2[ i ];
```

refining: vectorization with intrinsics

```
#elif defined ( __AVX__ ) || defined ( __AVX2__ )

#define V_DSIZE 4
typedef __m256d vd;
vd _dzero_ = {0, 0, 0, 0};
#define MUL_ADD( A1, A2, R ) _mm256_fmadd_pd( (A1), (A2), (R) )
```

```
vd sum_ __attribute__ ((aligned(64)));
sum_ = _dzero_;

vd * restrict A1 __attribute__ ((aligned(64)));
vd * restrict A2 __attribute__ ((aligned(64)));
A1 = (vd*)array1 ;
A2 = (vd*)array2 ;

int N__ = N / V_DSIZE;
int N_ = N__ * V_DSIZE;

tstart = CPU_TIME;

#pragma ivdep
for( int i = 0; i < N__; i++)
    sum_ = MUL_ADD( A1[i], A2[i], sum_ );

for ( int i = N_; i < N; i++)
    sum += array1[ i ] * array2[ i ];
```

Pipeline: hands-on

Codes

You find the code snippets presented here on our GitHub, with some comments about compilation.

Compile and run them with different options (and possibly different compilers) and try to understand what happens on your laptop and/or on Ulisse facility.

Outline

- Memory optimizations
- Loops optimizations
- Other optimizations
 - > prefetching
 - > macros as templates
 - > SIMD instructions
 - > access locality in NUMA systems

Prefetching

Prefetching

Modern CPUs have the capability of pre-emptively bring from memory into cache levels data that **will be needed shortly afterwards**.

They can do that following some speculative algorithm based on the current execution flow and assuming spatial locality and temporal locality.

Both *data* and *instructions* can be pre-fetched.

Pre-fetching may be both hardware-based and software-based (tipically the compiler insert pre-fetching instructions at compile-time).

Prefetching

From the point of view of the programmer, there are 2 possible ways to deal with prefetching

- **Explicit:** you explicitly insert a pre-fetching directive
very difficult to be achieved effectively: the directive must be inserted timely but not too early (data eviction) or too late (load latency)
- **Induced:** you consciously arrange data layout and execution flow so that to make it obvious to the compiler what to prefetch

Prefetching

Explicit prefetching: a simple example

```
int mybsearch(int *data, int N, int Key)
{
    int register low = 0;
    int register high = N;
    int register mid;

    while(low <= high) {
        mid = (low + high) / 2;

        if(data[mid] < Key)
            low = mid + 1;
        else if(data[mid] > Key)
            high = mid-1;
        else
            return mid;
    }

    return -1;
}
```

Prefetching

Explicit prefetching: a simple example

```
int mybsearch(int *data, int N, int Key)
{
    int register low = 0;
    int register high = N;
    int register mid;

    while(low <= high) {
        mid = (low + high) / 2;
        __builtin_prefetch (&data[(mid + 1 + high)/2], 0, 3);
        __builtin_prefetch (&data[(low + mid - 1)/2], 0, 3);
        if(data[mid] < Key)
            low = mid + 1;
        else if(data[mid] > Key)
            high = mid-1;
        else
            return mid;
    }

    return -1;
}
```

Prefetching

Explicit prefetching: a simple example

```
int mysearch(int *data, int N, int Key)
{
    int register low = 0;
```

```
luca@GGG:~/code/HPC_LECTURES/prefetching% ./prefetching_off
performing 13421772 lookups on 134217728 data..
```

set-up data.. set-up lookups..

start cycle.. time elapsed: 20.7534

```
luca@GGG:~/code/HPC_LECTURES/prefetching% ./prefetching_on
performing 13421772 lookups on 134217728 data with prefetching enabled..
```

set-up data.. set-up lookups..

start cycle.. time elapsed: 12.6204

```
    else if(data[mid] < Key)
        high = mid-1;
    else
        return mid;
}

return -1;
}
```

Prefetching

Explicit prefetching: a simple example

```
Samples: 71K of event 'cpu/mem-loads,ldlat=30/P', Event count (approx.): 13901140
Overhead      Samples  Memory access
 71,08%        42196  Local RAM hit
 24,14%        17022  LFB hit
  4,11%        10967  L3 hit
  0,63%         1714   L1 hit
  0,02%          75    L2 hit
  0,01%          15    L3 miss
  0,00%           1    Uncached hit
```

Read perf-report man page and Linux Man Page perf-report

```
__builtin_prefetch (&data[(mid + 1 + high)/2], 0, 3);
__builtin_prefetch (&data[(low + mid - 1)/2], 0, 3);
```

```
Samples: 61K of event 'cpu/mem-loads,ldlat=30/P', Event count (approx.): 11720387
```

```
Overhead      Samples  Memory access
 68,74%        29450  LFB hit
 27,04%        28208  L1 hit
  2,72%         909   Local RAM hit
  1,29%        2983   L3 hit
  0,20%         346   L2 hit
```

```
    }  
    return -1;
```

Prefetching



Usage of direct prefetching directive is highly uncertain, since it is difficult to spot the exact point – both in the code and in the execution – where to place them.

Moreover, the “exact point” is very likely dependent on the system you run on, and then it is susceptible to change significantly.

It is normally much safer to re-organize your code so to have **prefetching by pre-loading**.

Prefetching

Prefetching by preloading

```
elem a = elements[0]
for ( i = 0; i < 4*N_4; i+= 4 )
{
    elem e = elem[i+4]; // non-blocking miss
    elem b = elem[i+1]; // cache-hit
    elem c = elem[i+2]; // cache-hit
    elem d = elem[i+3]; // cache-hit
    Elaborate(a);
    Elaborate(b);
    Elaborate(c);
    Elaborate(d);
    a = e;
}
```

Prefetching: hands-on

A basic example for practicing with preloading

```
for ( i = 0; i < N; i++ )  
    sum += array[ i ];
```

You find code snippets with different flavours of prefetching-by-preloading technique on our GitHub, with some comments about compilation.

Compile and run them with different options (and possibly different compilers) and try to understand what happens on your laptop and/or on Ulisse facility.

Macros as templates

MACROS as templates

Standard **libc** has a lot of very good well optimized routines for many general purposes.

However, many routines are meant to work on a general structure without any precise knowledge of its internals.
Typically you call a **libc** routine like that;

```
libc_routine( void* base, size_t size, size_t n,  
              int (*)dealer(void *, ...) );
```

However, in this way the compiler can not switch on several optimizations because the data structure is opaque.

MACROS as templates

A possible cure for this is (in C) to write a MACRO that behaves like a template.

For instance, if you put the following code in a .h file

```
int MY_WONDER(A, B, C, ...) { \
    MY_DATA_TYPE doing something
    ..some stuff here..
    DEALER(A, B, C, ...)
    ..some other stuff.. }
```

MACROS as templates

And you define MY_DATA_TYPE and DEALER and whatever else you need just before including the .h file

```
#define MY_DATA_TYPE some_type
```

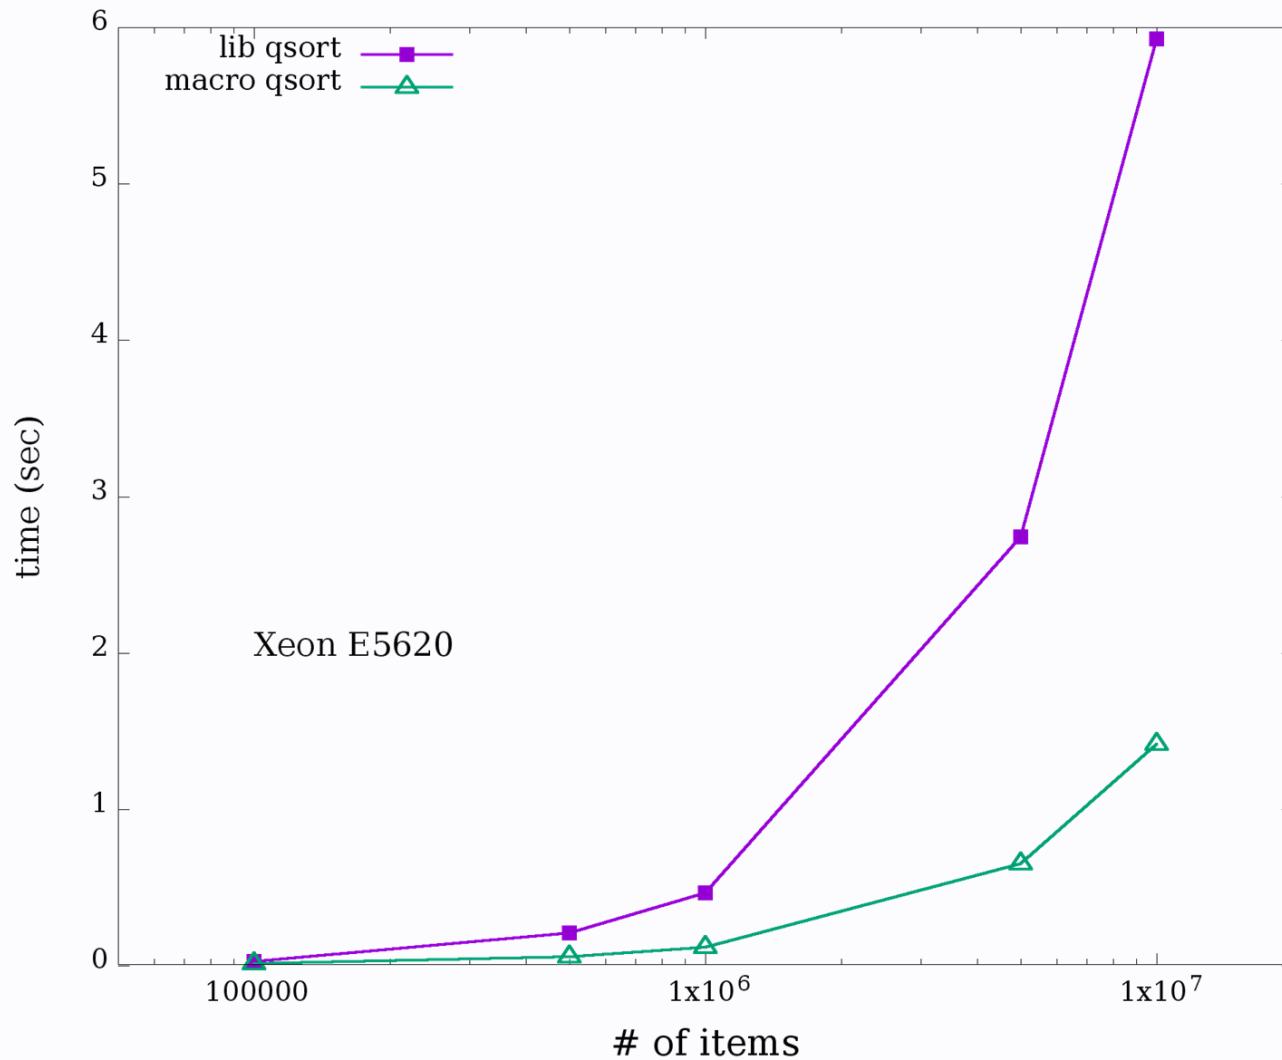
```
#define DEALER(A, B, C, ...) { \  
    ..some other stuff.. }
```

```
#include "my_routines.h"
```

You get something that resembles a template that the compiler can optimize.

MACROS as templates

Going practical, let's try a quicksort (live):



Access locality in NUMA sys

Access locality in ccNUMA systems

Let's start from a code snippet of the classic triad benchmark:

```
// initialization of arrays  
B[i] = ... ; C[i] = ...; D[i] = ...;  
  
#pragma omp parallel for  
for (cc = 0; cc < SIZE; cc++)  
    A[i] = B[i] + C[i] × D[i];
```

The initialization happens in a non-parallel region, while the loop where threads actually access data is in the **omp** region.

That means that data will be placed in the memory/cache of the LD0 (**first-touch policy**) and threads from other LDs will all have to queue and wait – very likely determining severe cache inefficiencies – to access the memory on LD0.

Access locality in ccNUMA systems

A good policy is then to parallelize the data initialization section as well:

```
#pragma omp parallel
B[i] = ... ; C[i] = ...; D[i] = ...;

#pragma omp parallel for
for (cc = 0; cc < SIZE; cc++)
    A[i] = B[i] + C[i] × D[i];
```

or, more in general, to initialize/read the data first-touching them in the same way they will be accessed.

[**tech note:** all that is true if the array is large enough so that each thread chunk is bigger than a page size. For smaller data set, consider to make them thread-private, creating local copies for each thread]

The winding road towards optimization

