FHPC/P1.2 course:

Lecture 1: Introduction to HPC (second part)

Stefano Cozzini

CNR/IOM and eXact-lab srl

MHPC
Master in High Performance Computing

Agenda of the previous lecture

- Prologue: why and where HPC ?
- What is HPC ?
  - Definitions&metrics
- Component of a HPC infrastructure
- HPC Concepts
  - Parallel programming paradigms
  - Evolution of paradigms
  - Ahmdal law / Gustafson law
  - Strong/weak scalability
- HOMEWORK&LABS

SISSA ICTP

MHPC
Master in High Performance Computing

Agenda: for today

- What is HPC infrastructure ?
  - Supercomputers & HPC Cluster
  - CPUs and Accelerators
  - Network/storage
- Software stack for HPC
  - Middleware: queue systems
  - Libraries/ Compiler/ performance Tools

SISSA ICTP

HPC Cluster Computer Architecture

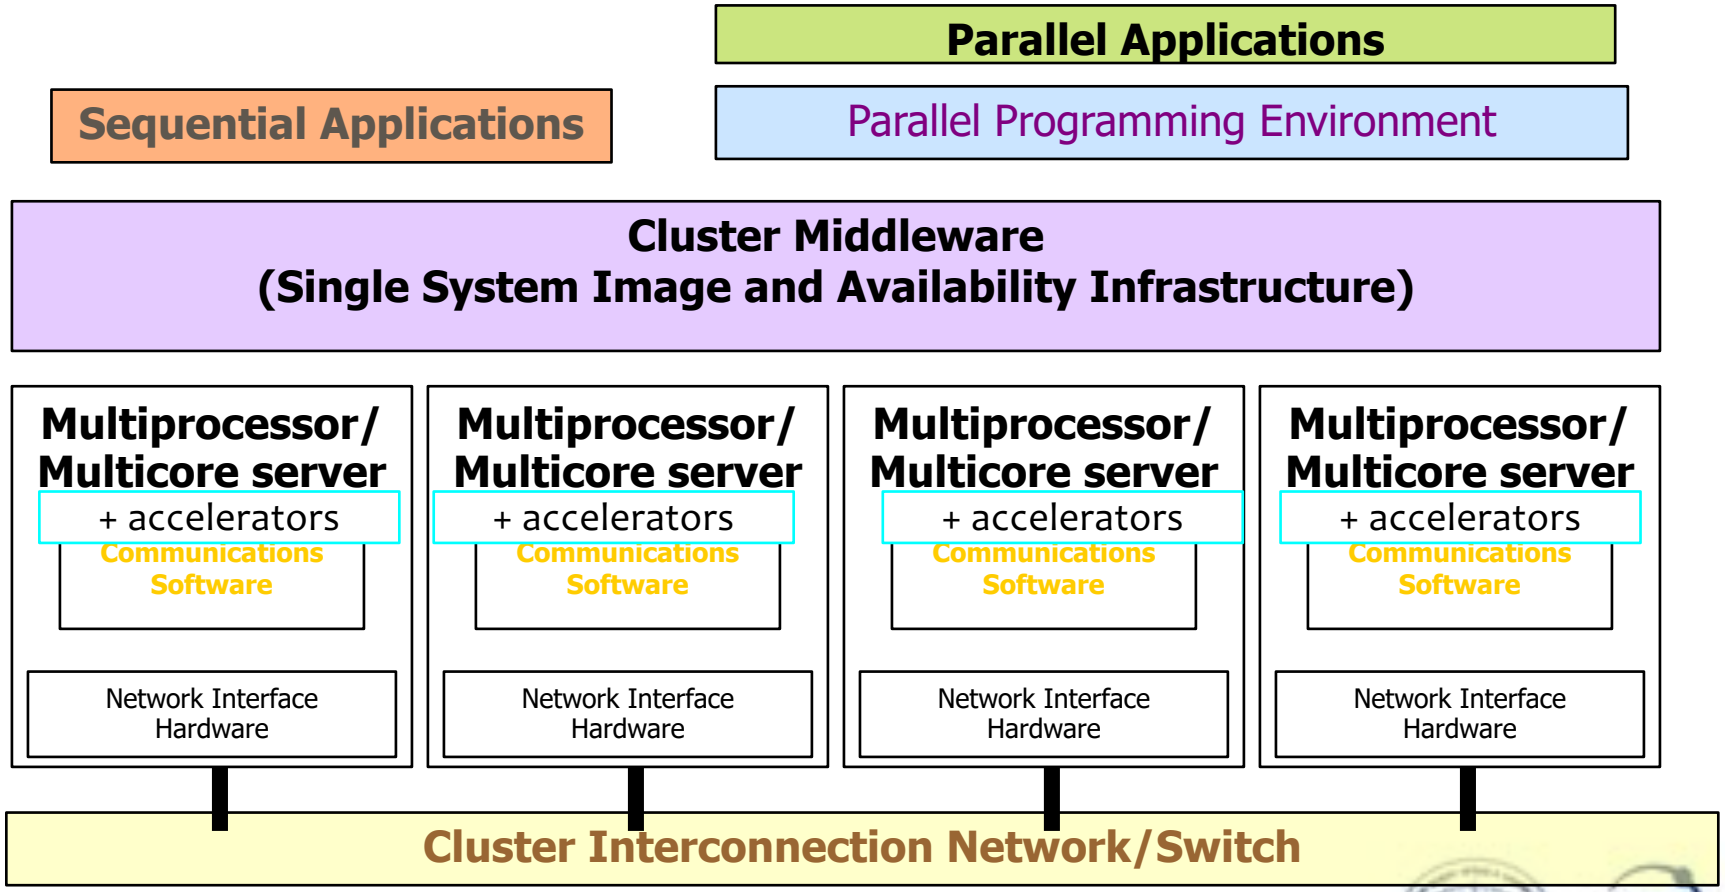

MHPC
Master in High Performance Computing

A modern node picture (Ullysses node)

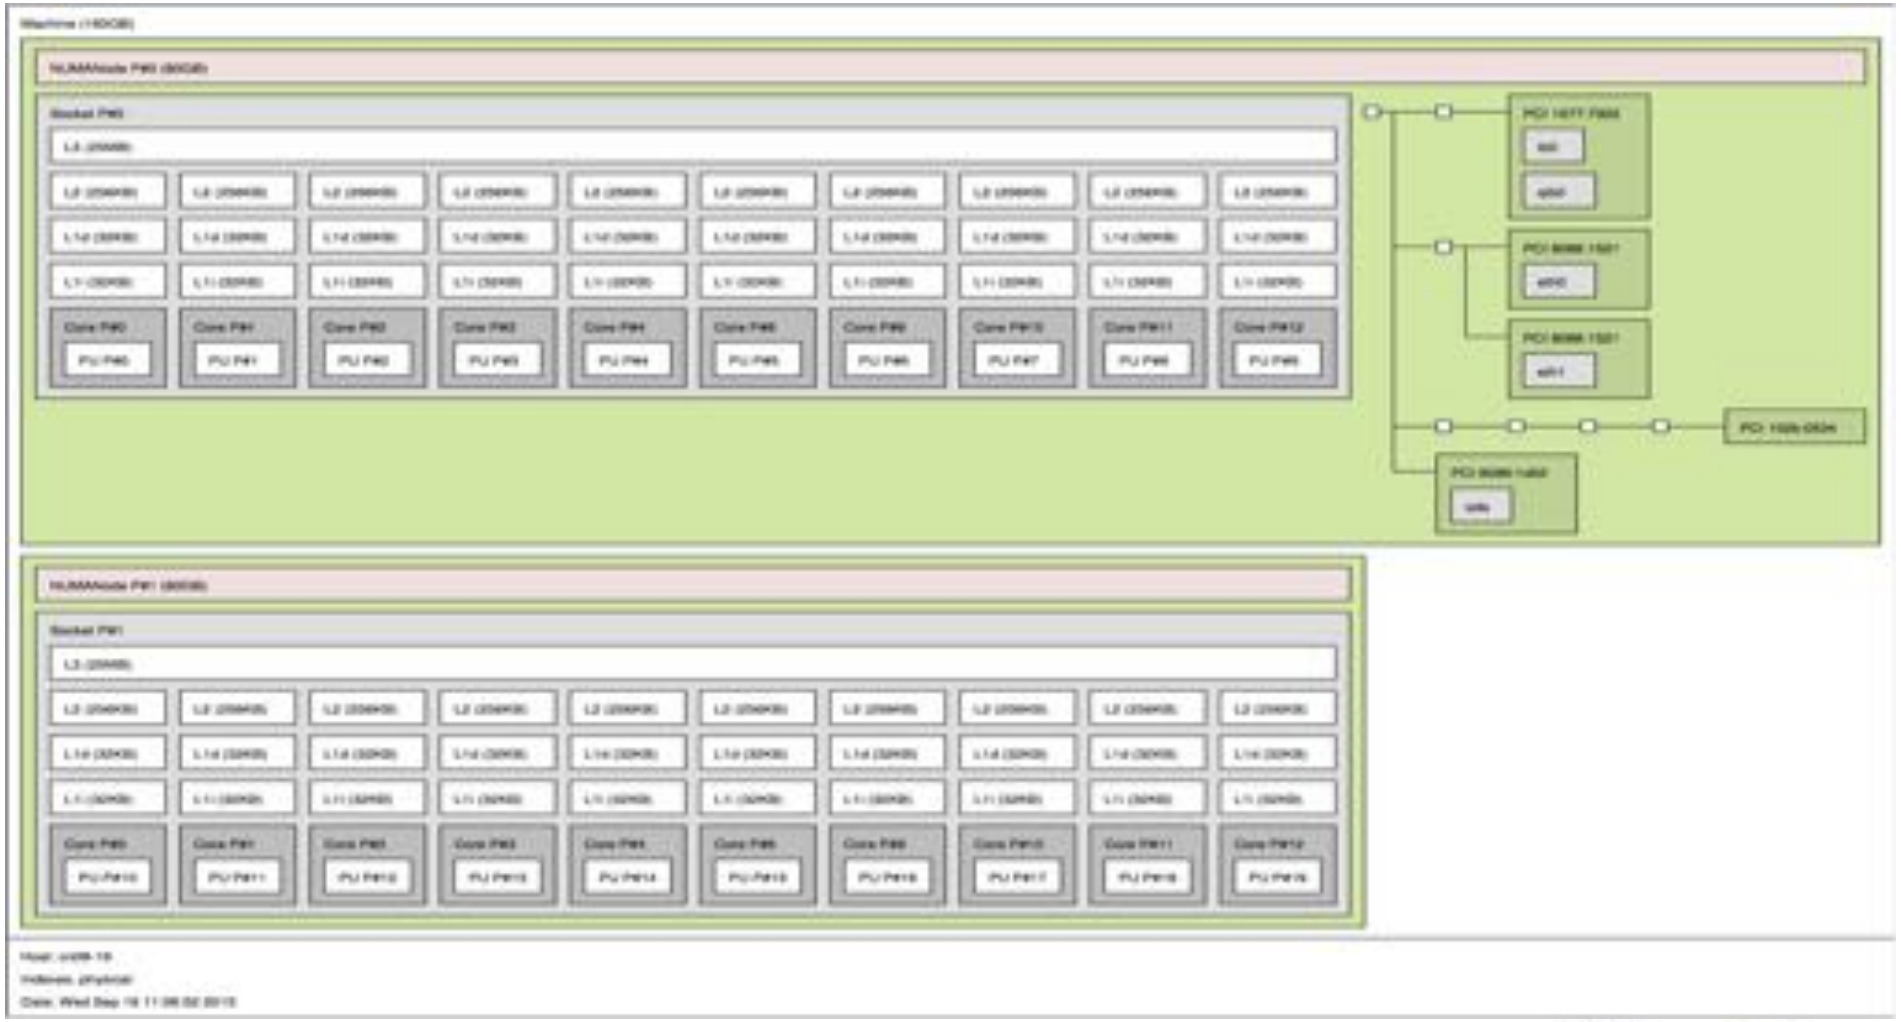

MHPC
Master in High Performance Computing

About HPC jargon

- Multiprocessor = server with more than 1 CPU
- Multicore= a CPU with more than 1 core

Processor = CPU =socket

BUT SOMETIME:

Processor= core

a process for each processor (i.e. each core)

SISSA ICTP

MHPC
Master in High Performance Computing

Elements of the clusters

- Several computers, nodes, often in special cases for easy mounting in a rack
- One or more networks (interconnects) to hook the nodes together
- Storage facilities.

SISSA ICTP

A node of modern HPC cluster

1U box

1 or 2 accelerators

A shared memory machine (SMP or NUMA)

MHPC
Master in High Performance Computing

Some times also blades.. In racks

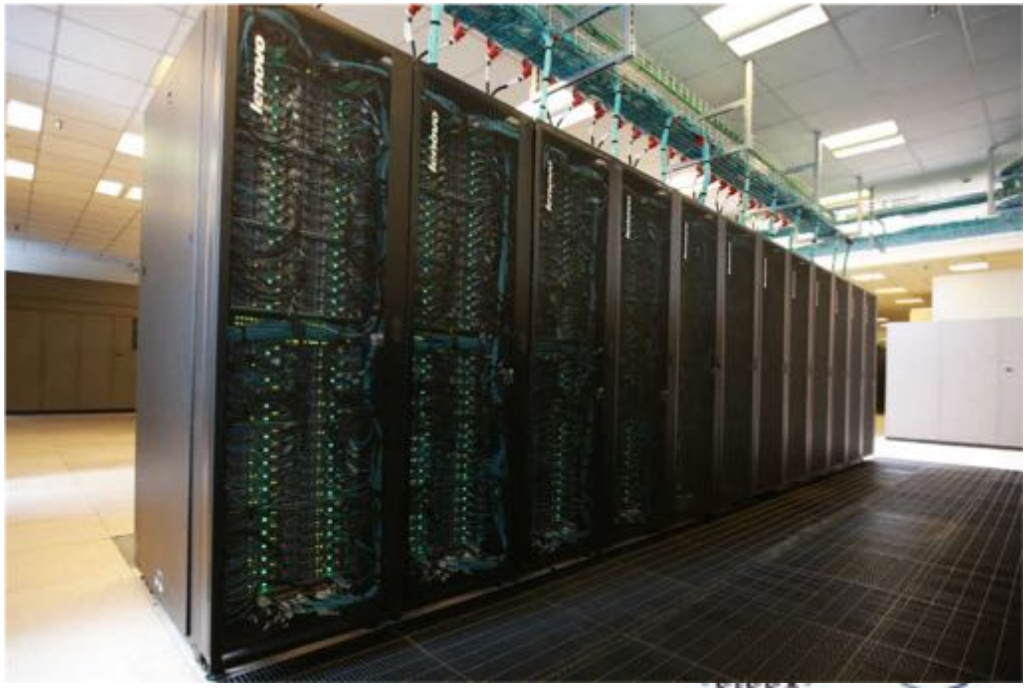

MHPC
Master in High Performance Computing

Building blocks of a cluster

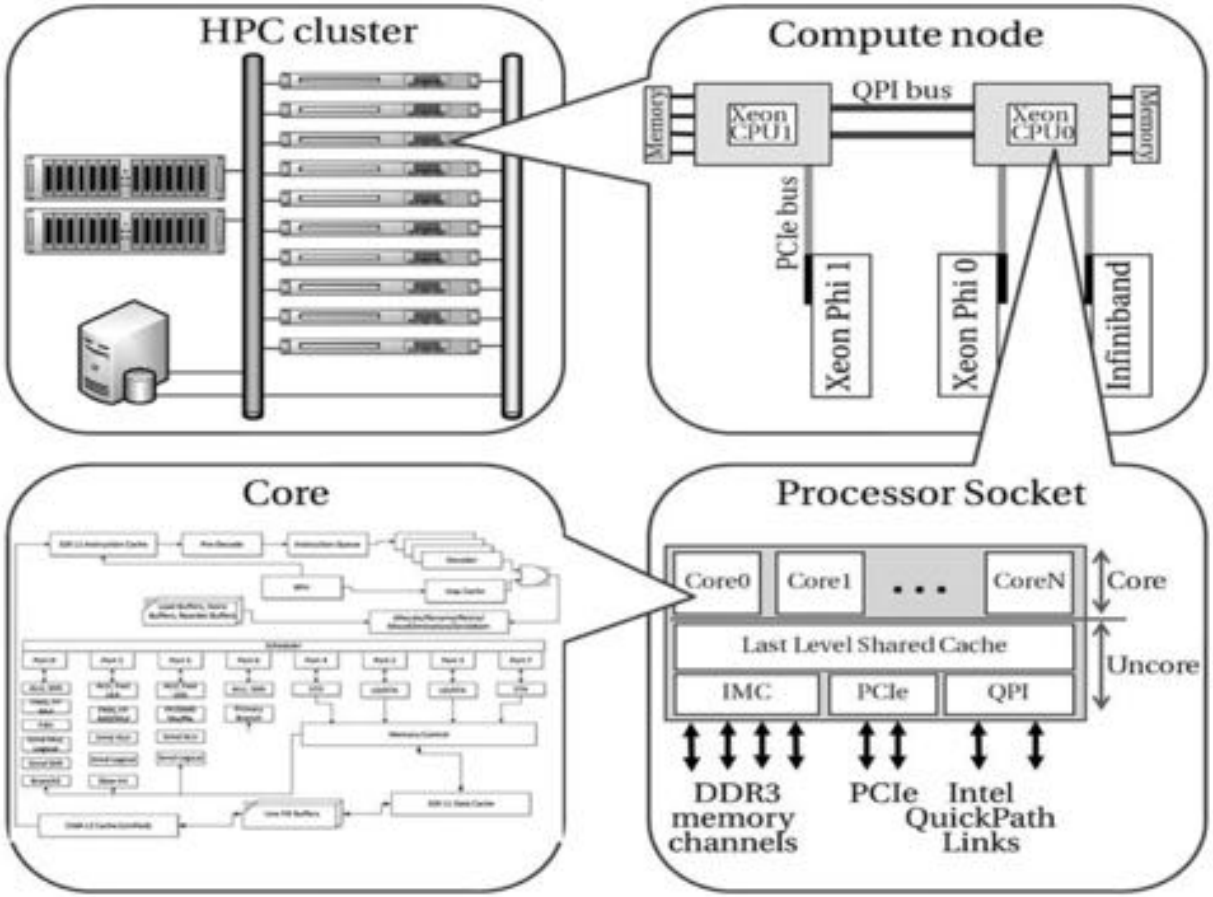

SISSA
ICTP

The motherboard components

Dual socket CPU

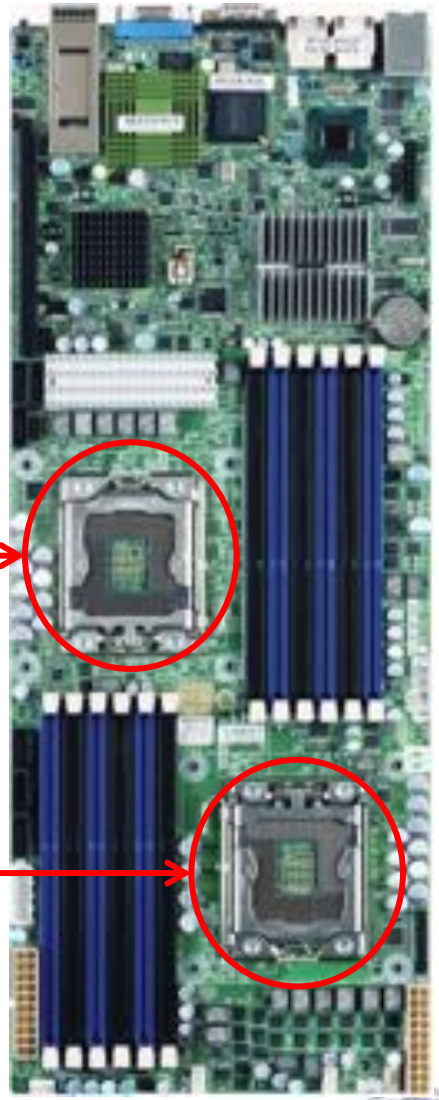

The motherboard components

DDR3 Ram 12 slots

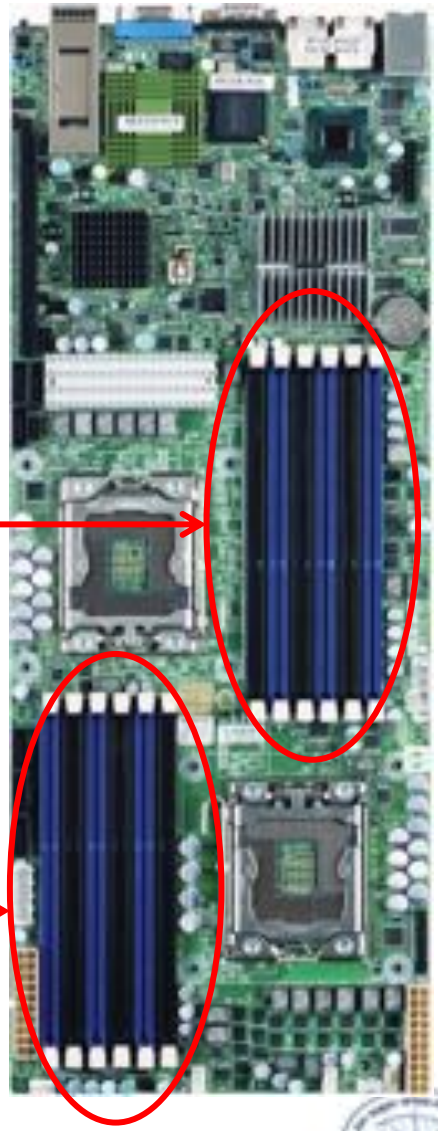

MHPC
Master in High Performance Computing

The motherboard components

Northbridge

SISSA ICTP

The motherboard components

PCI 8x 1 slot

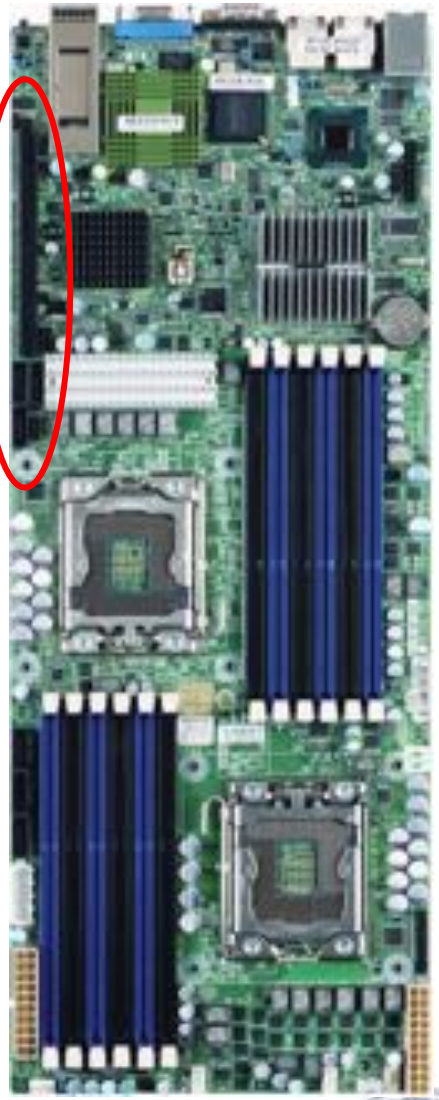

MHPC
Master in High Performance Computing

The motherboard components

Gb Ethernet dual port

SISSA ICTP

MHPC
Master in High Performance Computing

The motherboard components

Infiniband 4x QDR

SISSA ICTP

MHPC
Master in High Performance Computing

Motherboard components

Gb Ethernet dual port

PCI 8x 1 slot

Infiniband 4x QDR

Dual socket CPU

DDR3 Ram 12 slots

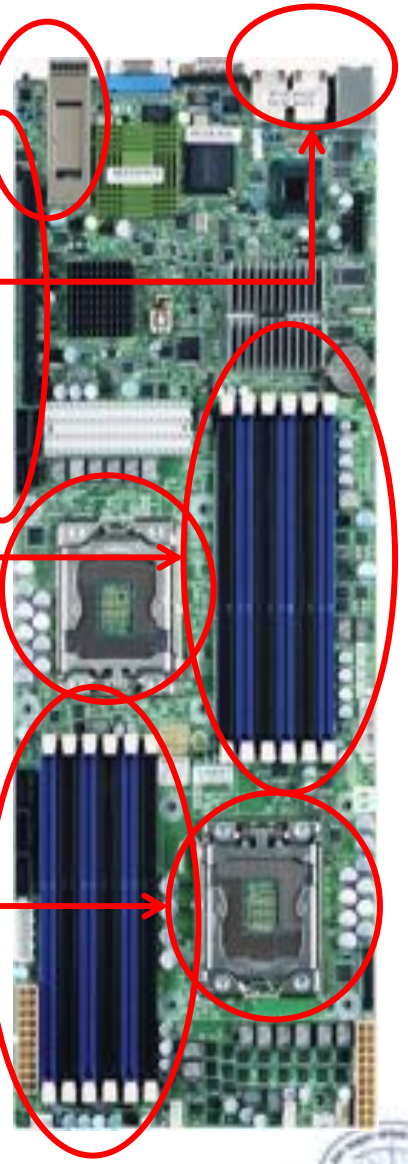

Slides from P.Altoe E4company

SISSA ICTP

MHPC
Master in High Performance Computing

Which CPUs on TOP500 system ?

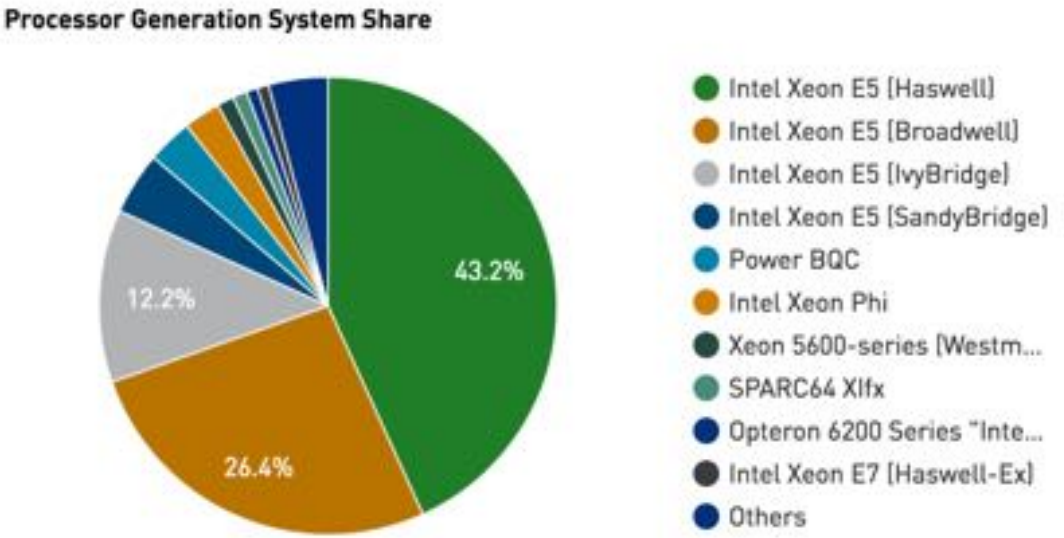

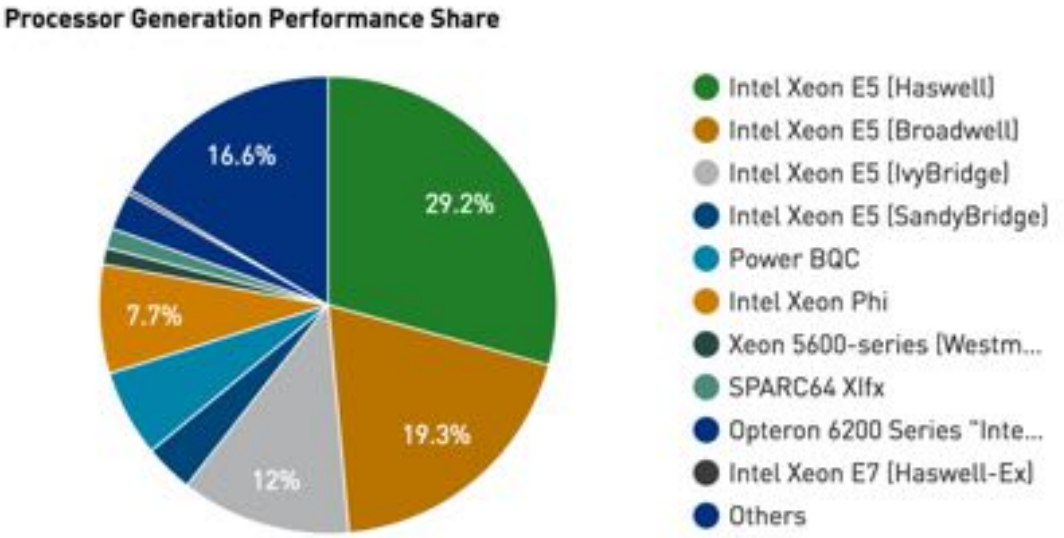

About network for cluster

- The performance of the network cannot be ignored
  - Latency: Initialization time before data can be sent
  - Per-link Peak Bandwidth: Maximum data transmission rate (varies with packet size)
  - Topology: how the network is done.

Clustering topology

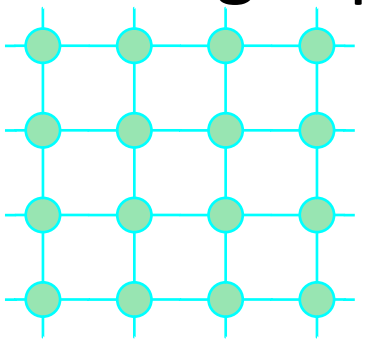

2D Mesh

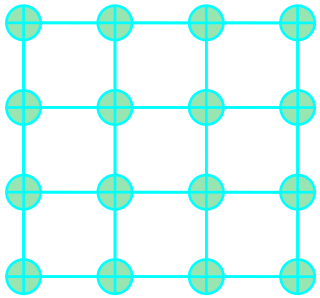

2D Torus

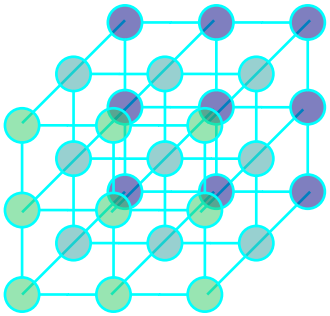

3D Mesh

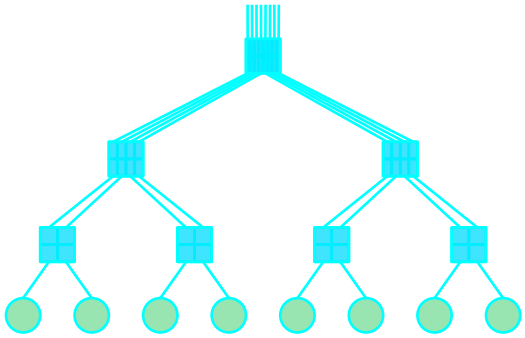

FAT TREE

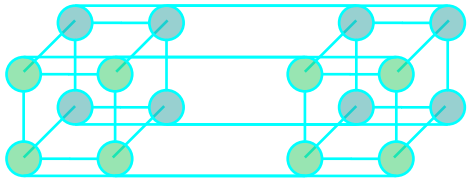

Hypercube (4-cube)

Latency&bandwidth

NETWORK	Latency	Bandwidth (GB/sec)
Gigabit	70-40	~ 0.125
10G	<5	~1.250
Infiniband 4DDR	~1.5/1.9	~ 3.2
Infiniband FDR	<1.0	~ 5

What is the UNIT OF MEASURE OF LATENCY ?

Microseconds: 3 order of magnitude larger than unit of measure of FP operations

HPC cluster logical structure..

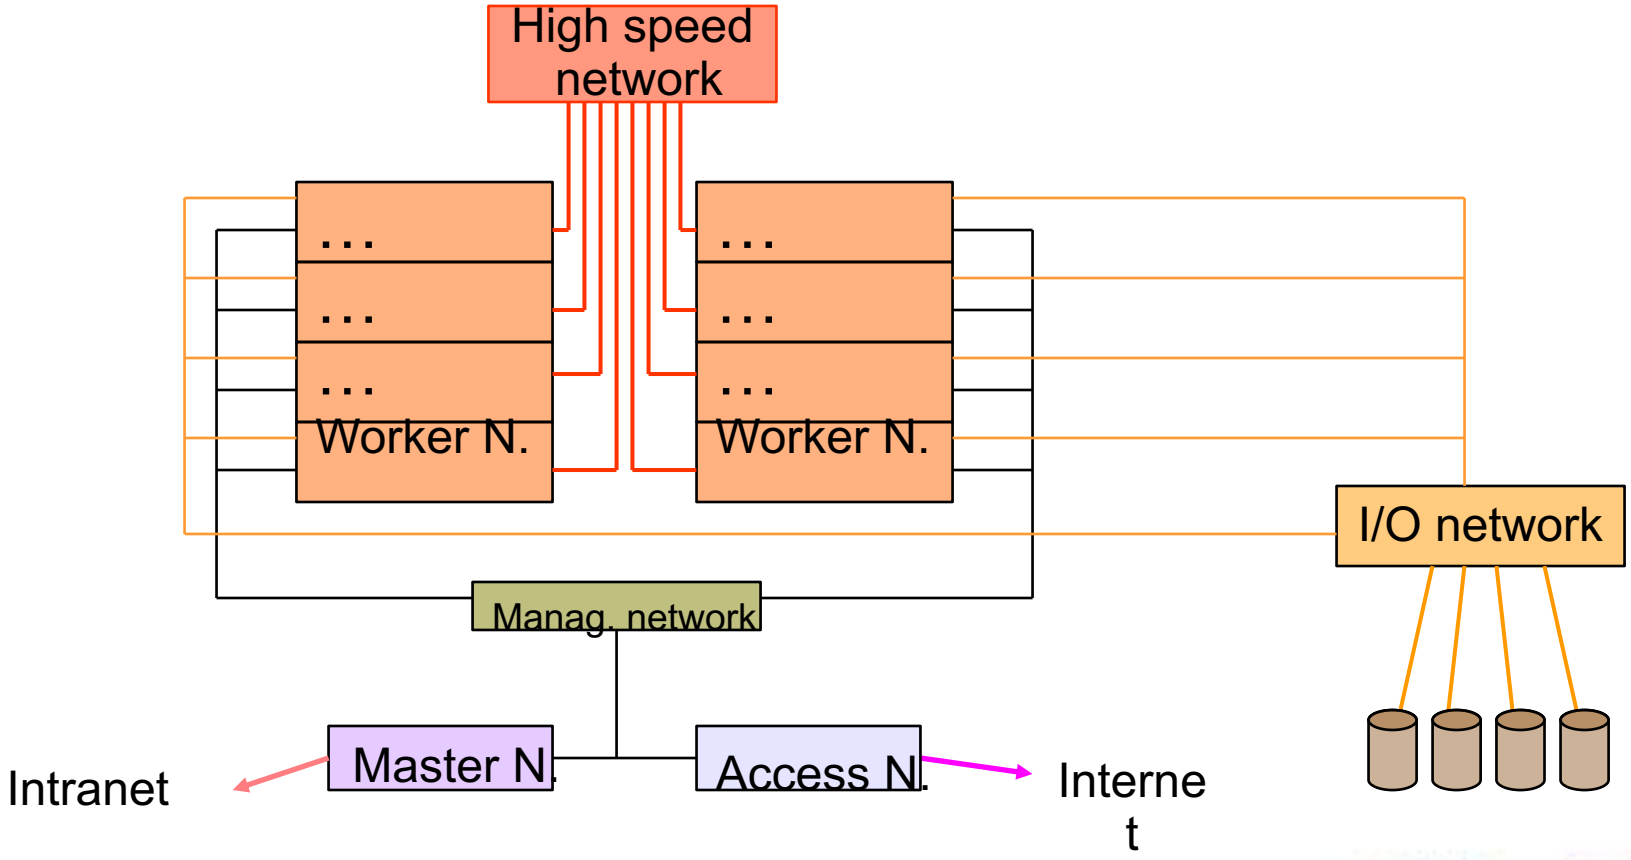

HPC cluster: 3 kind of network

- HIGH SPEED NETWORK
  - parallel computation
  - low latency /high bandwidth
  - Usual choices: Infiniband…
- I/O NETWORK
  - I/O requests (NFS and/or parallel FS)
  - latency not fundamental/ good bandwidth
  - GIGABIT is ok
- Management network
  - management traffic
  - any standard network

MHPC
Master in High Performance Computing

Network in Top500

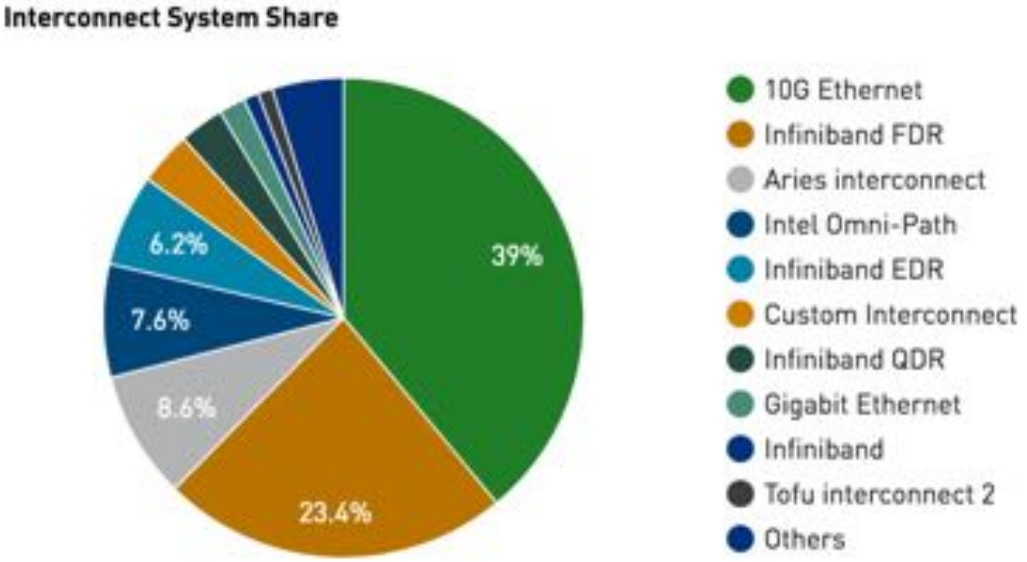

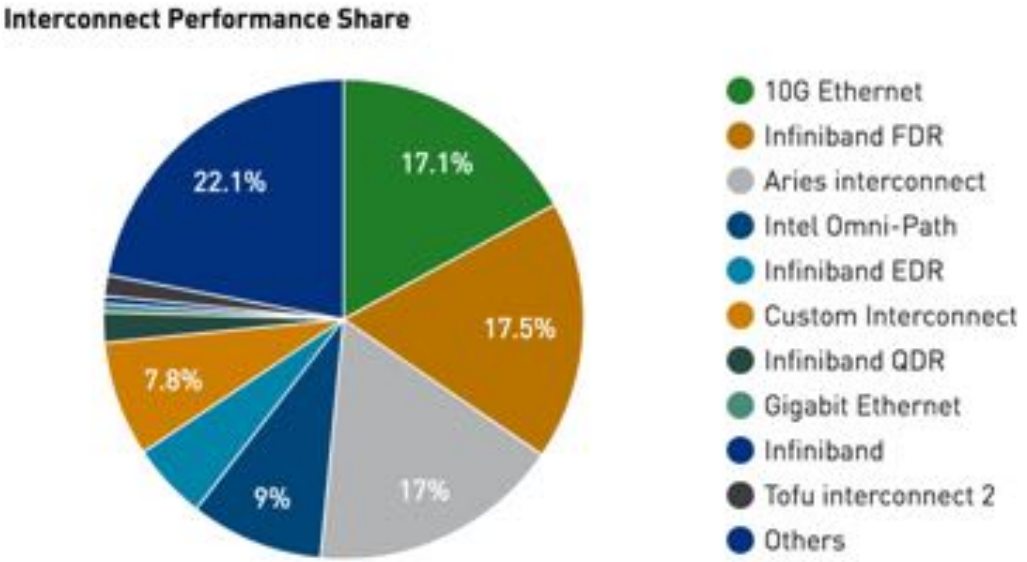

MHPC
Master in High Performance Computing

Accelerators: GPU

- Co-processors or accelerators have been around for a while
- Big burst in its adoption in HPC when Nvidia released CUDA (2006).
- GPGPUs or simply GPUs work in a different way to conventional CPUs. Emphasis on stream processing.
- Acceleration can be significant but depends on application.
- Nvidia market leader with astonishing performance..

SISSA ICTP

MHPC
Master in High Performance Computing

GPU PERFORMANCE COMPARISON

	P100	V100	Ratio
DL Training	10 TFLOPS	120 TFLOPS	12x
DL Inferencing	21 TFLOPS	120 TFLOPS	6x
FP64/FP32	5/10 TFLOPS	7.5/15 TFLOPS	1.5x
HBM2 Bandwidth	720 GB/s	900 GB/s	1.2x
STREAM Triad Perf	557 GB/s	855 GB/s	1.5x
NVLink Bandwidth	160 GB/s	300 GB/s	1.9x
L2 Cache	4 MB	6 MB	1.5x
L1 Caches	1.3 MB	10 MB	7.7x

SISSA ICTP

Accelerators: Intel PHI (MIC)

- Also an accelerator but more similar to a conventional multicore CPU.
- • Cores connected in a ring topology.
- • No need to write CUDA or OpenCL as Intel compilers will compile Fortran or C code for the MIC.

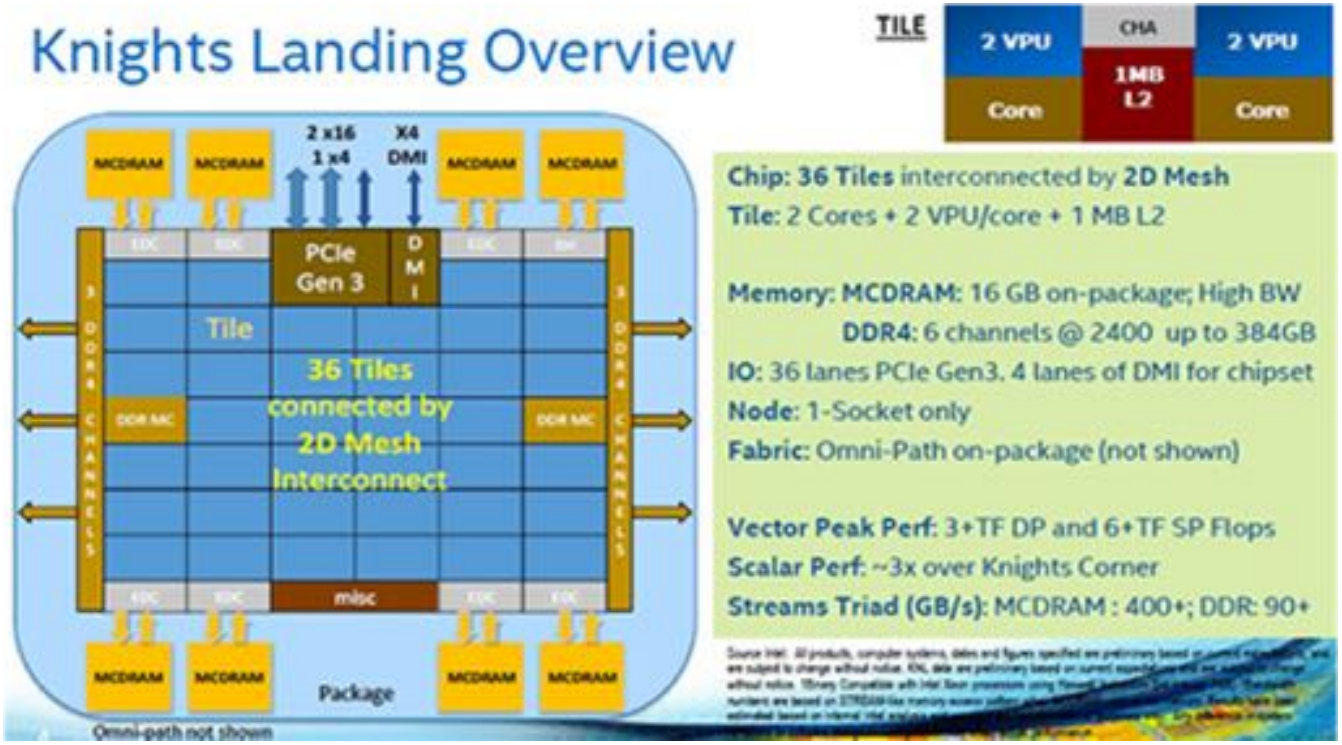

MHPC
Master in High Performance Computing

Accelerators in Top500

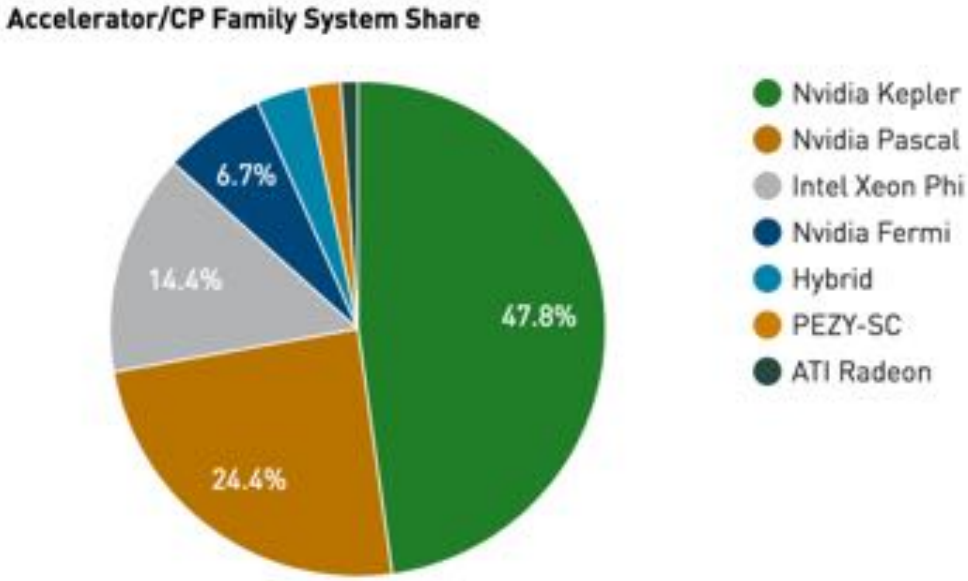

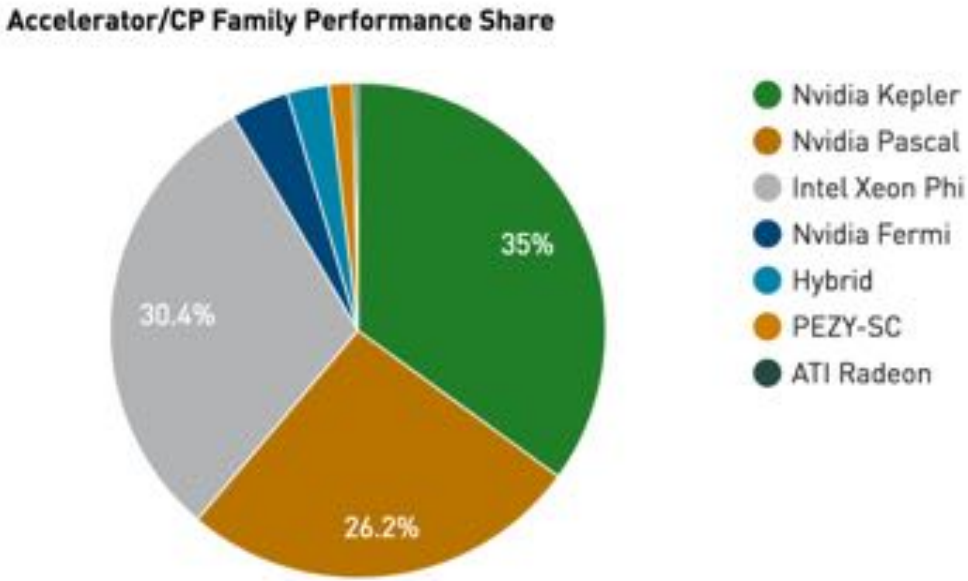

MHPC
Master in High Performance Computing

Why accelerators ?

- GPUs and MIC mandatory in HPC because of high performance and efficiency (i.e. Flops/watt).
- they are mainly to be attached to host CPUs via the PCIe bus (a standard PC-like connection).
- Both device families have limitations:
  - low device memory
  - slow transfer rate via PCIe link
  - difficulty in programming (particularly CUDA).
  - speedup is highly application and data dependent.
- New model are standalone models (e.g Knight's Landing) and/or and with faster connections (Nvlink).

SISSA ICTP

Last but not least: Storage

- High Speed Storage is required for HPC
  - Parallel Filesystem is mandatory:
    - Lustre/GPFS/BeeGFS etc..
- Hierarchical storage is also a solution:
  - Hierarchical storage management (HSM) is a data storage technique, which automatically moves data between high-cost and low-cost storage media.
    - First layer: SSD
    - Second layer : parallel FS
    - Third layer: SAN
    - Fourth layer: Tapes

MHPC
Master in High Performance Computing

Cluster example

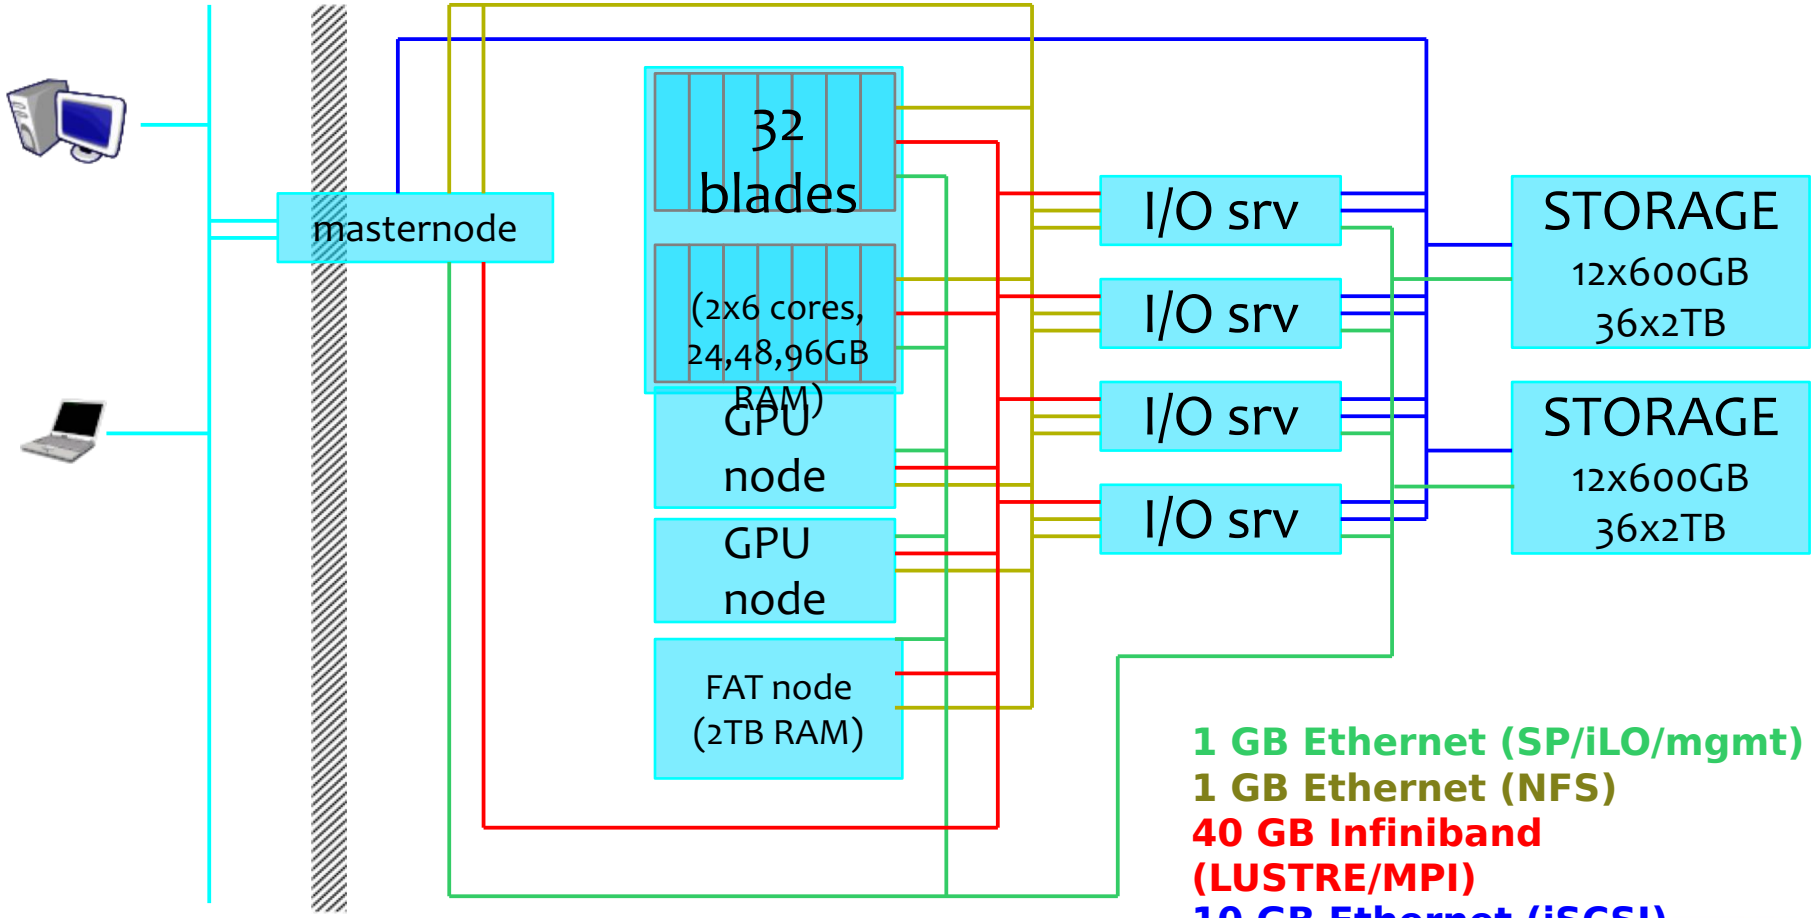

SISSA ICTP

System stack of a supercomputing [from ref 1]

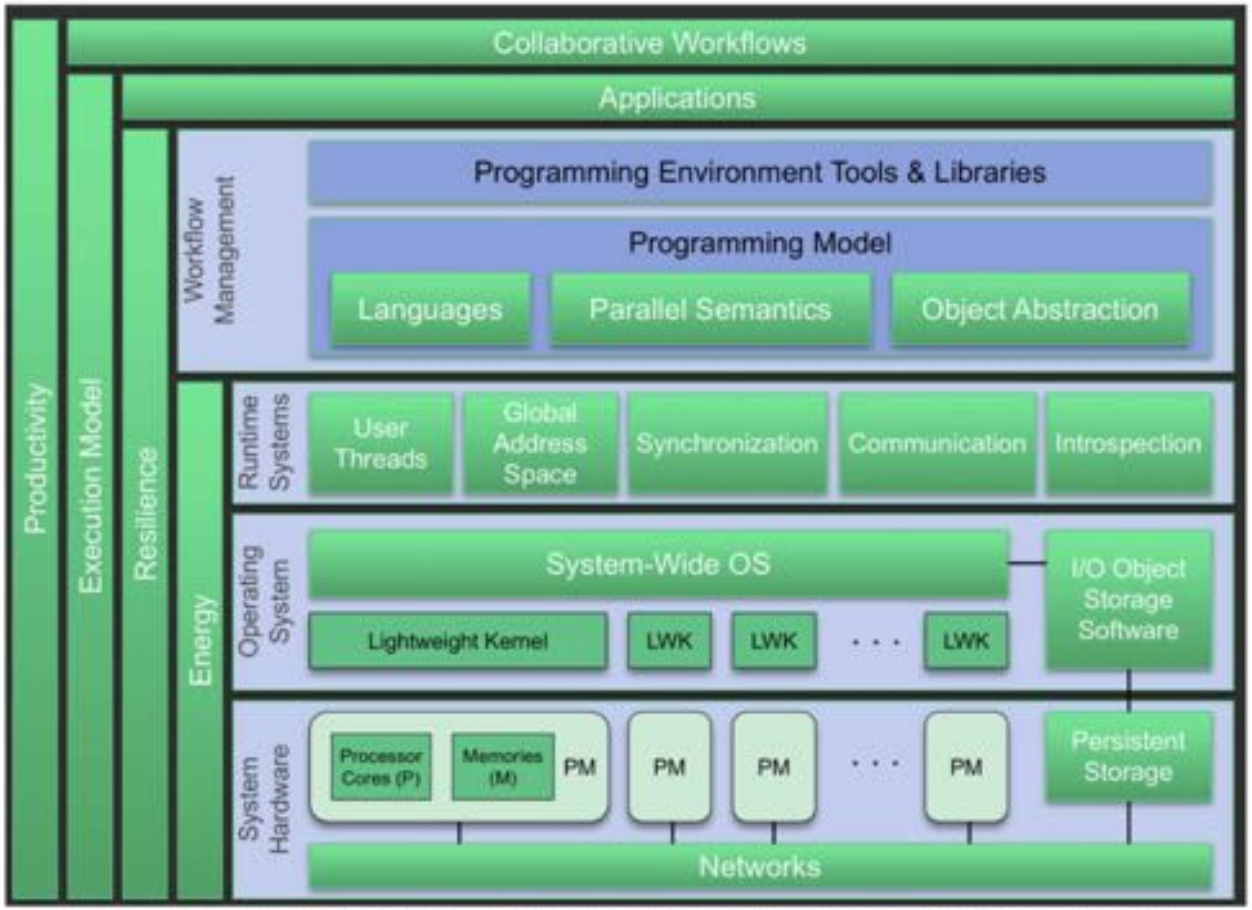

Figure 1.9 The system stack of a general supercomputer consists of a system hardware layer and several software layers. The first software layer is the operating system, encompassing both resource management and middleware to access input/output (I/O) channels. Higher software layers include runtime systems and workflow management.

MHPC
Master in High Performance Computing

HPC platform: the software stacks

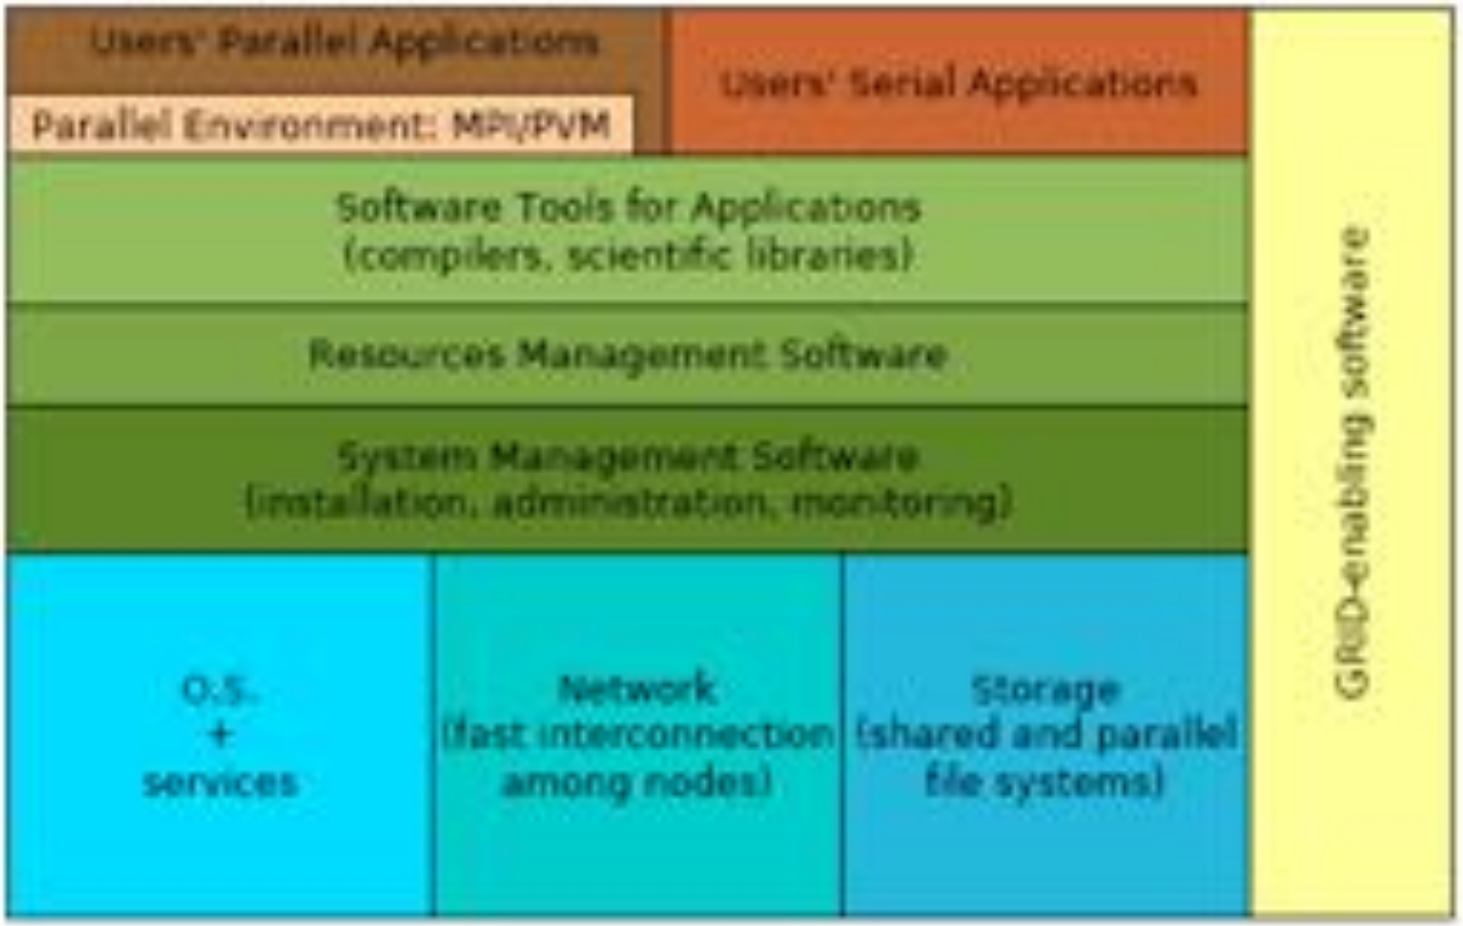

ICTP

the software stacks: sys. adm...

Cluster middleware: Middleware Design Goals

- Complete Transparency (Manageability):
  - Lets us see a single cluster system..
    - Single entry point, ftp, ssh, software loading...
- Scalable Performance:
  - Easy growth of cluster
    - no change of API & automatic load distribution.
- Enhanced Availability:
  - Automatic Recovery from failures
    - Employ checkpointing & fault tolerant technologies

Cluster middleware

Administration software:

- user accounts
- NTP/NFS/ etc…

Resource management and scheduling software (LRMS)

- Process distribution
- Load balance
- Job scheduling of multiple tasks

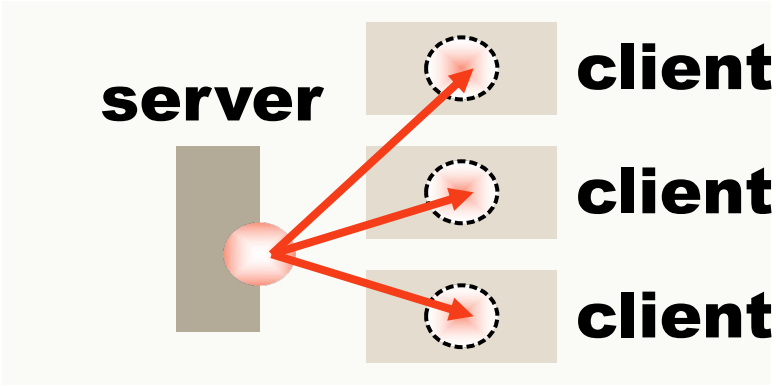

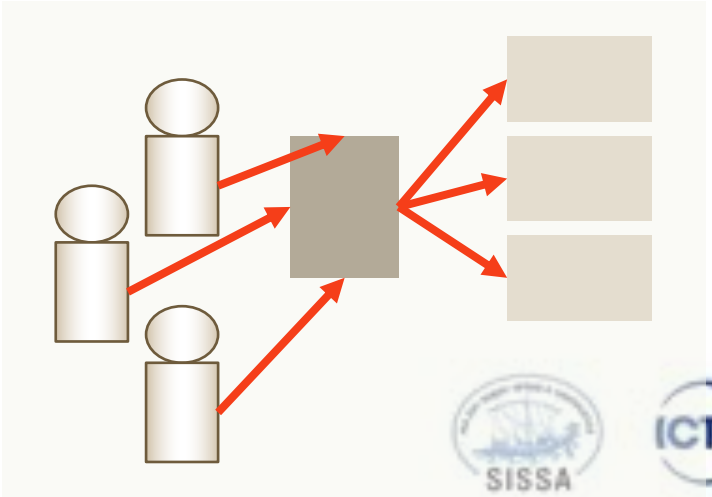

MHPC
Master in High Performance Computing

How much does it cost a computational infrastructure ?

- It is not just a matter of HW…
- Total Cost of Ownership is the right way to calculate the budget for an HPC infrastructure..

SISSA ICTP

MHPC
Master in High Performance Computing

Total Cost of Ownership

- It is the sum of all of the costs that a customer incurs during the lifetime of a technology solution.
- In the High Performance Computing (HPC) field, the Total Cost of Ownership is normally referred to the data center costs.
- Cost to the owner to build, operate and maintain the data center.
- Cost of Services delivered should be computed taking into account TCO.

SISSA ICTP

What should be included in the TCO for HPC ?

- Investment, operation and maintenance costs:
  - Hardware: servers, storage, networking, cabling, etc.
  - Electrical equipment: power distribution units, UPS, generators, etc.
  - Cooling systems: air conditioners, water cooling, etc.
- Infrastructure for the data center, power adaptation issues, etc.
- Energy consumption of the hardware and cooling systems
- Software licenses
- Human resources
- Maintenance

MHPC
Master in High Performance Computing

HOW can I reduce TCO ?

SISSA ICTP

Cloud computing is the answer ?

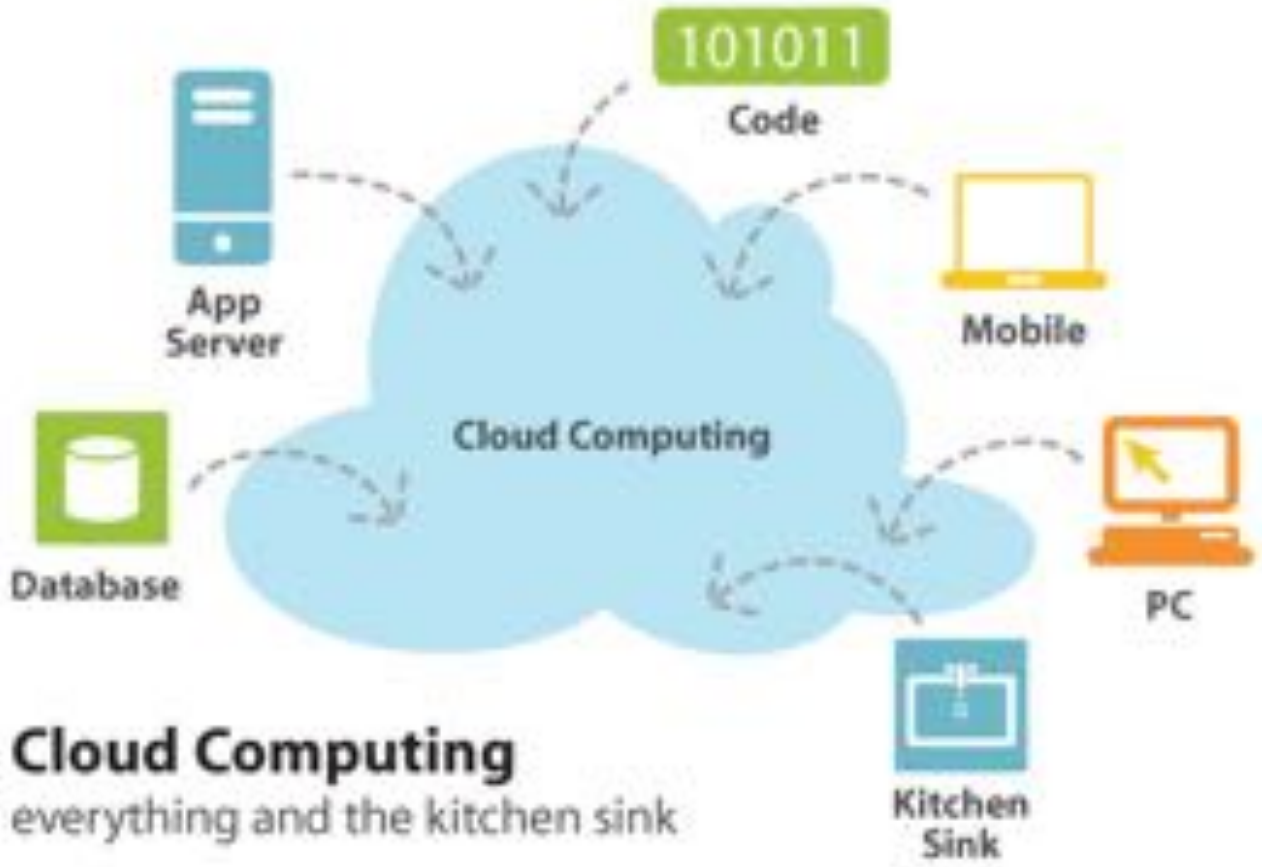

MHPC
Master in High Performance Computing

cloud approach

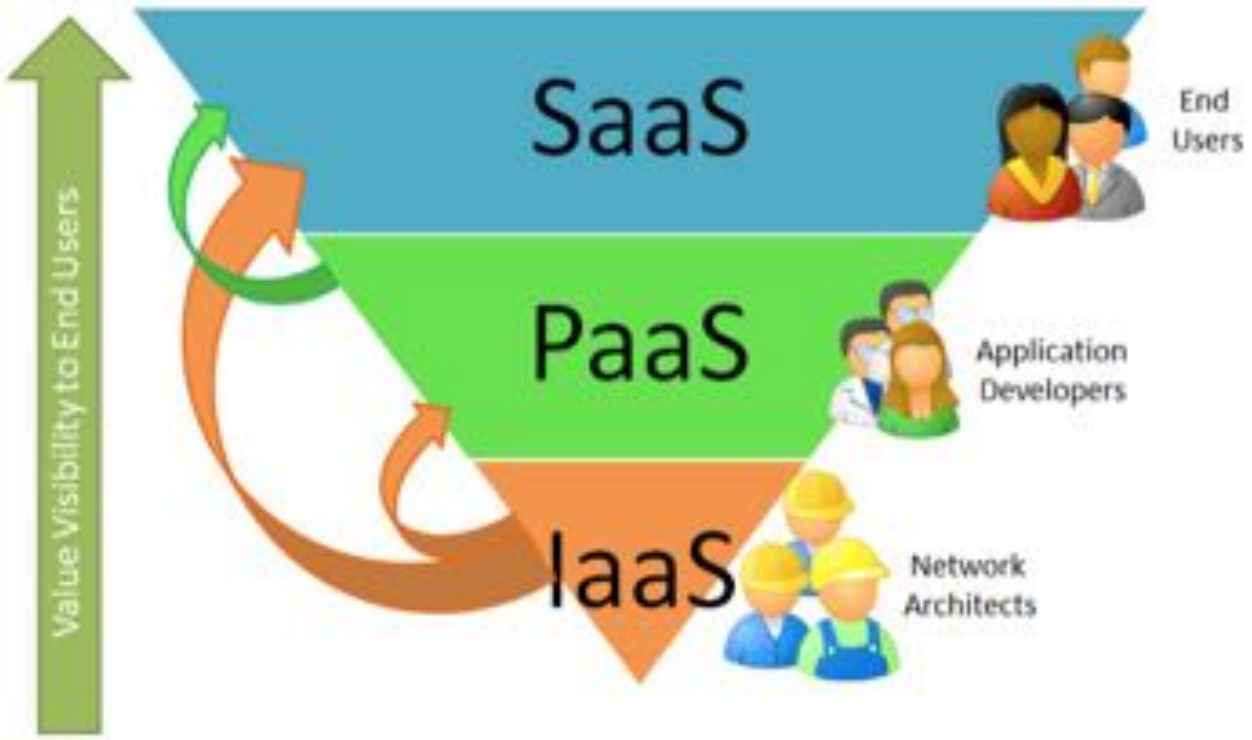

SISSA ICTP

MHPC
Master in High Performance Computing

The dream

- Cloud computing offers almost unlimited storage and instantly available and scalable computing resources,
- All the above at a reasonable metered cost.. (pay per use)
- However…
- the use of a typical cloud needs a bit of care..
- Remote HPC services can range from shared HPC clusters to fully virtualized cloud environments.

SISSA ICTP

Cloud computing and HPC

- *The case for HPC in the cloud is growing stronger, but still has a way to go, especially for the more traditional HPC segments in the public sector.*

- https://www.hpcwire.com/2018/03/15/how-the-cloud-is-falling-short-for-research-computing/

HPC on cloud..

- cloud computing represented about 2% of the HPC market by total revenue in 2016.
- About 35% of HPC users make occasional use of public cloud resources.
- A number of vendors already exist within the industry providing HPC in the cloud solutions.

MHPC
Master in High Performance Computing

HPC cloud providers

- AMAZON WEB SERVICES (AWS)
- GOOGLE CLOUD PLATFORM
- MICROSOFT AZURE
- IBM SPECTRUM COMPUTING
- PENGUIN COMPUTING ON DEMAND (POD)

SISSA ICTP

MHPC
Master in High Performance Computing

Conclusions

- HPC is about performance but not only
- Supercomputers are clusters !
- Clusters have many different components
- Parallel programming is needed to use HPC systems at best
- Several options/tools are available and sometime more than one approach is needed at the same time
- There are a lot of other lectures where all what we discussed in this first lectures will be analyzed in details

SISSA ICTP