

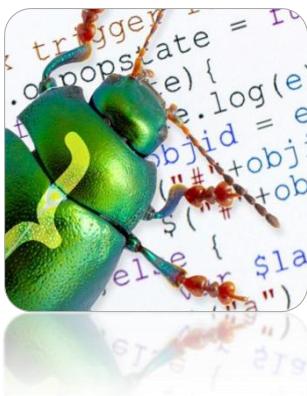
# Foundations of High Performance Computing

Luca Tornatore

Debugging with gdb

*serial, single-threaded, applications*

# Outline



Introducing  
Debugging

GDB  
basic  
How To



GUI  
for GDB



# Introducing debugging



# Smithsonian



It has long been recognized and documented that insects are the most diverse group of organisms, meaning that the numbers of species of insects are more than any other group. In the world, some 900 thousand different kinds of living insects are known. This representation approximates 80 percent of the world's species. The true figure of living species of insects can only be estimated from present and past studies. Most authorities agree that there are more insect species that have not been described (named by science) than there are insect species that have been previously named. Conservative estimates suggest that this figure is 2 million, but estimates extend to 30 million. In the last decade, much attention has been given to the entomofauna that exists in the canopies of tropical forests of the world. From studies conducted by Terry Erwin of the Smithsonian Institution's Department of Entomology in Latin American forest canopies, the number of living species of insects has been estimated to be 30 million. Insects also probably have the largest biomass of the terrestrial animals. At any time, it is estimated that there are some 10 quintillion (10,000,000,000,000,000,000) individual insects alive.

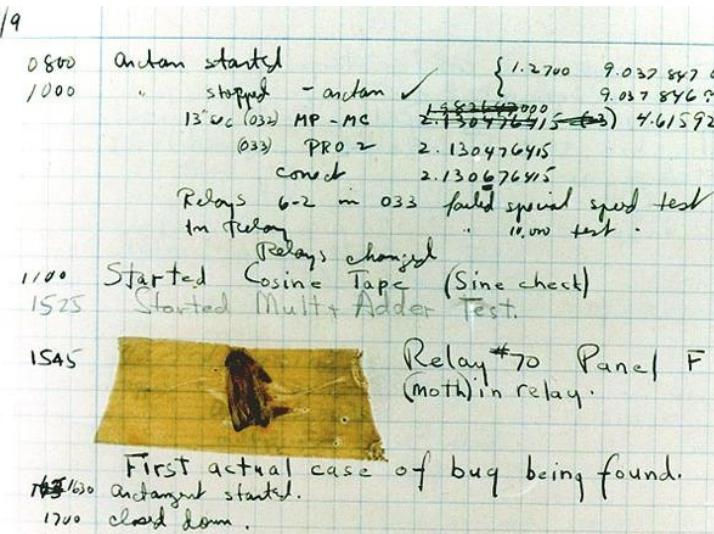


# Introducing debugging

The usage of *bug* and, hence, of *de-bugging* referred to programming is a long-standing tradition, whose origin is difficult to trace back.

An often-told story is about Mark-II calculator located at Harvard: On Sept. 9<sup>th</sup>, 1945, a technician found a moth in a relay that caused a flaw in a “program” execution.

Adm. Grace Hopper is reported to have written in its diary about that as “first actual case of bug being found”



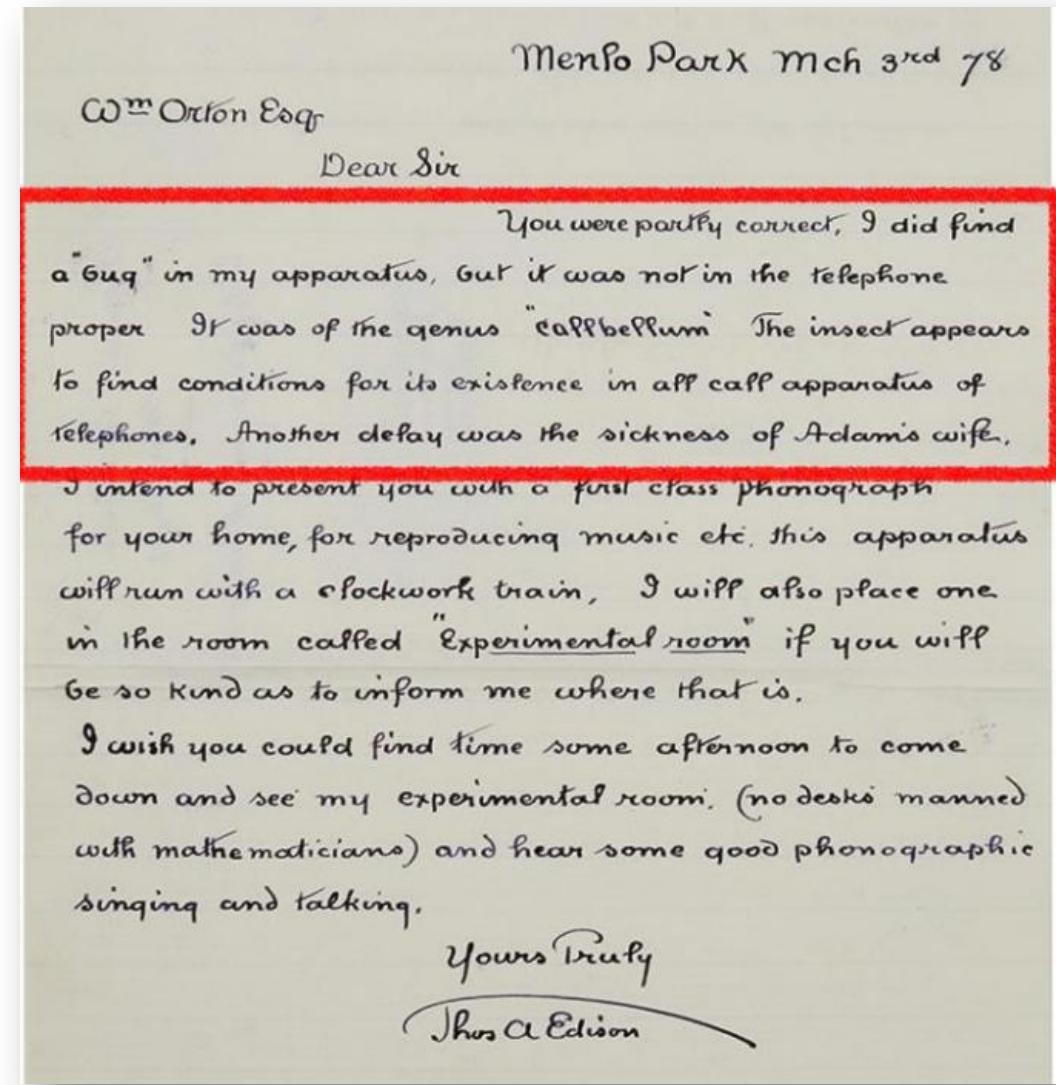


# Introducing debugging

However, Thomas Edison found another actual bug in one of his phones, as he reports in a letter to an associate.

He later wrote:

"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that "**Bugs**"—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."





# Introducing debugging

It is difficult to trace back how this term has been poured into computer programming jargon.

However, already early in the 60s' it was appearing in technical papers without need of explanation.

The immortal I. Asimov used it in a 1944 short robot story "Catch that rabbit", and his incredible influence contributed much to make the term popular.

More funny infos:

[IEEE Annals of the History of Computing](#)  
( Volume: 20 , Issue: 4 , Oct-Dec 1998 )

ask for the PDF if you're interested and do not have access

## Stalking the Elusive Computer Bug

PEGGY ALDRICH KIDWELL

*From at least the time of Thomas Edison, U.S. engineers have used the word "bug" to refer to flaws in the systems they developed. This short word conveniently covered a multitude of possible problems. It also suggested that difficulties were small and could be easily corrected. IBM engineers who installed the ASSC Mark I at Harvard University in 1944 taught the phrase to the staff there. Grace Murray Hopper used the word with particular enthusiasm in documents relating to her work. In 1947, when technicians building the Mark II computer at Harvard discovered a moth in one of the relays, they saved it as the first actual case of a bug being found. In the early 1950s, the terms "bug" and "debug," as applied to computers and computer programs, began to appear not only in computer documentation but even in the popular press.*

### Introduction

S talking computer bugs—that is to say, finding errors in computer hardware and software—occupies and has occupied much of the time and ingenuity of the people who design, build, program, and use computers.<sup>1</sup> Early programmers realized this with some distress. Maurice Wilkes recalls that in about June of 1949,

I was trying to get working my first non-trivial program, which was one for the numerical integration of Airy's dif-

published in the *Annals* in 1981. Here the time is given as the summer of 1945, but the computer is the Mark II, not the Mark I. A photograph shows the moth taped in the logbook, labeled "first actual case of bug being found." Small problems with computers have been called bugs ever since.<sup>5</sup>



# Introducing debugging

"It has been just so in all of my inventions. The **first step** is an intuition and comes with a burst, **then difficulties arise**—this thing gives out and [it is] then that "**Bugs**"—as such little faults and difficulties are called—show themselves and **months of intense watching, study and labor are requisite** before commercial success or failure is certainly reached."

T. Edison

That maybe is too dramatic and emphatic, but the point to get from Edison is that the *de-bugging activity* is an inherent and intrinsic one in software development



# Introducing debugging

Provided that

- you'll have close encounters with bugs in your life;
- de-bugging is a fundamental and unavoidable part of your work;

what is the best way to proceed ?



# Introducing debugging

## 1. Do not insert bugs

*Highly encouraged, but rarely works*



# Introducing debugging

## 1. Do not insert bugs

*Highly encouraged, but rarely works*

## 2. Add `printf` statements everywhere

*Highly discouraged, but sometimes works.*

*In case of memory problems – typically due to pointers chaos – you may see that the problem “disappear”, or changes its appearance, when you insert a new `printf`*



# Introducing debugging

## 1. Do not insert bugs

*Highly encouraged, but rarely works*

## 2. Add `printf` statements everywhere

*Highly discouraged, but sometimes works.*

*In case of memory problems – typically due to a os*

## 3. Use a DEBUGGER

*That is definitely the best choice and, fortunately,  
the subject of this lecture.*

**gdb** is almost certainly the best free, extremely feature-rich command-line debugger ubiquitously available on \*nix systems.

# Outline



Introducing  
Debugging

```
if(Q > 0)
{
    switch(Q)
    {
        case(1): // row y = 0 and/or plane symmetry
            if(subregions[1][BOTTOM][y_] == 0)
                // this subregion in quadrant 1 contains the
                // set-up corners for seed sub-region gen
                SBL[y_] = 0, SBL[x_] = Nmesh - subregion
                STR[y_] = 1, STR[x_] = Nmesh - subregion
                // find horizontal extension in this quad
                Np = STR[x_] - SBL[x_];
                // allocate memory for seeds in this strl
                SEED_y0 = (unsigned int*)malloc(sizeof(unsigned
                if(!internal.mimic_original_seedtable)
                {
                    // ...
                }
                else
                {
                    // ...
                }
            }
        }
}
```

GDB  
How To



GUI  
for GDB



Using  
GDB

# What GDB is useful for

There are 3 basic usages<sup>(\*)</sup> of GDB:

1. Debugging a code  
*best if it has been compiled with -g*
2. Inspecting a code crash through a core file
3. Debugging / inspecting a running code

(\*) Highly advised: learn keyboard commands. Although many GUI exist (we'll see some later), keyboard is still the best productivity tool.

```

if(x > 0)
{
    switch(y)
    {
        case 1: // row y = 0 and/or plane symmetry
            if(skelepred[0][y][0] == 0) // this subregion is quadrant 1 contains
            {
                // set min corners for next sub-region
                SKEL_X[0] = 0, SKEL_X[1] = 0 - nloop;
                SKEL_Y[0] = 0, SKEL_Y[1] = 0 - nloop;

                // find horizontal extension in this qu
                nloop = SKEL_X[1] - SKEL_X[0];
            }
            // allocate memory for nodes to this st
            SKED_Y[0] = -nloop;
            SKED_Y[1] = 0;
            SkeleAlloc[SKED_Y[0]:SKED_Y[1]] = 1;
        break;
        case 2: // internal, minc, original, seetable
            {

```

# Using GDB

# What GDB is useful for

# 1.Debugging a code



# Compiling with dbg infos

In order to include debugging information in your code, you need to compile it with `-g` family options (read the gcc manual for complete info):

**-g**

produce dbg info in O.S. native format

**-ggdb**

produce gdb specific extended info, as much as possible

**-glevel**

default level is 2. 0 amounts to no info, 1 is minimal, 3 includes extra information (for instance, macros expansion) – this allows macro expansion; add **-gdwarf-n** in case, where possible, where *n* is the maximum allowed (4)

**-ggdblevel**

you can combine the two to maximize the amount of useful info generated



Using  
GDB

# Debugging a code : start

You just start your code under gdb control:

```
%> gdb program
```

You may define already the arguments needed by your program at invocation:

```
%> gdb --args program arg1 arg2 ... argN
```

Or you can define the arguments from within the gdb session:

```
%> gdb program
```

```
Reading symbols from program...done.
```

```
(gdb) set args arg1 arg2 ... argN
```

```
(gdb) run
```

► live demo with [gdb\\_try\\_args.c](#)



```
GNU gdb (Ubuntu 8.0-0ubuntu1) 8.0.0.20140412
Copyright (C) 2013 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change it and/or redistribute it
under the terms of the GNU General Public License version 3 or later.
There is NO WARRANTY, to the extent permitted by law.
Type "show copying" and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.

```

```
break main
Breakpoint 1 at 0x401010: file main.c, line 10.

```

```
run
Starting program: /home/robin/Downloads/segmentation/main
[Thread debugging using libthread_db enabled]
Using host libthread_db library "/lib/x86_64-linux-gnu/libthread_db.so.1".

```

```
Breakpoint 1, 0x000000401010 in main ()
```

Using  
GDB

# Debugging a code : run and stop

You may just want the code to run, for instance to reach the point of a seg fault:

```
(gdb) run
```

Or, you may want to stop it from the beginning to have full control of each step:

```
(gdb) break main
(gdb) run
```

Or you may already know what is the problematic point to stop at:

```
(gdb) break location
(gdb) run
```



Using  
GDB

# Debugging a code : run and stop

**Breakpoints** are a key concept in debugging. They are stopping point at which the execution interrupts and the control is given back to you, so that you can inspect the memory contents (variables values, registers values, ... ) or follow the subsequent execution step by step.

You can define a breakpoint in several way

(gdb) break

(gdb) break  $\pm$ offset

(gdb) filename:linenum

(gdb) break functionname

insert a break at the current pos

insert a break *offset* lines after/before the current line

insert a break at *linenum* of file *filename*

insert a break at the entry point of function *functionname*

There are more options,  
just check the manual



Using  
GDB

# Debugging a code : run and stop

A break point may be defined as dependent on a given condition:

```
(gdb) break my_function if (arg1 > 3 )
```

This sets a breakpoint at function *my\_function* : the condition will be evaluated each time the point is reached, and the execution is stopped only if it is true.

Condition can be any valid expression.

```
(gdb) info break
```

gives you informations on active breakpoints.

```
(gdb) delete [n]
      clear [location]
      <disable | enable>   see the manual
```

```
[1000x1000]
```

```
using namespace std;
```

```
void f1() { // row i = 0 and/or plane separator
```

```
if (neighborCount(MINIMUS, i, j) == 4)
```

```
{ this->neighborCount[quadrant][i][j] = 4;
```

```
SHL[x][j] = 0, SHL[x][j] = max - neighborCount(MINIMUS, i, j);
```

```
SHR[x][j] = 0, SHR[x][j] = max - neighborCount(MAXIMUS, i, j);
```

```
if (row == 0 & plane == 0) { // boundary condition in this case
```

```
RP = SHL[x][j] + SHR[x][j];
```

```
LP = SHR[x][j] - RP; LP = abs(LP);
```

```
SHL[x][j] = (SHL[x][j] * max) / LP;
```

```
SHR[x][j] = (SHR[x][j] * max) / LP;
```

```
SHL[x][j] = (SHL[x][j] * max) / LP;
```

```
SHR[x][j] = (SHR[x][j] * max) / LP;
```

Using  
GDB

# Debugging a code : run and stop

You can define a list of commands to be executed when a given breakpoint is reached:

```
(gdb) break my_function if (arg1 > 3 )  
Breakpoint 1 at 0x.....: file blabla.c, line 42  
(gdb) command 1
```

Type commands for when breakpoint 1 is hit, one per line.

End with a line saying just "end".

```
> print arg1  
> print another_useful_variable  
> x/10wd a_global_integer_array
```



Using  
GDB

# Debugging a code : run and stop

When you have the control of the program execution, you can decide how to proceed:

(gdb) cont [c]

(gdb) cont *count-ignore*

(gdb) next [n] | *count*  
nexti

(gdb) step [s] | *count*  
stepi

(gdb) until [u] | *count*

(gdb) advance *location*

continue until the end / next stop

continue ignoring the next *count-ignore* stops (for instance, occurrences of a bp)

continue to the next src line *in the current stack frame*

continue to the next src line

continue until a src line past the current one is reached in the current stack fr.

continue until the specified location is reached

► live demo with **gdb\_try\_breaks.c**



```
1000 > n
1000 |
```

Using  
GDB

# Debugging a code : run and stop

You can also rewind your execution step-by-step

```
(gdb) reverse-continue [rc]
(gdb) reverse-step [count]
      reverse-stepi
(gdb) reverse-next [count]
      reverse-nexti
(gdb) set exec-direction <verse | forward >
```

```
void() {  
    int i, j;  
    for (i = 0; i < N; i++) {  
        for (j = 0; j < N; j++) {  
            if (i == j) {  
                M[i][j] = 1;  
            } else if (i + j == N - 1) {  
                M[i][j] = 1;  
            } else {  
                M[i][j] = 0;  
            }  
        }  
    }  
}
```

Using  
GDB

# Debugging a code : source list

Often, when you are debugging, you may have the need of looking at either the source lines or at the generated assembler:

(gdb) list *linenum*

print src lines around line *linenum* of the current source file

(gdb) list *function*

print the source lines of *function*

(gdb) list *location*

print src lines around *location*

(gdb) set listszie *count*

control the number of src lines printed

► live demo with **gdb\_try\_breaks.c**



Using  
GDB

# Debugging a code : stack examination

Examining the stack is often of vital importance. With GDB you can have a quick and detailed inspection of all the stack frames.

(gdb) backtrace [args]

*n*

*-n*

*full*

where, info stack

print the backtrace of the whole stack

print only the *n* innermost frames

print only the *n* outermost frames

print local variables value, also

additional aliases

► live demo with **gdb\_try\_breaks.c**



Using  
GDB

# Debugging a code : memory examination

Accessing to the content of memory is a fundamental ability of a debugger. You have several different ways to do that:

(gdb) print *variable*  
p/F *variable*

print the value of *variable*  
print *variable* in a different format  
(x, d, u, o, t, a, c, f, s)  
see then manual for advanced location

(gdb) x/FMT *address*

Explore memory starting at address  
*address*

-> see at live demo how to use this

(gdb) display *expr*  
display/fmt *expr*  
display/fmt *addr*

Add *expr* to the list of expressions to  
display each time your program stops

► live demo with **gdb\_try\_breaks.c**



Using  
GDB

# Debugging a code : memory examination

Memory can be searched to find a particular value, of a given size

```
(gdb) find [/sn] start,  
end, val1 [,val2, ...]
```

```
      find [/sn] start,  
+len, val1 [,val2, ...]
```

Search memory for a particular sequence of bytes.  
*s* is the size of type to be searched  
*n* is the max number of occurrences

```
[1000x1000]
```

```
main() {  
    int i, j, k;  
    int x[10];  
    int y[10];  
    int z[10];  
    ...  
    for (i = 0; i < 10; i++) {  
        x[i] = i * 10;  
        y[i] = i * 10;  
        z[i] = i * 10;  
    }  
    ...  
    for (i = 0; i < 10; i++) {  
        for (j = 0; j < 10; j++) {  
            for (k = 0; k < 10; k++) {  
                if (i == j || j == k || k == i) {  
                    x[i] = y[j] + z[k];  
                }  
            }  
        }  
    }  
    ...  
    for (i = 0; i < 10; i++) {  
        printf("%d\n", x[i]);  
    }  
}
```

Using  
GDB

# Debugging a code : registers examination

Examining registers may be also useful (although it's something that only quite advanced users can conceive)

```
(gdb) info registers  
(gdb) info vector  
(gdb) print $rsp  
(gdb) x/10wd $rsp  
(gdb) x/10i $rip
```

print the value of all registers  
print the content of vector registers  
print value of the stack pointer  
print values of the first 10 4-bytes integers on the stack  
print the next 10 asm instructions

► live demo with **gdb\_try\_breaks.c**



Using  
GDB

# Debugging a code : macro expansion

If there are macros in your code, they can be expanded, provided that you compiled the code with the appropriate option:

`-g3 [gdb3] [-dwarf-4]`

(gdb) macro expand *macro*

shows the expansion of macro *macro*;  
*expression* can be any string of tokens

(gdb) info macro [-a|-all]  
*macro*

shows the current (or all) definition(s) of  
*macro*

(gdb) info macros *location*

shows all macro definitions effective at  
*location*



Using  
GDB

# Debugging a code : watch points

You can set *watchpoints* (aka “keep an eye on this and that”) instead of breakpoints, to stop the execution whenever a value of an expression / variable/memory region changes

(gdb) `watch variable`

(gdb) `watch expression`

(gdb) `watch -l expression`

(gdb) `watch -l expression`  
[mask *maskvalue*]

(gdb) `rwatch [-l] expr`  
[mask *mvalue*]

keep an eye onto *variable*

stops when the value of expression changes (the scope of variables is respected)

Interpret *expression* as a memory location to be watched

a mask for memory watching: specifies what bits of an address should be ignored (to match more addresses)

stops when the value of *expr* is read

```
[1000x1000]
```

```
case(1): // row i = 0 and plane connects  
if (neighborCount(MINIMUS, x_1) == 4)  
{  
    this->neighbor_in_quadrant_3 = true;  
  
    MIN(x_1) = 0, MIN(x_2) = max - neighborCount(MINIMUS, x_1);  
    MAX(x_1) = max, MAX(x_2) = max - neighborCount(MINIMUS, x_1);  
    if (neighborCount(MAXIMUS, x_1) > 0)  
        max = MIN(x_1) + neighborCount(MAXIMUS, x_1);  
    else  
        max = MAX(x_1) - neighborCount(MAXIMUS, x_1);  
  
    tdiff_yt = (max - min) / neighborCount(MAXIMUS, x_1);  
}  
else  
    tdiff_yt = (max - min) / neighborCount(MINIMUS, x_1);
```

Using  
GDB

# What GDB is useful for

## 2. Inspecting a code crash through a core file



Using  
GDB

# After a crash : having the core file

It happens that you have code crashes in conditions not easily reproducible when you debug the code itself, for a number of reasons.

However, the O.S. can dump the entire “program status” on a file, called the *core file*:

```
luca@GGG:~/code/tricks% ./gdb_try_watch
no arguments were given, using default: 100

something wrong at point 2
[1]    8435 segmentation fault (core dumped)  ./gdb_try_watch
luca@GGG:~/code/tricks% ls -l core
-rw----- 1 luca luca 413696 nov  8 15:45 core
luca@GGG:~/code/tricks% 
```

In order to allow it to dump the core, ou have to check / set the core file size limit:

```
%> ulimit -c [size limit in KB]
```

```
void() {  
    int i, j;  
    for (i = 0; i < 8; i++) {  
        for (j = 0; j < 8; j++) {  
            if (i == 0 || i == 7 || j == 0 || j == 7) {  
                cout << "X";  
            } else if (i + j == 7) {  
                cout << "O";  
            } else if (i + j == 14) {  
                cout << "X";  
            } else {  
                cout << " ";  
            }  
        }  
        cout << endl;  
    }  
}
```

Using  
GDB

# After a crash: inspect the core file

Once you have a core, you can inspect it with GDB

```
%> gdb executable_name core_file_name
```

Or

```
%> gdb executable_name  
(gdb) core core_file_name
```

The first thing to do, normally, is to unwind the stack frame to understand where the program crashed:

```
(gdb) bt full
```

► live demo with [gdb\\_try\\_watch.c](#)

```
[1000x1000] 1
using(gdb)

```

Using  
GDB

# What GDB is useful for

## 3. Debugging a running process

```
mpirun -n 4 ./matrixmult
Using GDB
```

# Running processes: attaching

In order to debug a running process, you can simply attach gdb to it:

```
%> gdb
(gdb) attach process-id
```

and start searching it to understand what is going on

► live demo with **[gdb\\_try\\_attach.c](#)**

# Outline



Introducing  
Debugging



GDB  
How To



GUI  
for GDB



# GDB graphical user interfaces

1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS
5. DDD
6. NEMIVER
7. ECLIPSE, NETBEANS, CODEBLOCKS



# GDB built-in tui

You can start gdb with a text-user-interface:

```
%> gdb -tui
```

Or you can activate/deactivate it from gdb itself:

(gdb) ctrl-x a	change focus
(gdb) ctrl-x o	shows assembly windows
(gdb) ctrl-x 2	display src and commands
(gdb) layout src	display assembly and commands
asm	display src, asm and commands
split	display registers window
regs	



# GDB built-in tui

```
gdb_try_breaks.c
372  {
373
B+> 374      if ( argc > 1 )
375          // arg 0 is the name of the program itself
376          {
377              printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378              for ( int i = 1; i < argc; i++ )
379              {
380                  printf( "\targument %d is : %s\n", i, *(argv+i) );
381              }
382              printf( "\n" );
383          }
384
385      else
386
387          printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );
388
389
390      int arg1;
391
392      if ( argc > 1 )
393          arg1 = atoi( *(argv+1) );
394
395      else
396          arg1 = DEFAULT_ARG1;
397
398      int ret;
399
400      ret = function 1( arg1 );
```

```
native process 8943 In: main
(gdb) l
360     in /home/luca/code/tricks/gdb_try_breaks.c
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks

Breakpoint 1, main (argc=1, argv=0x7fffffffdaa8) at gdb_try_breaks.c:374
(gdb) 
```

L374 PC: 0x555555554d38



```
gdb_try_breaks.c
B+> 374      if ( argc > 1 )
375          // arg 0 is the name of the program itself
376          {
377              printf( "\nexploring my %d argument%c:\n", argc-1, (argc>2)?'s':' ' );
378              for ( int i = 1; i < argc; i++ )
379              {
380                  printf( "\targument %d is : %s\n", i, *(argv+i) );
381              }
382              printf( "\n" );
383          }
384      else
385
386
387      printf ( "no arguments were given, using default: %d\n\n", DEFAULT_ARG1 );
B+> 0x555555554d38 <main+15>    cmpl    $0x1,-0x14(%rbp)
0x555555554d3c <main+19>    jle     0x555555554db7 <main+142>
0x555555554d3e <main+21>    cmpl    $0x2,-0x14(%rbp)
0x555555554d42 <main+25>    jle     0x555555554d4b <main+34>
0x555555554d44 <main+27>    mov     $0x73,%edx
0x555555554d49 <main+32>    jmp     0x555555554d50 <main+39>
0x555555554d4b <main+34>    mov     $0x20,%edx
0x555555554d50 <main+39>    mov     -0x14(%rbp),%eax
0x555555554d53 <main+42>    sub     $0x1,%eax
0x555555554d56 <main+45>    mov     %eax,%esi
0x555555554d58 <main+47>    lea     0x2bf(%rip),%rdi      # 0x5555555501e
0x555555554d5f <main+54>    mov     $0x0,%eax
0x555555554d64 <main+59>    callq   0x555555554680 <printf@plt>
0x555555554d69 <main+64>    movl    $0x1,-0xc(%rbp)
native process 9102 In: main                                         L374  PC: 0x555555554d38
(gdb) break main
Breakpoint 1 at 0xd38: file gdb_try_breaks.c, line 374.
(gdb) r
Starting program: /home/luca/code/tricks/gdb_try_breaks

Breakpoint 1, main (argc=1, argv=0xffffffffdaa8) at gdb_try_breaks.c:374
(gdb) layout split
(gdb) 
```



# GUI for GDB

## GDB dashboard

<https://github.com/cyrus-and/gdb-dashboard>

 live demo



GUI for  
GDB

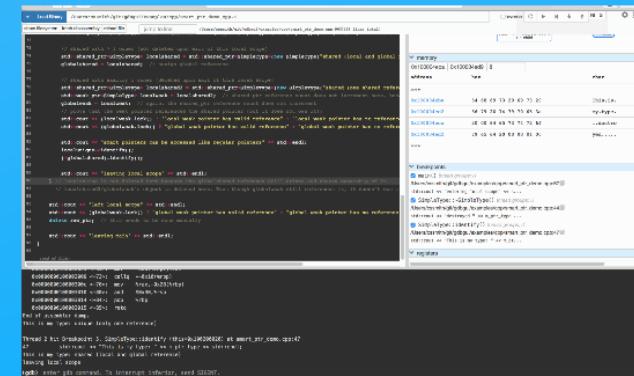
# GDBGUI

[gdbgui](#) Download Docs Examples Screenshots Videos About Chat GitHub



Browser-based debugger for C, C++, go, rust, and more  
gdbgui turns this

```
>>> gdb
GNU gdb (Ubuntu 7.11.1-0ubuntu1-16.5) 7.11.1
Copyright (C) 2016 Free Software Foundation, Inc.
License GPLv3+: GNU GPL version 3 or later <http://gnu.org/licenses/gpl.html>
This is free software: you are free to change and redistribute it.
There is NO WARRANTY, to the extent permitted by law. Type "show copying"
and "show warranty" for details.
This GDB was configured as "x86_64-linux-gnu".
Type "show configuration" for configuration details.
For bug reporting instructions, please see:
<http://www.gnu.org/software/gdb/bugs/>.
Find the GDB manual and other documentation resources online at:
<http://www.gnu.org/software/gdb/documentation/>.
For help, type "help".
Type "apropos word" to search for commands related to "word".
>>>
```



[Download Now »](#)

[View demo](#)

<https://gdbgui.com/>

[▶ live demo](#)



# GDB graphical user interfaces

1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS GUD
5. DDD
6. NEMIVER
7. ECLIPSE, NETBEANS, CODEBLOCKS



# GDB graphical user interfaces

1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS GUD
5. DDD
6. NEMIVER
7. ECLIPSE, NETBEANS, CODEBLOCKS



# GDB graphical user interfaces

1. GDB text-user-interface
2. GDB DASHBOARD
3. GDBGUI
4. EMACS GUD
5. DDD
6. NEMIVER
7. ECLIPSE, NETBEANS, CODEBLOCKS